

# GOAL-BASED AGENTS: SOLVING PROBLEMS BY SEARCHING

## CHAPTER 3, SECTIONS 1-4

## Problem solving agents

Form of **goal-based agent** that formulates the problem of **reaching a goal** in its environment, searches for a **sequence of actions** solving the problem, and executes it.

Assumptions about the task environment:

- ◇ static
- ◇ single agent
- ◇ deterministic
- ◇ fully observable

*do not have to constantly interacting with environment*


**offline** (or open-loop) problem solving is suitable under those assumptions; the entire sequence solution can be **executed “eyes closed.”**

*as you can surely succeed given the correct model.*

# Outline

- ◇ Problem formulation
- ◇ Problem formulation examples
- ◇ Tree search algorithm
- ◇ Uninformed search strategies

## Example: Romania



## Example: Romania

On holiday in Romania; currently in Arad.

Flight leaves tomorrow from Bucharest

### Formulate problem:

initial state: in Arad

goal: be in Bucharest

states: various cities

actions: drive between cities

### Search for a solution:

sequence of drive actions or equivalently (in this case)

sequence of cities, e.g., Arad, Sibiu, Fagaras, Bucharest

## Problem formulation

A **problem** is defined by four items:

**initial state** e.g., “at Arad”

**successor function**  $S(x)$  = set of action–state pairs  
e.g.,  $S(\text{Arad}) = \{ \langle \underset{\text{action}}{\text{Arad}} \rightarrow \underset{\text{state}}{\text{Zerind}}, \text{Zerind} \rangle, \dots \}$

**goal test**, can be

**explicit**, e.g.,  $x = \text{“at Bucharest”}$

**implicit**, e.g.,  $\text{HasAirport}(x)$


**path cost** (additive)

e.g., sum of distances, number of actions executed, etc.

$c(x, a, y)$  is the **step cost**, assumed to be  $\geq 0$

A **solution** is a sequence of actions  
leading from the initial state to a goal state

## Example: Vacuum world




states??

actions??

## goal test??

path cost??

# Example: Vacuum world



states??: dirt presence in each room and robot location (ignore dirt amounts)


actions??

goal test??

path cost??



## Example: Vacuum world




states??: dirt presence in each room and robot location (ignore dirt amounts)

actions??: *Left, Right, Suck, NoOp*

goal test??

path cost??

## Example: Vacuum world




states??: dirt presence in each room and robot location (ignore dirt amounts)

actions??: *Left, Right, Suck, NoOp*

goal test??: = no dirt

path cost??

## Example: Vacuum world



states??: dirt presence in each room and robot location (ignore dirt amounts)

actions??: *Left*, *Right*, *Suck*, *NoOp*

goal test??: = no dirt

path cost??: 1 per action (0 for *NoOp*)

## Selecting a state space

Real world is absurdly complex

⇒ state space must be **abstracted** for problem solving // extract only elements

(Abstract) state = set of real states

(Abstract) action = complex combination of real actions

e.g., “Arad → Zerind” represents a complex set  
of possible routes, detours, rest stops, etc.

For guaranteed realizability, **any** real state “in Arad”  
must get to **some** real state “in Zerind”

(Abstract) solution =

set of real paths that are solutions in the real world

Abstraction should be “easier” than the original problem!

important for solving  
the problem.

## Example: The 8-puzzle

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State

states??

actions??

goal test??

path cost??

// just consider the action of the blank space

## Example: The 8-puzzle

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State

states??: integer locations of tiles (ignore intermediate positions)

actions??

goal test??

path cost??

## Example: The 8-puzzle

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State

states??: integer locations of tiles (ignore intermediate positions)

actions??: move blank left, right, up, down (ignore unjamming etc.)

goal test??

path cost??

## Example: The 8-puzzle

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State

states??: integer locations of tiles (ignore intermediate positions)

actions??: move blank left, right, up, down (ignore unjamming etc.)

goal test??: = goal state (given)

path cost??



## Example: The 8-puzzle

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State

states??: integer locations of tiles (ignore intermediate positions)

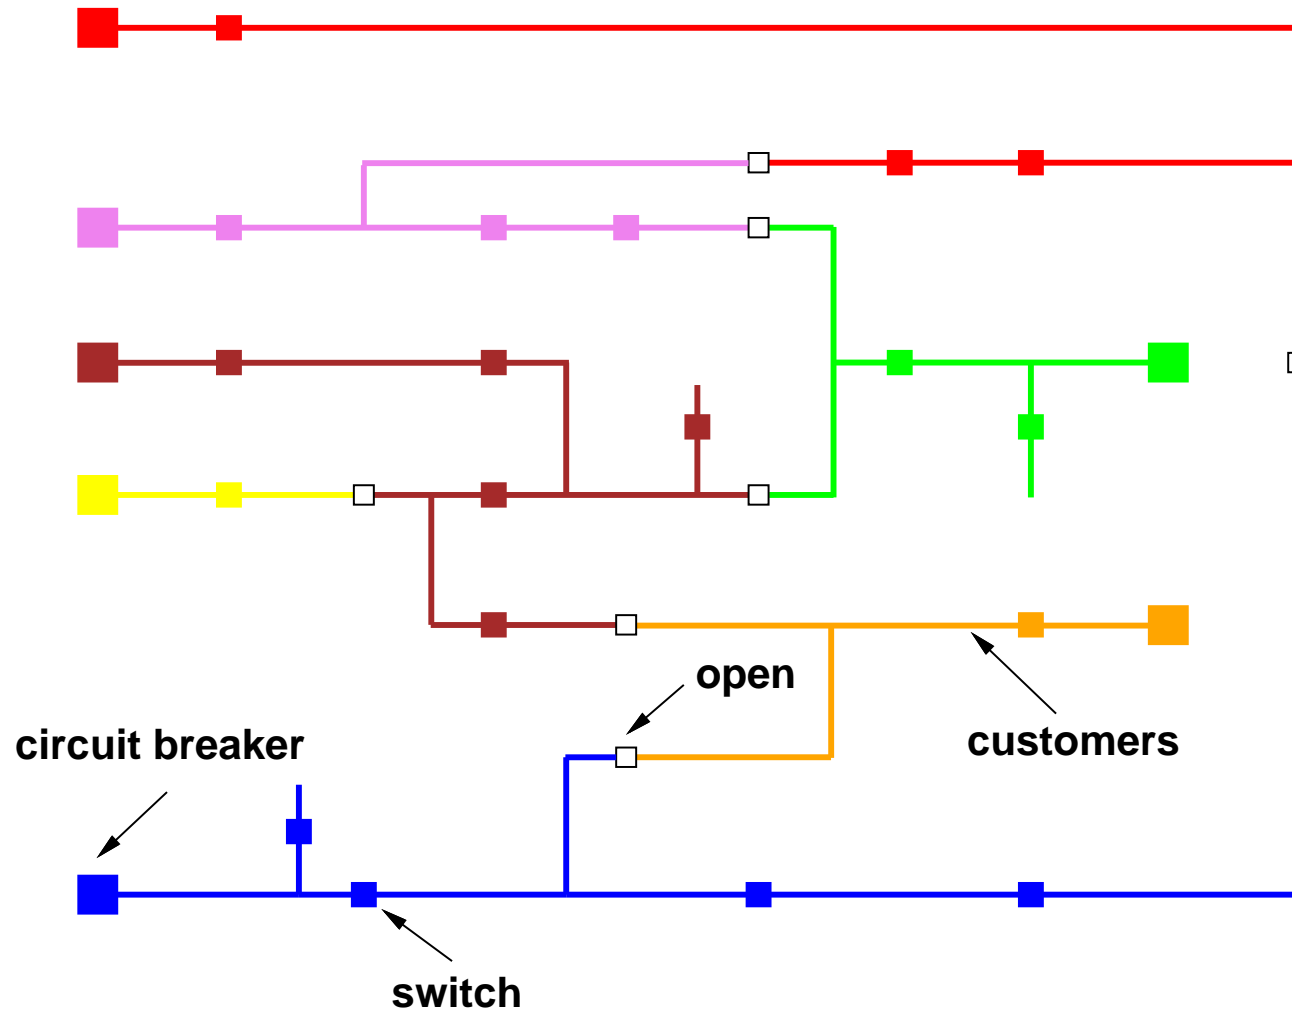
actions??: move blank left, right, up, down (ignore unjamming etc.)

goal test??: = goal state (given)


path cost??: 1 per move

[Note: optimal solution of  $n$ -Puzzle family is NP-hard]

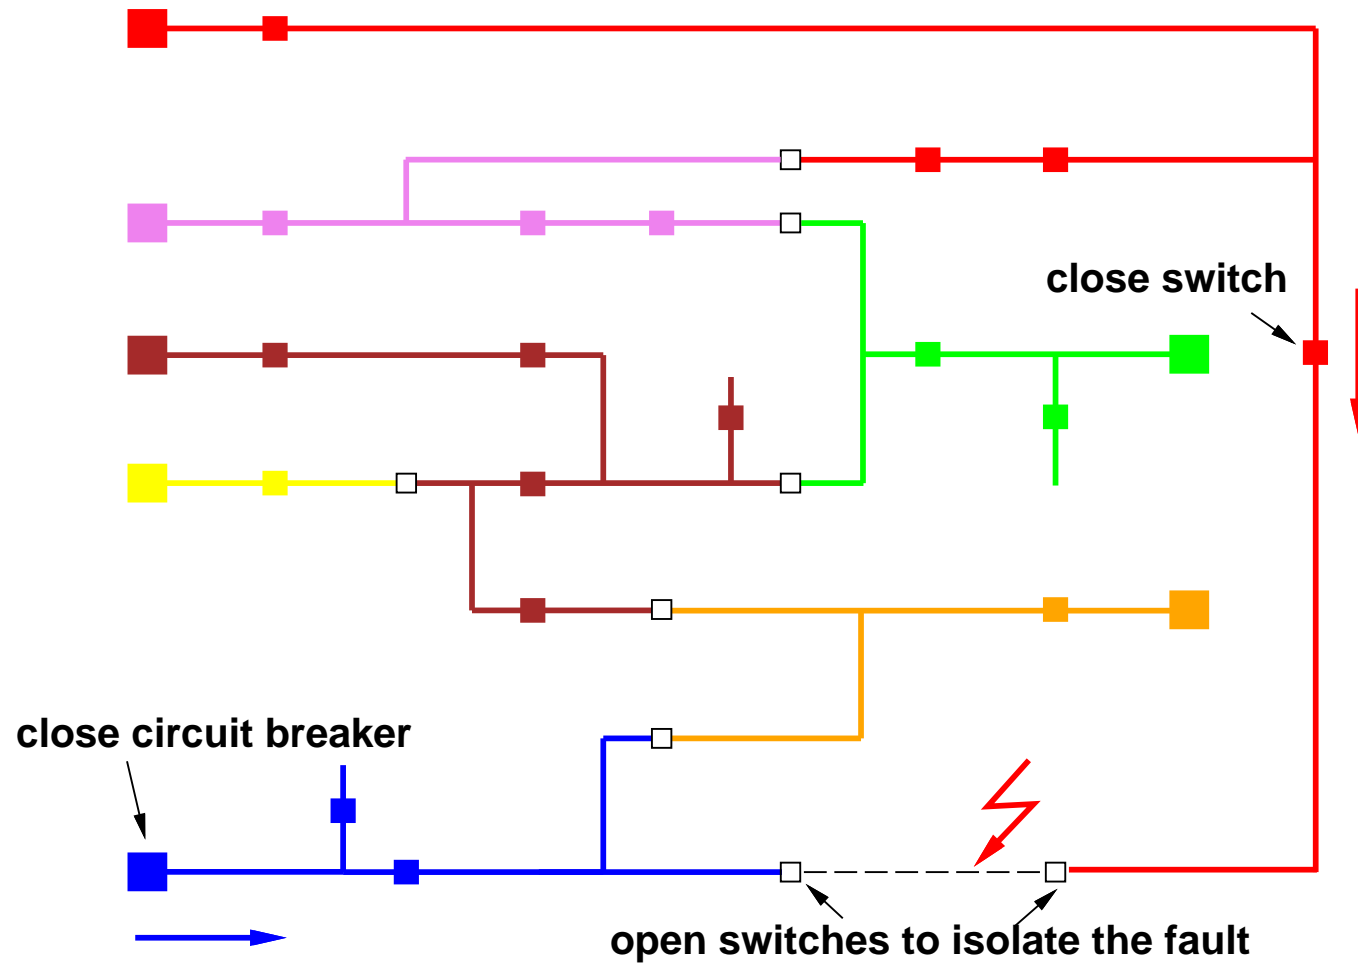
## Example: power supply restoration



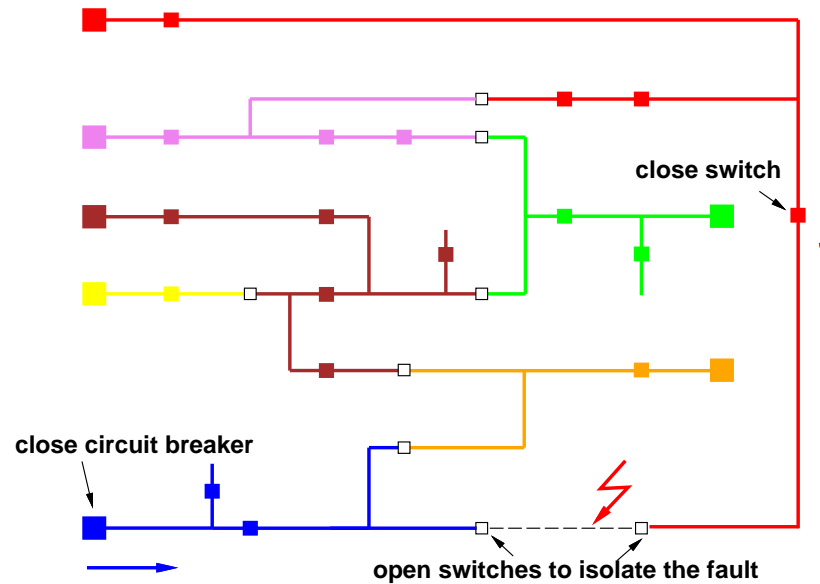
## Example: power supply restoration



## Example: power supply restoration



## Example: power supply restoration




states??: connections, status faulty/non-faulty of the lines,  
positions open/closed of the switches and circuit-breakers,

actions??: open/close a switch or a circuit-breaker

goal test??: resupply all non-faulty lines

path cost??: number of actions

## Example: power supply restoration



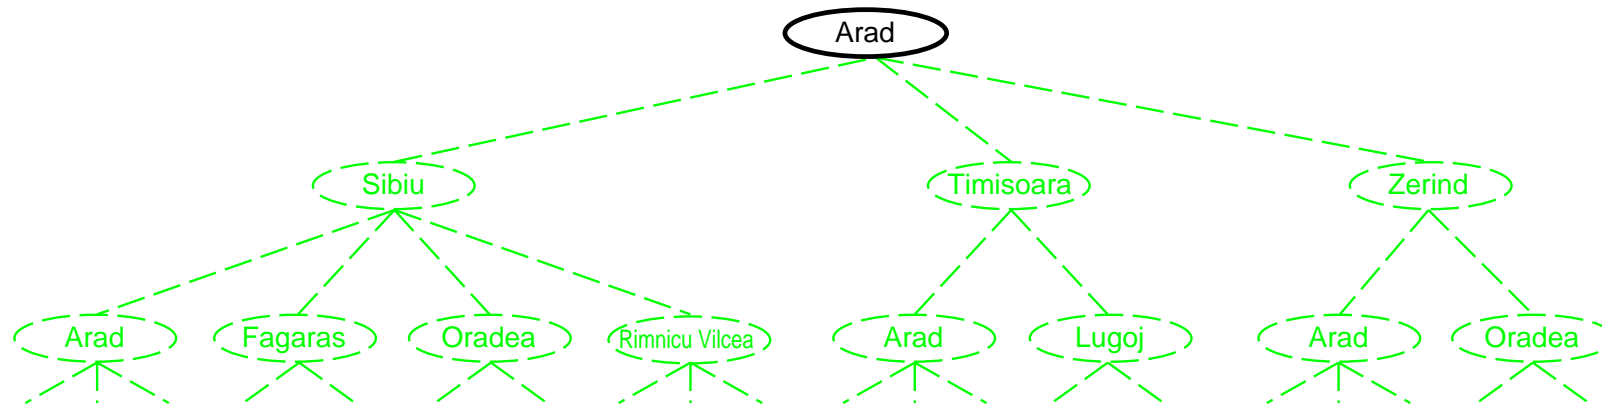
states??: connections, status faulty/non-faulty of the lines,  
positions open/closed of the switches and circuit-breakers,  
power consumed on each line, capacity of circuit-breakers and lines

actions??: open/close a switch or a circuit-breaker without exceeding capacity

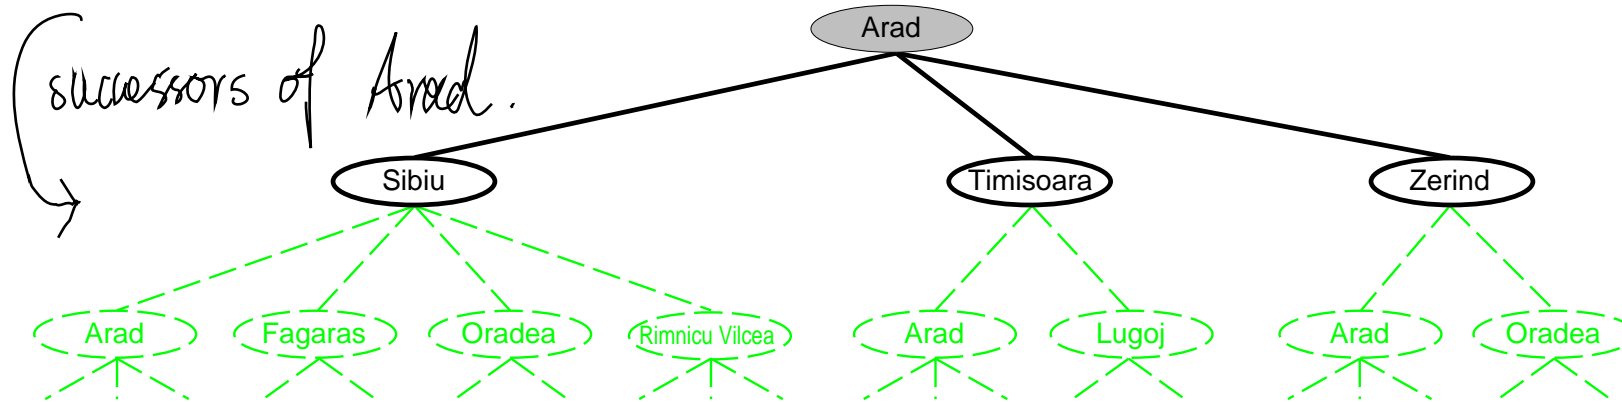
goal test??: resupply all non-faulty lines

path cost??: number of actions, power margins

# Tree search example

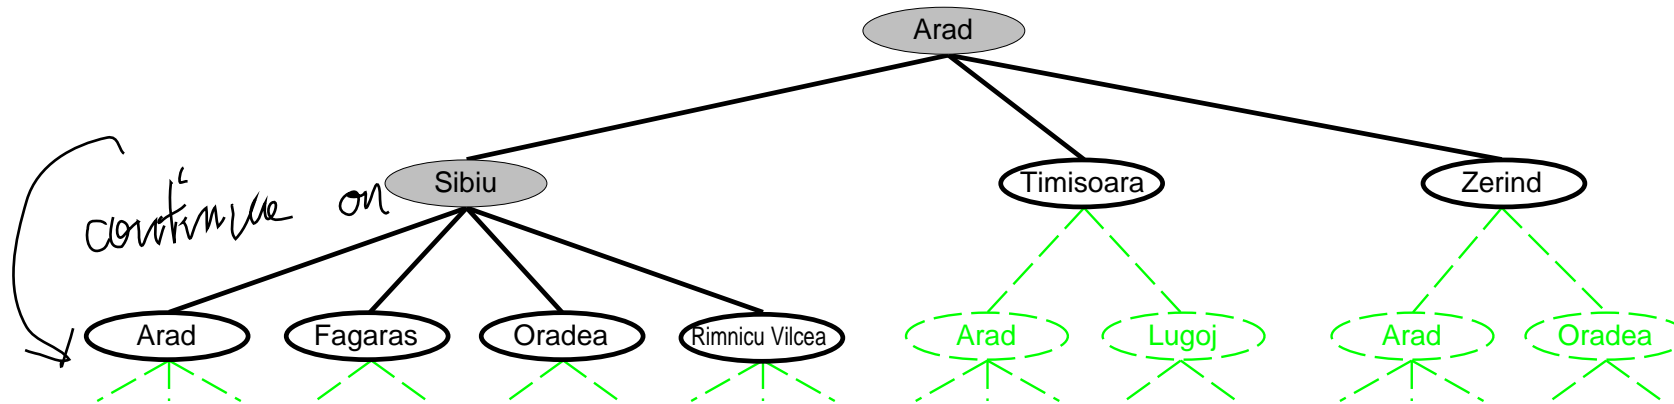


# Tree search example





# Tree search example



⋮  
until we have a goal.

# Tree search algorithm

Basic idea:

offline, simulated exploration of state space  
by generating successors of already-explored nodes  
(a.k.a. **expanding** nodes)

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion on the frontier then return failure
    choose a frontier node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the frontier of the tree
  end
```


## Implementation: states vs. nodes

A **state** is a (representation of) a physical configuration

A **node** is a data structure constituting part of a search tree

includes **state**, **parent**, **children**, **depth**, **path cost**  $g(n)$

States do not have parents, children, depth, or path cost!



State  
↓  
Expand Current Nodes  
↓  
Form New State w/  
New Nodes

The EXPAND function creates new nodes, filling in the various fields and using the SUCCESSORFN of the problem to create the corresponding states.

**Frontier** implemented as priority queue of nodes ordered according to strategy

## Implementation: general tree search

**function** TREE-SEARCH(*problem*, *frontier*) **returns** a solution, or failure

*frontier*  $\leftarrow$  INSERT(MAKE-NODE(INITIAL-STATE[*problem*]), *frontier*)

**loop do**

*emp*    **if** *frontier* is empty **then return** failure

*node*  $\leftarrow$  REMOVE-FRONT(*frontier*)     $\ll$  Clear the previous state

*res*    **if** GOAL-TEST(*problem*, STATE(*node*)) **then return** *node*

*next*    *frontier*  $\leftarrow$  INSERTALL(EXPAND(*node*, *problem*), *frontier*)     $\ll$  insert the new state

**function** EXPAND(*node*, *problem*) **returns** a set of nodes

*successors*  $\leftarrow$  the empty set

**for each** *action*, *result* **in** SUCCESSOR-FN(*problem*, STATE[*node*]) **do**

*s*  $\leftarrow$  a new NODE

PARENT-NODE[*s*]  $\leftarrow$  *node*; ACTION[*s*]  $\leftarrow$  *action*; STATE[*s*]  $\leftarrow$  *result*

PATH-COST[*s*]  $\leftarrow$  PATH-COST[*node*] + STEP-COST(STATE[*node*], *action*,  
*result*)

DEPTH[*s*]  $\leftarrow$  DEPTH[*node*] + 1     $\ll$  depth count for tree search

add *s* to *successors*

**return** *successors*

# Uninformed search strategies

A strategy is defined by picking the **order of node expansion**

This is the order used for the **priority queue** implementing the frontier

*distance / measurement* **Uninformed** strategies use only the information available in the definition of the problem

*from side* Breadth-first search

*to side of something* Uniform-cost search

Depth-first search

Depth-limited search

Iterative deepening search


# Breadth-first search

Expand shallowest unexpanded node

## Implementation:

*frontier* is a FIFO queue, i.e., new successors go at end

"first in, first out"




# Breadth-first search

Expand shallowest unexpanded node

## Implementation:

*frontier* is a FIFO queue, i.e., new successors go at end




# Breadth-first search

Expand shallowest unexpanded node

## Implementation:

*frontier* is a FIFO queue, i.e., new successors go at end






# Breadth-first search

Expand shallowest unexpanded node

## Implementation:

*frontier* is a FIFO queue, i.e., new successors go at end




# Breadth-first search

Expand shallowest unexpanded node

## Implementation:

*frontier* is a FIFO queue, i.e., new successors go at end




# Breadth-first search

Expand shallowest unexpanded node

## Implementation:

*frontier* is a FIFO queue, i.e., new successors go at end




# Breadth-first search

Expand shallowest unexpanded node

## Implementation:

*frontier* is a FIFO queue, i.e., new successors go at end



## Search strategies

A strategy is defined by picking the **order of node expansion**

Strategies are evaluated along the following dimensions:

**completeness**—does it always find a solution if one exists?

**solution optimality**—does it always find a least-cost solution?

**time complexity**—number of nodes generated (or expanded)

**space complexity**—maximum number of nodes in memory

Time and space complexity are measured in terms of

$b$ —maximum branching factor of the search tree

$d$ —depth of the shallowest solution

$m$ —maximum depth of the state space (may be  $\infty$ )


# Properties of breadth-first search

Complete??

# Properties of breadth-first search

Complete?? Yes (if  $b$  and  $d$  are finite)

Time??



depth

#nodes

0

$b^0$

1

$b^1$

2=d

$b^2 = b^d$

3=d+1


$b^3 - b = b^{d+1} - b$

# Properties of breadth-first search

Complete?? Yes (if  $b$  and  $d$  are finite)

Time??  $1 + b + b^2 + b^3 + \dots + b^d + (b^{d+1} - b) = O(b^{d+1})$ , i.e., exp. in  $d$

Space??





# Properties of breadth-first search


Complete?? Yes (if  $b$  and  $d$  are finite)

Time??  $1 + b + b^2 + b^3 + \dots + b^d + (b^{d+1} - b) = O(b^{d+1})$ , i.e., exp. in  $d$

Space??  $O(b^{d+1})$

Optimal??

$b = 10$ , 1 million node/sec, 1Kb/node,  $d=12$  would take 13 days and 1 petabyte of memory.



# Properties of breadth-first search


Complete?? Yes (if  $b$  and  $d$  are finite)

Time??  $1 + b + b^2 + b^3 + \dots + b^d + (b^{d+1} - b) = O(b^{d+1})$ , i.e., exp. in  $d$

Space??  $O(b^{d+1})$

Optimal?? Yes if cost = 1 per step; not optimal in general

$b = 10$ , 1 million node/sec, 1Kb/node,  $d=12$  would take 13 days and 1 petabyte of memory.



depth	#nodes
0	$b^0$
1	$b^1$
2=d	$b^2 = b^d$
3=d+1	$b^3 - b = b^{d+1} - b$

## Uniform-cost search

Expand least-cost unexpanded node

### Implementation:

*frontier* = queue ordered by path cost, lowest first

Equivalent to breadth-first if step costs all equal

Complete?? Yes, if step cost  $\geq \epsilon$

Time?? # of nodes with  $g \leq C^*$ ,  $O(b^{1+\lceil C^*/\epsilon \rceil})$   
where  $C^*$  is the cost of the optimal solution

Space?? # of nodes with  $g \leq C^*$ ,  $O(b^{1+\lceil C^*/\epsilon \rceil})$

Optimal?? Yes—nodes expanded in increasing order of  $g(n)$


# Depth-first search

Expand deepest unexpanded node

## Implementation:

*frontier* = LIFO queue, i.e., put successors at front

"last in first out"




# Depth-first search

Expand deepest unexpanded node

## Implementation:

*frontier* = LIFO queue, i.e., put successors at front




# Depth-first search

Expand deepest unexpanded node

**Implementation:**

*frontier* = LIFO queue, i.e., put successors at front




# Depth-first search

Expand deepest unexpanded node

## Implementation:

*frontier* = LIFO queue, i.e., put successors at front




# Depth-first search

Expand deepest unexpanded node

**Implementation:**

*frontier* = LIFO queue, i.e., put successors at front






# Depth-first search

Expand deepest unexpanded node

## Implementation:

*frontier* = LIFO queue, i.e., put successors at front




# Depth-first search

Expand deepest unexpanded node

## Implementation:

*frontier* = LIFO queue, i.e., put successors at front




# Depth-first search

Expand deepest unexpanded node

**Implementation:**

*frontier* = LIFO queue, i.e., put successors at front




# Depth-first search

Expand deepest unexpanded node

**Implementation:**

*frontier* = LIFO queue, i.e., put successors at front




# Depth-first search

Expand deepest unexpanded node

**Implementation:**

*frontier* = LIFO queue, i.e., put successors at front



# Properties of depth-first search

Complete??

# Properties of depth-first search

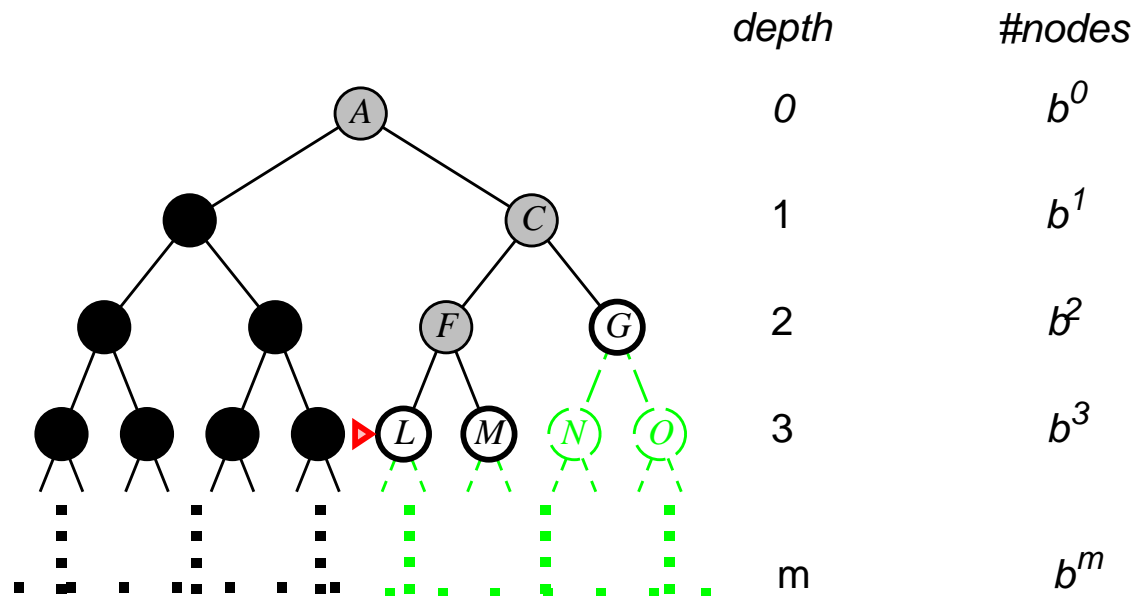


Complete?? No: fails in infinite-depth spaces, spaces with loops

Modify to avoid repeated states along path

⇒ complete in finite spaces

Time??



# Properties of depth-first search

Complete?? No: fails in infinite-depth spaces, spaces with loops


Modify to avoid repeated states along path

⇒ complete in finite spaces

Time??  $O(b^m)$ : terrible if  $m$  is much larger than  $d$

but if solutions are dense, may be much faster than breadth-first

Space??





# Properties of depth-first search

Complete?? No: fails in infinite-depth spaces, spaces with loops

Modify to avoid repeated states along path


⇒ complete in finite spaces

Time??  $O(b^m)$ : terrible if  $m$  is much larger than  $d$

but if solutions are dense, may be much faster than breadth-first

Space??  $O(bm)$ , i.e., linear space! (deepest node+ancestors+their siblings)

Optimal??



# Properties of depth-first search

Complete?? No: fails in infinite-depth spaces, spaces with loops

Modify to avoid repeated states along path


⇒ complete in finite spaces

Time??  $O(b^m)$ : terrible if  $m$  is much larger than  $d$

but if solutions are dense, may be much faster than breadth-first

Space??  $O(bm)$ , i.e., linear space! (deepest node+ancestors+their siblings)

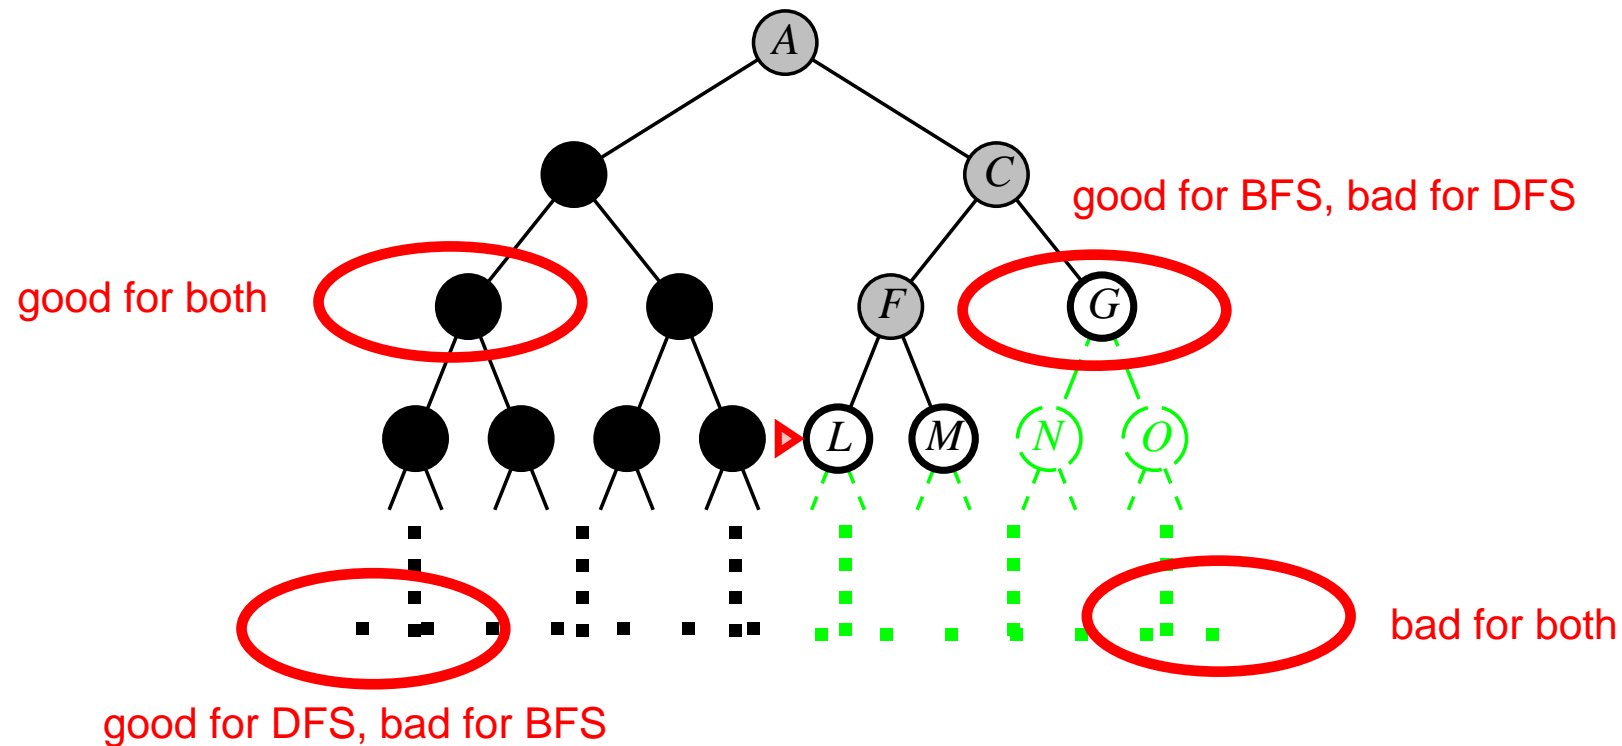
Optimal?? No



# Breadth-first versus depth-first search

Use **breadth-first** search when there exists **short** solutions.

Use **depth-first** search when there exists **many** solutions.



## Eternity II Puzzle

Use **breadth-first** search when there exists **short** solutions.

Use **depth-first** search when there exists **many** solutions.



2 million dollar prize! Few deep solutions

## Iterative deepening search

Combines advantages of breadth-first and depth-first search:

- completeness

- returns shallowest solution

- use linear amount of memory

Performs a **series of depth limited depth-first searches**

```
function ITERATIVE-DEEPENING-SEARCH(problem) returns solution/failure
  inputs: problem, a problem
  for depth  $\leftarrow$  0 to  $\infty$  do
    result  $\leftarrow$  DEPTH-LIMITED-SEARCH(problem, depth)
    if result  $\neq$  cutoff then return result
  end
```

## Depth-limited search

= depth-first search with depth limit  $l$ ,  
i.e., nodes at depth  $l$  have no successors

### Recursive implementation:

```

function DEPTH-LIMITED-SEARCH(problem, limit) returns soln/fail/cutoff
    RECURSIVE-DLS(MAKE-NODE(INITIAL-STATE[problem]), problem, limit)

function RECURSIVE-DLS(node, problem, limit) returns soln/fail/cutoff
    cutoff-occurred?  $\leftarrow$  false
    if GOAL-TEST(problem, STATE[node]) then return node
    else if DEPTH[node] = limit then return cutoff
    else for each successor in EXPAND(node, problem) do
        result  $\leftarrow$  RECURSIVE-DLS(successor, problem, limit)
        if result = cutoff then cutoff-occurred?  $\leftarrow$  true
        else if result  $\neq$  failure then return result
    if cutoff-occurred? then return cutoff else return failure
  
```

cutoff: no solution within the depth limit, failure: the problem has no solution





# Iterative deepening search $l = 0$

Limit = 0



# Iterative deepening search $l = 1$


Limit = 1





# Iterative deepening search $l = 2$

Limit = 2



# Properties of iterative deepening search

Complete??

## Properties of iterative deepening search

Complete?? Yes (if  $b$  and  $d$  are finite)

Time??

## Properties of iterative deepening search

Complete?? Yes (if  $b$  and  $d$  are finite)

Time??  $(d + 1)b^0 + db^1 + (d - 1)b^2 + \dots + b^d = O(b^d)$

Space??

## Properties of iterative deepening search

Complete?? Yes (if  $b$  and  $d$  are finite)

Time??  $(d + 1)b^0 + db^1 + (d - 1)b^2 + \dots + b^d = O(b^d)$

Space??  $O(bd)$

Optimal??

## Properties of iterative deepening search

Complete?? Yes (if  $b$  and  $d$  are finite)

Time??  $(d + 1)b^0 + db^1 + (d - 1)b^2 + \dots + b^d = O(b^d)$

Space??  $O(bd)$

Optimal?? Yes, if step cost = 1

Can be modified to explore uniform-cost tree

Numerical comparison for  $b = 10$  and  $d = 5$ , solution at far right leaf:

Time: IDS does better because other nodes at depth  $d$  are not expanded

$$N(\text{IDS}) = 6 + 50 + 400 + 3,000 + 20,000 + 100,000 = 123,456$$

$$N(\text{BFS}) = 1 + 10 + 100 + 1,000 + 10,000 + 100,000 + 999,990 = 1,111,101$$

BFS can be modified to apply goal test when a node is **generated**

Space: IDS does much better  $N(\text{IDS}) = 50$ ,  $N(\text{BFS}) \simeq 1,000,000$

## Summary of algorithms

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening
Complete?	Yes*	Yes*	No	Yes, if $l \geq d$	Yes*
Time	$b^{d+1}$	$b^{1+\lceil C^*/\epsilon \rceil}$	$b^m$	$b^l$	$b^d$
Space	$b^{d+1}$	$b^{1+\lceil C^*/\epsilon \rceil}$	$bm$	$bl$	$bd$
Optimal?	Yes*	Yes	No	No	Yes*

## Summary

### ◇ Problem solving agents

formulate a problem, search off-line for a solution, execute it

### ◇ Problem formulation

initial state, successor function, goal test, path cost

### ◇ Example problems

traveling around romania, 8-puzzle, power supply restoration

### ◇ Tree search algorithms

build and explore a tree, strategy picks up the order of node expansion

### ◇ Implementation

select the first node on the frontier, test for goal, expand

### ◇ (Uninformed) strategies (breadth first, uniform cost, ...)

characterised by their completeness, optimality, complexity

### ◇ Iterative deepening complete, finds the shallowest solution

uses only linear space and no more time than uninformed strategies