

BEYOND CLASSICAL SEARCH

CHAPTER 4, SECTION 5

Outline

- ◇ On-line search in unknown environments
- ◇ On-line search problem formulation
- ◇ Learning Real-Time A*

On-line search and unknown environments

Off-line search commits to a solution and executes it

On-line search algorithms interleave computation and action:

search \rightarrow execute \rightarrow observe \rightarrow search \rightarrow ...

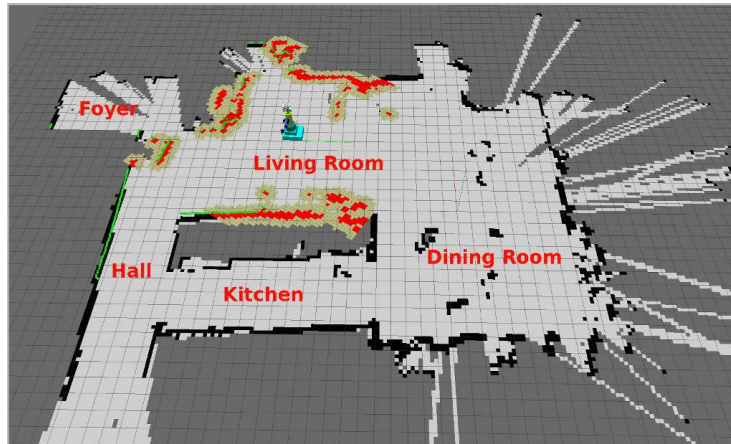
When is it useful???

- ◇ good idea in **dynamic** environments
- ◇ good idea in **time-constrained** environments
- ◇ necessary in **unknown** environments [exploring & learning]

Think of baby gradually discovering how the world works

Applications

- ◇ building a map of an unknown building
- ◇ search and rescue applications
- ◇ submarine robot exploration of Europa



On-line search problem formulation

An on-line search problem consists of:

- ◇ $\text{ACTIONS}(s)$: returns the set of legal actions in state s [effects not known]
- ◇ The step cost function $c(s, a, s')$: returns the cost of going from s to s' via action a . This cannot be used until the agent has tried a in s and knows that s' is the outcome. [learnt]
- ◇ $\text{GOAL-TEST}(s)$: returns true if s is a goal state

Observability: the current percept identifies the current state [Assume percept is good enough]

The agent cannot determine $\text{RESULT}(s, a)$ except by being in s , executing a and observing the result! Degree of ignorance might be reduced for some applications then memorize


The agent might have access to some (admissible) heuristic function $h(s)$ [assume it has]

[Assume] No dead-end: the goal must be reachable from every state
no algorithm can avoid dead-ends in all environments


↑
estimate

Local search algorithms



Hill-climbing using $h(s)$ can quickly get stuck in a local minimum.



Random walk will eventually find a goal but will often take exponentially many steps.



Random vs LRTA*



LRTA*: Learning Real-Time A*

Learning action results

Store a table $result(s, a)$ recording the result of actions tried

When a is first tried in s , record the resulting state

Learning estimated cost to the goal \Rightarrow heuristic

Store and update a table $H(s)$ recording the estimate of visited states

When s is first met, initialise $H(s) = h(s)$ for some admissible $h(s)$

When s is left, update its cost estimate $H(s)$ to $\min_a C(s, a)$

A* [just like A* which the initial state ^{of path} is not the start state] estimate when met. replace this est. when we leave it

$C(s, a)$ is the estimated cost of reaching the goal from state s via action a :

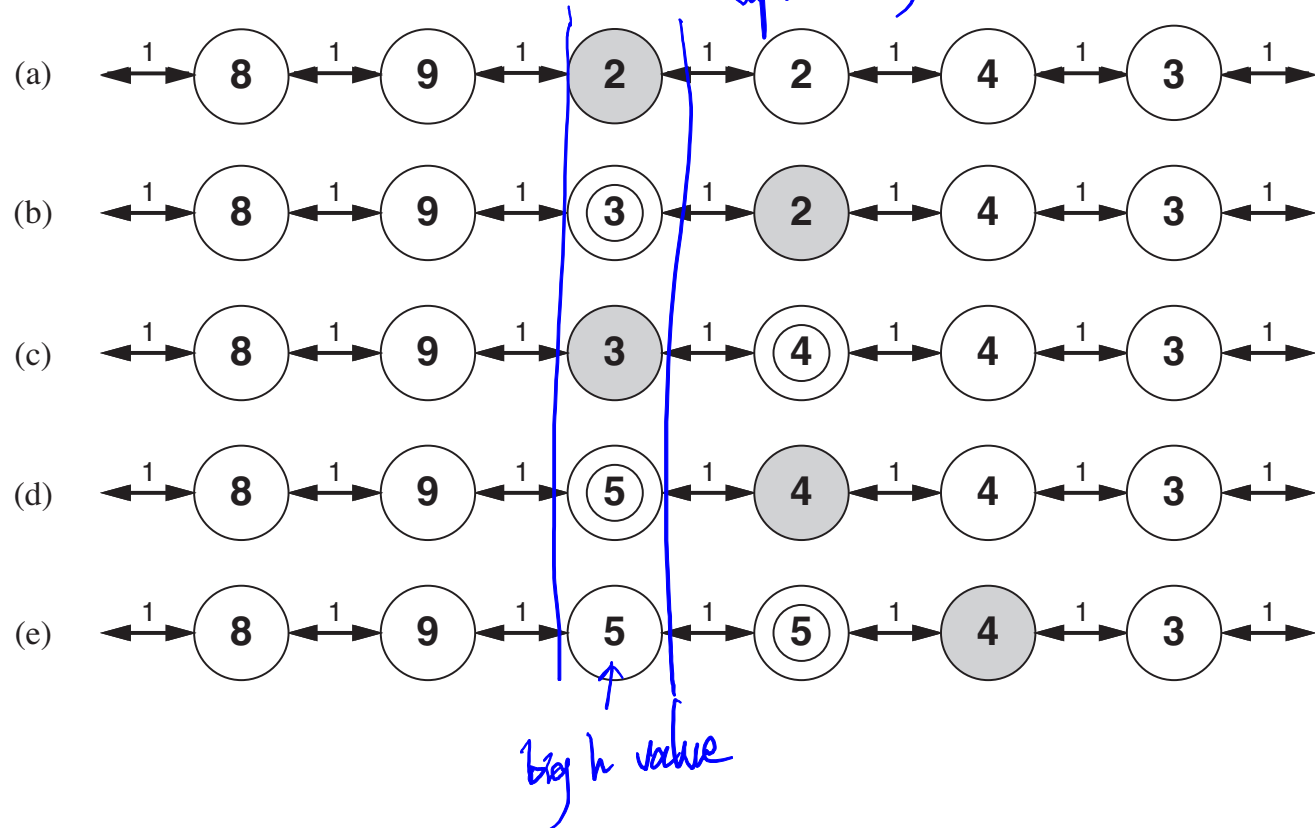
$$C(s, a) = \begin{cases} c(s, a, s') + H(s') & \text{if } result(s, a) \text{ is known to be } s' \\ h(s) & \text{otherwise} \end{cases}$$

Take the best local move (real-time)

In state s select the action a minimising the estimated cost $C(s, a)$

Update of cost estimates

Cost estimates **improve and converge to the true cost** (over a number of trials) This allows escaping local maxima.



LRTA* Algorithm

function LRTA*-AGENT(*cur*) **returns** an action

inputs: *cur*, a percept that identifies the current state

static: *prev, action*, the previous state and action, initially null

result, table indexed by state and action, initially empty

H, table of cost estimates indexed by state, initially empty

if GOAL-TEST(*cur*) **then return** *stop*

if *cur* is a new state **then** $H[cur] \leftarrow h(cur)$ / heuristic

if *prev* is not null **then**

$result[prev, action] \leftarrow cur$

$H[prev] \leftarrow \min_{a \in ACTIONS(prev)} EST-COST(\underline{prev}, a, result[prev, a], H)$ ← update

$action \leftarrow \operatorname{argmin}_{a \in ACTIONS(cur)} EST-COST(\underline{cur}, a, result[cur, a], H)$

prev ← *cur*

return *action*



function EST-COST(*s*, *a*, *s'*, *H*) **returns** a cost estimate

if *s'* is undefined **then return** $h(s)$

else return $c(s, a, s') + H[s']$

Δ

Example



result function


		0
5	4	1
4	3	2

cost estimation H

2	1	0
3	2	1
4	3	2

heuristic h

Example



"oh I have actions A & B,
take one at random"

result function


4		

cost estimation H

2	1	0
3	2	1
4	3	2

heuristic h

Example



↑		

result function



3		
4		

cost estimation H


2	1	0
3	2	1
4	3	2

heuristic h


Example



result function





cost estimation H




heuristic h


Example



result function




cost estimation H




heuristic h

Example



"now I have estimations"



result function



3	2	
4	3	

cost estimation H


2	1	0
3	2	1
4	3	2

heuristic h


Example



result function





cost estimation H



heuristic h

Example



result function



5	4	
4	3	

cost estimation H

2	1	0
3	2	1
4	3	2

heuristic h

Example



result function



5	4	
4	3	

cost estimation H

2	1	0
3	2	1
4	3	2

heuristic h

Example



result function



5	4	
4	3	

cost estimation H

2	1	0
3	2	1
4	3	2

heuristic h

Example



result function



5	4	
4	3	

cost estimation H

2	1	0
3	2	1
4	3	2

heuristic h

Example



result function



5	4	
4	3	2

cost estimation H

2	1	0
3	2	1
4	3	2

heuristic h

Example



result function



5	4	1
4	3	2

cost estimation H

2	1	0
3	2	1
4	3	2

heuristic h

Example



result function

		0
5	4	1
4	3	2

cost estimation H

2	1	0
3	2	1
4	3	2

heuristic h

Properties of LRTA*

Complete?? Yes if the state space is finite (and there are no dead-ends)

Time?? $O(|S|^3)$

Space?? $O(|S| \times |A|)$ to store Result

Optimal?? no, but if h is admissible, converges to the optimal over repeated trials

“actually pretty good”

Summary

Exploration problems arise when the agent does not know the states or physics of its environment.

On-line search can build a map and find a goal state if there are no dead-ends.

Hill-climbing can get stuck in local minima.

LRTA* updates heuristic information to escape local minima and find the goal in much fewer steps than random walks.