


KNOWLEDGE REPRESENTATION AND REASONING: SOLVING CONSTRAINT SATISFACTION PROBLEMS

CHAPTER 7

Constraint Satisfaction Problems



- ◇ Binary constraint network $\gamma = \langle V, D, C \rangle$
- V a finite set of variables v_1, \dots, v_n
 - D a set of [finite] sets D_{v_1}, \dots, D_{v_n}
 - C a set of binary relations $\{C_{u,v} \mid u, v \in V, u \neq v\}$
 $C_{u,v} \subseteq D_u \times D_v$

Outline of the lecture

- ◇ Recall constraint networks and backtracking search
- ◇ Tightening CSPs by learning from mistakes
- ◇ Problem structure: constraint graphs
- ◇ Symmetry
- ◇ CSPs and optimisation
- ◇ Summary

Recall


- improvement*
- ◇ Pure backtracking [*Partial Tree Construction → Go through partial assignments*]
 - If the current partial assignment is consistent
 - Choose a variable, assign each value from its domain in turn
 - Search the resulting sub-tree

- ◇ Forward checking [*special case of AC — can consider in this way*]
 - Prune values from neighbour variables if they are not supported by the assigned one


- ◇ Arc consistency
 - Prune similarly for all pairs of values related by a constraint // *stronger Version*
- ◇ Variable ordering and value ordering heuristics important for real efficiency

↑
recall KR 3


Learning from mistakes




Learning from mistakes



Learning from mistakes




Learning from mistakes



[only apply FC here]


- ◇ This assignment is consistent but can't be extended to a solution

Learning from mistakes



◇ This assignment is consistent but can't be extended to a solution

Learning from mistakes




◇ The previous assignment must be wrong

not counting the last green one, which was forced

so remember the earlier choices, and don't do it again!

[BRBR → wipe out]

Learning from mistakes



◇ Actually, we're going to backtrack further

so the bottom line was no good.

Remember that combination $(v_1:b, v_2:r, v_3:b)$ as a disallowed triple of a (3-ary) constraint

[BRBR thrown out, remember BRB]

◇ Never repeat a mistake: don't backtrack twice for the same reason


only for these variables

Constraint learning: notes

- ◇ Learned constraints may be added to the network or kept separately *[if these are binary, we can add them]*
- ◇ A separate store of nogoods is usual, as they are usually large
 - May add binary ones to the network and store the rest
 - Data structures matter: indexing for rapid inference is important
- ◇ Every branch may add another nogood, so there are too many of them
 - Storage requires exponential space
- ◇ Hence common to have a strategy for forgetting them *solution*
 - e.g. let the longest ones lapse after a while
 - or just keep the “tail” and discard when backtracking leaves the region where it applies *depends on situations*
- ◇ Constraint learning useful for CSP solvers; essential for SAT solvers

remembering vs. memory spacing

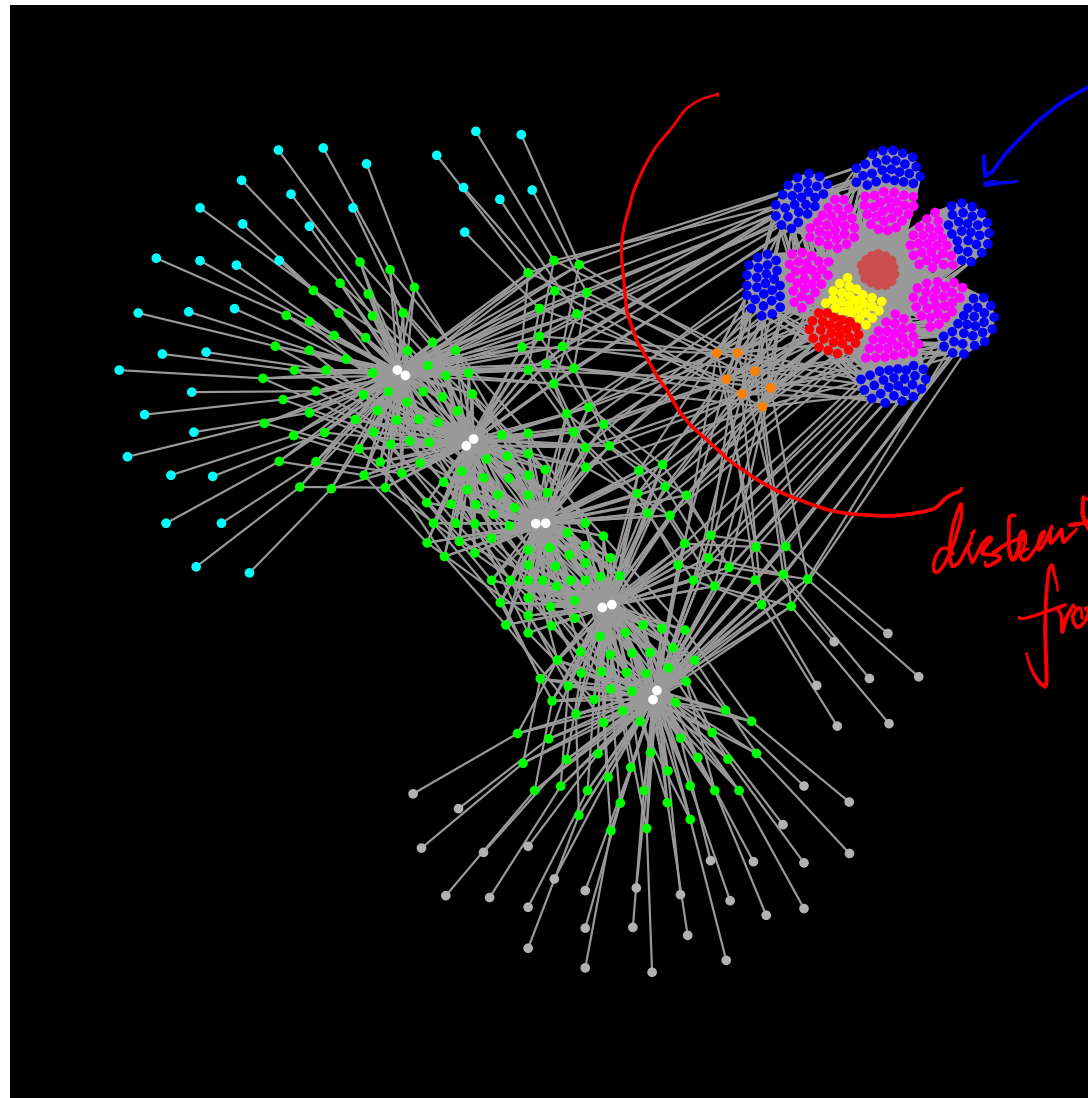
Constraint graphs



- ◇ Some decision variables are related by constraints, some are not
- ◇ Hence we may consider the graph where
 - vertices are decision variables
 - edges are constraints
- ◇ Graph contains information about the structure of the problem
- ◇ Great for visualisation, as well as automated reasoning

"shape of the problem"

PSR Constraint Graph



calculations

Power Supply Grid
Problem


distant
from plans.

green : switches

white : plans

orange : circuit breakers

PSR Constraint Graph




Can observe properties of the encoding from the graph:


- The plan (white) only loosely communicates with the calculations (blue and magenta)
- There is no direct information flow from the calculations at one step to the calculations at the next, even though most of the distribution grid is the same
- The first line of switches (green), next to the circuit breakers (orange) has a special status in the CSP. This may be worth investigating.

*you can see
some features
of your problem.*

Constraint graphs: notes

- ◇ The examples are static views. Dynamic ones animated to show the search can also be very useful.
- ◇ The dual graph, where the vertices are the constraints and an edge between two constraints means they share a variable, gives yet another view.
- ◇ So does the bipartite graph with variable nodes and constraint nodes.
- ◇ Constraint graphs are not specific to binary CSPs: they can be useful in analysing logical descriptions of given domains, in SAT solving or in automated reasoning generally.
-  Note: the constraint graph only shows which decision variables are connected. It is not affected by whether the problem has solutions or not.

Tree-like constraint graphs



If the constraint graph is a tree, this is **always good news!**

We can always solve such a CSP efficiently:

Enforce arc consistency: if wiped out, you're done

Choose a vertex to be the root of the tree [any vertex]

Start assigning values at the root


Don't assign a value to a variable before its parent in the tree

Do forward checking at each step

all "bad values" pruned in time


The search will be backtrack-free.

Constraint graph: Longmult (SAT)




Verifying Algorithms

Constraint graph: logical calculus tester



Normal view: variables as vertices


Constraint graph: logical calculus tester



Dual view: constraints as vertices


/ slightly different
structure .

Constraint graph: logical calculus tester




Bipartite view: variables and constraints as vertices


Symmetry



Symmetry




Symmetry



◇ What would happen if we started with a different choice of colours?


Symmetry



- ◇ What would happen if we started with a different choice of colours?
- ◇ Exactly the same, but with the colours interchanged.
- ◇ So any solution to graph colouring with k colours can be re-labelled to give $k!$ solutions with the same colours in different orders.
- ◇ We say that the values in this problem are symmetric.


[pattern stays. labels changed]

Symmetry



◇ What would happen if we started at the top right?

Symmetry



- ◇ What would happen if we started at the top right?
- ◇ Exactly the same, but inverted.
- ◇ Any solution can be rotated or reflected in a diagonal to give an equivalent solution with variables interchanged.
- ◇ We say that this problem has a variable symmetry.

Using symmetry

- ◇ Is symmetry on our side or against us?
- ◇ It's our friend if we know about it and use it, but our enemy otherwise! *[waste of time] [not only solutions, but also partial assignments]*
- ◇ The bad part: if a problem has lots of symmetries, we can waste huge amounts of time searching symmetric (and equally empty) sub-spaces, or generating solutions that tell us nothing really new.
- ◇ The good part: if we explore one of these sub-spaces, we know we can prune all of the others without losing anything essential.
- ◇ Unlike arc consistency, etc, symmetry pruning can delete solutions, but it can never delete **all** of them.

Symmetry: how it's done

◇ Note that if solutions are symmetric, partial assignments have (at least) the same symmetries, so early pruning may be possible

~~✗~~ The usual method for removing symmetric sub-problems is to add symmetry-breaking constraints

◇ Formulae true of one (or some) of the symmetric solutions but false of the others.


⚡ will not destroy solvability

◇ E.g. we could add a constraint saying $v_1 = \text{blue} \wedge v_2 = \text{red}$.
— safe addition: if there are solutions, there's one satisfying this
— reduces two of the domains to singletons
— rules out 5 of the 6 symmetric solutions.

“wide range of choices”


◇ If we want to recover the missing solutions, that's possible without search.

Symmetry: an extreme case



- ◇ Suppose we have a **pigeonhole problem**: show that it's impossible to fit 10 pigeons in 9 pigeonholes (without overcrowding)
- ◇ A backtracking search will start assigning holes to pigeons, house 9 of them and discover that the tenth has nowhere to go.
- ◇ Then it will backtrack, try a different ninth pigeon, and find that there is still one left over ... etc.
- ◇ **9! backtracks, even with arc consistency; no solution.**

Symmetry: an extreme case



- ◇ But one pigeon looks just like another (to a CSP solver), and one hole looks just like another as well. *// obvious symmetry*
- ◇ So pigeon number 1 goes in hole number 1, without loss of generality.
- ◇ Assign hole 2 to pigeon 2, etc. Then pigeon 10 is homeless. The end!
- ◇ A good symmetry-breaker is $\forall x \forall y ((x < y) \rightarrow (\text{hole}(x) < \text{hole}(y)))$.
—would be true of 1 solution if there were just enough holes
- ◇ 9! branches reduced to 1.

Optimal Solutions

- ◇ Constraint solvers often asked to produce **optimal** solutions
 - though in practice, suboptimal but **good** solutions suffice
- ◇ Optimisation not treated (much) in this course
 - worth a course on its own
- ◇ However, we should note it, so:

Optimal Solutions

- ◇ Constraint solvers often asked to produce **optimal** solutions
 - in practice, suboptimal but **good** solutions usually suffice
- ◇ Optimisation not treated (much) in this course
 - worth a course on its own
- ◇ However, we should note it, so
- ◇ Two common ways of defining “better” or “worse” solutions:
 1. via an **objective function**: a quantity to be minimised (or maximised)
 2. via **soft constraints**: can be violated, but as little as possible
- ◇ The sum of soft constraint violations behaves as an objective function.

Optimal Solving

There are **many** techniques for solving problems optimally. The only one to be noted here is Depth First Branch and Bound (DFBB)

- ◇ The default search algorithm used by most FD solvers
- ◇ Easy to implement, generally applicable, complete
- ◇ Also functions well as an anytime method

Branch and Bound

- ◇ Use lower bound estimate L of the cost of solutions extending the current partial assignment
 - underestimates the objective function at each node
- ◇ Also use a bound B
 - overestimates the objective function (globally)
 - initialise to infinity (or a known overestimate)
- ◇ Traverse the search tree e.g. depth first
- ◇ Backtrack if $L \geq B$
- ◇ Each time a solution is found, set B to its objective value
- ◇ B is monotone decreasing as solutions are found
- ◇ So search tree branches tend to get shorter towards the end

hence
strictly
improving

DFBB: Intermediate solutions

- ◇ First solution is at the bottom of the leftmost (complete) branch
 - Fast: Likely to be found quickly
 - Dirty: Likely to be of low quality
- ◇ Always trying to improve on the best so far
 - Any improvement will do
- ◇ So DFBB produces a sequence of (strictly) improving solutions
- ◇ We can interrupt the search at any time
 - when the current solution is good enough
 - when a time limit expires
 - when the next process needs to start
 - when we just get fed up with waiting
- ◇ Intermediate solutions are valuable, because optimal ones can be very expensive to compute (and proofs of optimality even more expensive).

Summary

- ◇ Constraint (nogood) learning from wipeouts usually improves efficiency
 - Space (memory) is a limitation for nogood learning, so forgetting is also important
- ◇ Constraint graphs give information about problem structure
 - Certain constraint graphs (e.g. trees) indicate that problems are easy
- ◇ Value symmetry and variable symmetry are frequently present in CSPs
 - Pruning symmetric sub-spaces is a big winner where there is extensive symmetry
- ◇ Optimisation (minimising a cost or objective function) is usual for CSPs
- ◇ Depth First Branch and Bound is commonly used in FD solvers
 - Conveniently provides intermediate solutions of increasing goodness