


KNOWLEDGE REPRESENTATION AND REASONING: SOLVING CONSTRAINT SATISFACTION PROBLEMS

CHAPTER 6.2, 6.3

Constraint Satisfaction Problems



- ◇ Binary constraint network $\gamma = \langle V, D, C \rangle$
 - V a finite set of variables v_1, \dots, v_n
 - D a set of [finite] sets D_{v_1}, \dots, D_{v_n}
 - C a set of binary relations $\{C_{u,v} \mid u, v \in V, u \neq v\}$
 $C_{u,v} \subseteq D_u \times D_v$

Outline of the lecture

- ◇ Constraint modelling
- ◇ Inference
- ◇ Forward checking
- ◇ Variable and value ordering
- ◇ Arc consistency

Constraint Modelling

- ◇ Before any constraint solving can happen, the CSP must be **defined**
- ◇ Model must define V , D and C
- ◇ Explicit definition of C is painful, so use high-level description
- ◇ Hence we want to write a **logical** models
 - Write constraints as formulae of first order logic
 - Describe what would count as a solution to the problem
 - Compiler will turn this into a low-level constraint network
 - **Logical model is purely declarative: no algorithm!**
- ◇ Old style constraint programming: logic is implicit in the program

MiniZinc

We shall use the constraint modelling language MiniZinc

```
-----  
% N Queens Problem: place N queens on an NxN  
% chessboard so that no queen attacks another
```

```
int: N;  
array[1..N] of var 1..N: q;  
constraint forall (x,y in 1..N where x < y) (  
    q[x] != q[y] /\  
    abs(q[x]-q[y]) != y-x);  
solve satisfy;  
-----
```

MiniZinc is essentially first order logic with some syntactic sugar and basic support for arithmetic etc.

Minizinc

- ◇ MiniZinc model is completely solver-independent
- ◇ Also good to separate the “conceptual model” of the problem from data defining a specific problem instance
 - e.g. for the N queens, the data file could specify $N = 8$;
- ◇ MiniZinc gets transformed into a simple fragment “Flat Zinc”
- ◇ Flat Zinc can be turned into input code for many solvers
 - Finite domain (FD) solvers
 - Mixed integer programming (MIP) solvers
 - SAT solvers
 - Local search solvers
- ◇ Mapping into low-level data structures is internal to the solvers
- ◇ Logical model + default mapping: the Holy Grail

Recall Backtracking

function BACKTRACK(γ, a) **returns** solution, or “inconsistent”

if a is inconsistent with γ **then return** “inconsistent”

if a is total **then return** a

select variable v for which a is not defined

for each d in D_v **do**

$a' \leftarrow a \cup \{(v, d)\}$

$a'' \leftarrow \text{BACKTRACK}(\gamma, a')$

if $a'' \neq \text{“inconsistent”}$ **then return** a''

end


return “inconsistent”

call: BACKTRACK($\gamma, \{ \}$)

Backtracking: the Good and the Bad


- ◇ Better than exhaustive search: avoids enumerating many inconsistent (partial) assignments by detecting them as soon as they happen
- ◇ Once an inconsistent partial assignment is reached, all of its extensions are pruned
- ◇ **Advantages:**
 - Very simple to implement
 - Very fast (per node of the search tree)
 - Complete (always gives a decision)
- ◇ **Disadvantages:**
 - Does no reasoning except detecting actual inconsistency
 - Cannot look further ahead than the current state

Simple example: Graph colouring




- ◇ Given an undirected graph with n nodes, given k colours, assign a colour to each node so that no two adjacent nodes (with an edge between them) are the same colour.
- ◇ Representation using binary constraints is easy.
- ◇ Problem is NP-complete, so difficult in the worst case.

Inference




- ◇ Assign values from the bottom left corner, going across the rows

Inference




- ◇ Assign values from the bottom left corner, going across the rows
- ◇ Choose blue first

Inference




- ◇ Assign values from the bottom left corner, going across the rows
- ◇ Choose blue first
- ◇ Inconsistent

Inference




- ◇ Assign values from the bottom left corner, going across the rows
- ◇ Choose blue first
- ◇ Choose red next

Inference




- ◇ Assign values from the bottom left corner, going across the rows
- ◇ Choose blue first, then red
- ◇ Now nodes 5 and 6 must both be green

Inference




- ◇ Assign values from the bottom left corner, going across the rows
- ◇ Choose blue first, then red
- ◇ Nodes 5 and 6 must both be green

Inference




- ◇ Assign values from the bottom left corner, going across the rows
- ◇ Choose blue first, then red
- ◇ Nodes 5 and 6 must both be green

Inference




- ◇ Assign values from the bottom left corner, going across the rows
- ◇ Choose blue first, then red
- ◇ You are wasting your time!

Inference




- ◇ Assign values from the bottom left corner, going across the rows
- ◇ Choose blue first, then red
- ◇ You are wasting your time!

Inference




- ◇ Assign values from the bottom left corner, going across the rows
- ◇ Choose blue first, then red, then green
- ◇ It won't work!!

Inference




- ◇ Assign values from the bottom left corner, going across the rows
- ◇ Choose blue first, then red, then green

Inference




- ◇ Assign values from the bottom left corner, going across the rows
- ◇ Choose blue first, then red, then green

Inference



- ◇ Assign values from the bottom left corner, going across the rows
- ◇ Choose blue first, then red, then green

Inference



- ◇ Assign values from the bottom left corner, going across the rows
- ◇ Choose blue first. then red, then green
- ◇ So a way of detecting the problem early could save work.
So could assigning the green ones before the red one to their left

More about inference

- ◇ Inference in CSP solving: deducing additional constraints that follow from the already known constraints.
- ◇ Hence a matter of replacing γ by an equivalent and strictly tighter constraint network γ' .
- ◇ γ and γ' with the same variables are **equivalent** iff they have the same solutions.
- ◇ $\gamma' = (V, D', C')$ is **tighter** than $\gamma = (V, D, C)$ iff:
 - (i) For all $v \in V$, $D'_v \subseteq D_v$
 - (ii) For all $C_{u,v} \in C$, $C'_{u,v} \subseteq C_{u,v}$ γ' is **strictly tighter** than γ if it is tighter and $\gamma \neq \gamma'$
- ◇ Inference reduces the number of consistent partial assignments

How to use inference

Inference as offline pre-processing

- ◇ Just once before search starts
- ◇ Little runtime overhead, modest pruning power. Not considered here.
— but important in SAT solving, for instance

Inference during search

- ◇ At every recursive call of backtracking
- ◇ When backing up out of a search branch, retract any inferred constraints that were local to that branch because they depend on a
- ◇ Strong pruning power. May have large runtime overhead

Backtracking with inference

function BACKTRACK(γ, a) **returns** solution, or “inconsistent”

if a is inconsistent with γ **then return** “inconsistent”

if a is total **then return** a

$\gamma' \leftarrow$ a copy of γ

$\gamma' \leftarrow$ Inference(γ', a)

if exists v with $D'_v = \{ \}$ **then return** “inconsistent”

select variable v for which a is not defined

for each d in D'_v **do**

$a' \leftarrow a \cup \{(v, d)\}$

$a'' \leftarrow$ BACKTRACK(γ', a')

if $a'' \neq$ “inconsistent” **then return** a''

end


return “inconsistent”

◇ Inference sets $D_v = \{d\}$ for each $(v, d) \in a$ and then delivers a tighter equivalent network.

Forward Checking


- ◇ Inference: for all variables v and u where $a(v) = d$ is defined and $a(u)$ is undefined, set D_u to $\{d' : d' \in D_u, (d', d) \in C_{u,v}\}$.
- ◇ That is, remove from domains any value not consistent with those that have been assigned.
- ◇ Obviously sound: it does not rule out any solutions
- ◇ Can be implemented incrementally for efficiency: only necessary to consider v to be the variable which has just been assigned.
- ◇ Simple to implement and low computational cost
- ◇ Almost always pays off (unless subsumed by stronger inferences)

Forward Checking example




◇ As before, start in the bottom left corner and go across the rows

Forward Checking example




◇ Impossible values get removed from related domains

Forward Checking example




◇ So no inconsistent assignment actually gets reached

Forward Checking example




◇ We still don't make two-step inferences

Forward Checking example




◇ We still don't make two-step inferences

Forward Checking example




◇ Now there is a wipeout: a variable with an empty domain

Forward Checking example




◇ Backtrack and change – but now there is another wipeout

Forward Checking example




◇ So backtrack some more

Forward Checking example




◇ Continue to explore the branch

Forward Checking example




◇ Now some moves are forced

Forward Checking example




◇ Now some moves are forced, and still consistent

Forward Checking example




◇ Blue is the bad choice

Forward Checking example




◇ And ...

Forward Checking example



◇ And we're ...

Forward Checking example



◇ And we're done!

Making choices

function BACKTRACK(γ, a) **returns** solution, or “inconsistent”

if a is inconsistent with γ **then return** “inconsistent”

if a is total **then return** a

select some variable v for which a is not defined

for each d in D_v **in some order do**

$a' \leftarrow a \cup \{(v, d)\}$

$a'' \leftarrow \text{BACKTRACK}(\gamma, a')$

if $a'' \neq \text{“inconsistent”}$ **then return** a''

end

return “inconsistent”

call: BACKTRACK($\gamma, \{ \}$)

The size of the search space depends on the order in which we choose variables and values.

Variable ordering

- ◇ Common strategy: most constrained variable (aka “first-fail”)
Choose a variable with the smallest (consistent) domain
Minimise $|\{d \in D_v : a \cup \{(v, d)\} \text{ consistent}\}|$
- ◇ Minimises branching factor (at the current node)
- ◇ Extreme case: select variables with unique possible values first
 - Value is forced by the existing assignment
 - Obviously should be done in all cases
 - Compare unit propagation in SAT solving


Other variable ordering strategies

- ◇ Most constraining variable
Involved in as many constraints as possible
Maximise $|\{u \in V : a(u) \text{ undefined}, C_{u,v} \in C\}|$
- ◇ Seek biggest effect on domains of unassigned variables
Detect inconsistencies earlier, shortening search tree branches
- ◇ Others include history-dependent strategies
 - e.g. involved in a lot of (recent) conflicts
 - or selected many/few times before
- ◇ Random selection can also help, especially for tie-breaking

Value ordering


- ◇ Common strategy: least constraining value
Choose a value that won't conflict much with others
Minimise $|\{\{d' \in D_u : a(u) \text{ undefined}, C_{u,v} \in C, (d, d') \notin C_{u,v}\}\}|$
- ◇ Minimise useless backtracking below current node
- ◇ If no solutions, or if we want all solutions, value ordering doesn't matter: we have to go over the whole sub-tree anyway.
- ◇ If there is a solution, we may be lucky and find it without backtracking on this value choice

Forward Checking plus First-Fail




- ◇ Forward checking is rather weak on its own, but it combines well with the first-fail heuristic for variable ordering, to make a powerful technique.
- ◇ Unit propagation (selecting variables with singleton domains) is particularly important when forward checking is used.

Forward Checking plus First-Fail




- ◇ Impossible values get removed from related domains

Forward Checking plus First-Fail




◇ Note that there is only one value in the middle

Forward Checking plus First-Fail




- ◇ Colour that one green, as it has the smallest domain
- ◇ More domains are reduced to singletons

Forward Checking plus First-Fail




◇ Forced move

Forward Checking plus First-Fail




◇ Forced move

Forward Checking plus First-Fail




◇ Forced move

Forward Checking plus First-Fail




◇ Forced move

Forward Checking plus First-Fail



◇ Forced move

Forward Checking plus First-Fail




◇ Search was backtrack-free!


Arc Consistency

- ◇ A stronger inference rule: make all variables **arc consistent**
- ◇ Variable v is arc consistent with respect to another variable u iff for every $d \in D_v$ there is at least one $d' \in D_u$ such that $(d, d') \in C_{v,u}$. A CSP $\gamma = (V, D, C)$ is said to be arc consistent (AC) iff every variable in V is arc consistent with every other.
- ◇ Any $d \in D_v$ which has no support in D_u is incapable of being assigned to v in any solution, so it can be removed from D_v .
- ◇ Removing all unsupported values makes γ AC. This is clearly a valid constraint inference, as no solutions are lost.
- ◇ Enforcing AC subsumes both forward checking and unit propagation.


Arc Consistency



Arc Consistency




Arc Consistency




- ◇ Already done: since this is AC, the only possible assignment must be a solution.
- ◇ Now it's just a matter of filling in the values

Arc Consistency




- ◇ Already done: since this is AC, the only possible assignment must be a solution.
- ◇ Now it's just a matter of filling in the values

Arc Consistency




- ◇ Already done: since this is AC, the only possible assignment must be a solution.
- ◇ Now it's just a matter of filling in the values

Arc Consistency




- ◇ Already done: since this is AC, the only possible assignment must be a solution.
- ◇ Now it's just a matter of filling in the values

Arc Consistency




- ◇ Already done: since this is AC, the only possible assignment must be a solution.
- ◇ Now it's just a matter of filling in the values

Arc Consistency




- ◇ Already done: since this is AC, the only possible assignment must be a solution.
- ◇ Now it's just a matter of filling in the values

Arc Consistency



- ◇ Already done: since this is AC, the only possible assignment must be a solution.
- ◇ Now it's just a matter of filling in the values

Arc Consistency




◇ Search was backtrack-free—and all over at step 2

Arc consistency: AC-3


```
function REVISE( $\gamma, u, v$ ) returns modified  $\gamma$ 
  for each  $d \in D_u$  do
    if there is no  $d' \in D_v$  with  $(d, d') \in C_{u,v}$  then
       $D_u \leftarrow D_u \setminus \{d\}$ 
  end
  return  $\gamma$ 
```

```
function AC-3( $\gamma$ ) returns modified  $\gamma$ 
   $M \leftarrow \{(u, v), (v, u) : C_{u,v} \in C\}$ 
  while  $M \neq \{\}$  do
    remove some element  $(u, v)$  from  $M$ 
     $\gamma \leftarrow \text{REVISE}(\gamma, u, v)$ 
    if  $D_u$  has changed then
       $M \leftarrow M \cup \{(w, u) : C_{w,u} \in C, w \neq v\}$ 
  end
  return  $\gamma$ 
```

AC-3 Example




AC-3 Example




$$M = \{(3, 2), (5, 2), (6, 2)\}$$

AC-3 Example




$$M = \{(5, 2), (6, 2), (6, 3)\}$$

AC-3 Example




$$M = \{(6, 2), (6, 3), (4, 5), (6, 5), (8, 5), (9, 5)\}$$

AC-3 Example




$$M = \{(6, 3), (4, 5), (6, 5), (8, 5), (9, 5), (3, 6), (5, 6), (9, 6)\}$$

AC-3 Example




$$M = \{(6, 5), (8, 5), (9, 5), (3, 6), (5, 6), (9, 6), (7, 4), (8, 4)\}$$

AC-3 Example




$$M = \{(8, 5), (9, 5), (3, 6), (5, 6), (9, 6), (7, 4), (8, 4), (2, 6)\}$$

AC-3 Example




$$M = \{(9, 5), (3, 6), (5, 6), (9, 6), (7, 4), (8, 4), (2, 6), (4, 8), (7, 8), (9, 8)\}$$

AC-3 Example




$$M = \{(3, 6), (5, 6), (9, 6), (7, 4), (8, 4), (2, 6), (4, 8), (7, 8), (9, 8), (6, 9), (8, 9)\}$$

AC-3 Example




$$M = \{(5, 6), (9, 6), (7, 4), (8, 4), (2, 6), (4, 8), (7, 8), (9, 8), (6, 9), (8, 9), (2, 3)\}$$

AC-3 Example




$$M = \{(7, 4), (8, 4), (2, 6), (4, 8), (7, 8), (9, 8), (6, 9), (8, 9), (2, 3), (5, 9)\}$$

AC-3 Example




$$M = \{(8, 4), (2, 6), (4, 8), (7, 8), (9, 8), (6, 9), (8, 9), (2, 3), (5, 9, (8, 7))\}$$

AC-3 Example



$$M = \{(2, 6), (4, 8), (7, 8), (9, 8), (6, 9), (8, 9), (2, 3), (5, 9), (8, 7), (5, 8)\}$$

AC-3 Example



$$M = \{(9, 8), (6, 9), (8, 9), (2, 3), (5, 9), (8, 7), (5, 8), (4, 7)\}$$

Arc consistency: notes

- ◇ At every iteration, all arcs not in the set M are consistent
 - M contains the arcs that still need to be checked
 - M often implemented as a queue, but it doesn't have to be
- ◇ On termination, the network is AC
- ◇ Unlike forward checking, makes inferences from unassigned variables
- ◇ Arc consistency is widely used in modern CSP solvers
- ◇ Slower (per node) than forward checking, but prunes more

Summary

- ◇ **Variable orderings** in backtracking can dramatically reduce the size of the search tree. **Value orderings** don't, but they may lead to solutions earlier.
- ◇ **Inference** tightens γ without losing equivalence, during backtracking. This reduces the amount of search needed. The benefit in reduced tree size must be traded off against the time cost of the reasoning.
- ◇ **Forward checking** removes values conflicting with an assignment already made
- ◇ **Arc consistency** extends this to all variables, whether assigned or not. It is stronger than forward checking and unit propagation, but costs more to compute.