

KNOWLEDGE REPRESENTATION AND REASONING: CONSTRAINT SATISFACTION AND LOCAL SEARCH

CHAPTERS 4.1, 6.4

Constraint Satisfaction Problems



- ◇ Binary constraint network $\gamma = \langle V, D, C \rangle$
- V a finite set of variables v_1, \dots, v_n
 - D a set of [finite] sets D_{v_1}, \dots, D_{v_n}
 - C a set of binary relations $\{C_{u,v} \mid u, v \in V, u \neq v\}$
 $C_{u,v} \subseteq D_u \times D_v$

Outline of the lecture

- ◇ Recall [optimal] constraint solving
- ◇ Local Search in general
- ◇ Large Neighbourhood Search in particular
- ◇ Constraint modelling
- ◇ Summary

Recall

- ◇ CSP may be solved for **satisfiability** (any solution will do)
- ◇ May instead require **optimality** (best solution)
- ◇ **Objective** function (i.e. cost) measures badness of solutions
- ◇ FD solvers commonly use **Depth-First Branch and Bound** (DFBB)
 - Use a lower bound L on the objective and an upper bound B
 - Backtrack whenever $L \geq B$
 - Revise B whenever a solution is found
- ◇ DFBB fits well with backtracking CSP solution methods
- ◇ Gives a sequence of monotonically improving solutions

There is a problem

- ◇ Absolute optimality is often too hard to compute
- ◇ Need methods that scale up better and yield (probably) good solutions
- ◇ Enter **Local Search**
(actually a generic name for many search techniques)
- ◇ Widely used in industrial applications

Local Search


- ◇ The general idea: search in the space of **total** assignments
- ◇ An alternative to backtracking: not so rigidly defined
 - Usually involves randomness at some point
- ◇ Example (for satisfiability):
 - Start with a random assignment of values to all variables
(Of course, this doesn't satisfy all constraints)
 - Repeatedly:
 - Choose a variable (random choice is good)
 - Revise its value to minimise its constraint violations
 - Stop when all constraints are satisfied
- ◇ Optimisation version might remember the best assignment so far, and stop when the objective function hasn't improved for a while.

Iterative improvement algorithms

- ◇ Local search of this kind iteratively improves total assignments
- ◇ General idea: keep a single “current” state, try to improve it
 - perform **local** moves in the **neighbourhood** of the current state
 - no guarantee of completeness (may fail to find any solution)
 - no guarantee of optimality
 - no possibility of showing unsatisfiability
- ◇ However, method scales up better than complete search in many domains
- ◇ Small memory requirements, suitable for online as well as offline search

Example: Travelling Salesperson Problem

Start with any complete tour, perform pairwise exchanges to reduce tour duration




Variants of this approach get within 1% of optimal very quickly with thousands of cities



Example: n -queens

Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal



Move a queen in its column to reduce number of conflicts



$h = 5$



$h = 2$



$h = 0$

Variants of this approach almost always solve n -queens problems almost instantaneously for very large n , e.g., $n = 1$ **million**

Hill-climbing (or gradient ascent/descent)

Moves to the best neighbour


Climbing a mountain in the fog without a map: just go up!

```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  inputs: problem, a problem
  local variables: current, a node
                     neighbour, a node

  current ← MAKE-NODE(INITIAL-STATE[problem])
  loop do
    neighbour ← a highest-valued successor of current
    if VALUE[neighbour] ≤ VALUE[current] then return STATE[current]
    current ← neighbour
  end
```

Hill-climbing contd.

Useful to consider state space landscape



Random-restart hill climbing overcomes local maxima—trivially complete

Random sideways moves 😊 escape from shoulders 😞 loop on flat maxima

8 queens: simple HC has 14% success rate, HC + sideways moves 94%

3 million queens: HC + random restart + sideways moves needs < 1 mn

Simulated annealing

Idea: escape local maxima by allowing some “bad” moves
but gradually decrease their badness and frequency

```
function SIMULATED-ANNEALING(problem, schedule) returns a solution state
  inputs: problem, a problem
           schedule, a mapping from time to “temperature”
  local variables: current, a node
                    neighbour, a node
                    T, a “temperature” controlling prob. of downward steps

  current ← MAKE-NODE(INITIAL-STATE[problem])
  for t ← 1 to ∞ do
    T ← schedule[t]
    if T = 0 then return current
    neighbour ← a randomly selected successor of current
     $\Delta E \leftarrow \text{VALUE}[\textit{neighbour}] - \text{VALUE}[\textit{current}]$ 
    if  $\Delta E > 0$  then current ← neighbour
    else current ← neighbour only with probability  $e^{\Delta E/T}$ 
```

Population-based methods

- ◇ **Local beam search**: maintain a population of k states; add some neighbours at each step; delete the worst ones to keep the population stable
 - As usual, many variants exist
- ◇ **Genetic (evolutionary) algorithms**: define “crossover” between pairs of states in the population (compare genetics); offspring have some features of each parent; cull according to a “fitness” measure
 - Much research over several decades
 - Again, many variations on the general method

Population-based search and simulated annealing will not be covered in this course: COMP4660 / COMP8420 covers evolutionary algorithms in some detail

A hybrid: Large Neighbourhood Search

- ◇ Search in the space of **solutions** rather than **states**
- ◇ Given a current solution:
 - Destroy part of it by forgetting the values of some variables
 - See the problem of assigning values to those variables as a CSP
 - Solve [optimally] using a complete search method such as DFBB
- ◇ Alternates **destroy** and **repair** phases
 - Repair phase searches a neighbourhood of the current solution
 - Neighbourhood size remains tolerable, even if the problem is huge
- ◇ The old solution is still in the neighbourhood, so there is always a next solution available, given enough search
- ◇ May search for the **optimal** solution in the neighbourhood, or for an improvement on the previous solution, or just for **any** solution in the neighbourhood.

LNS: Example (1)

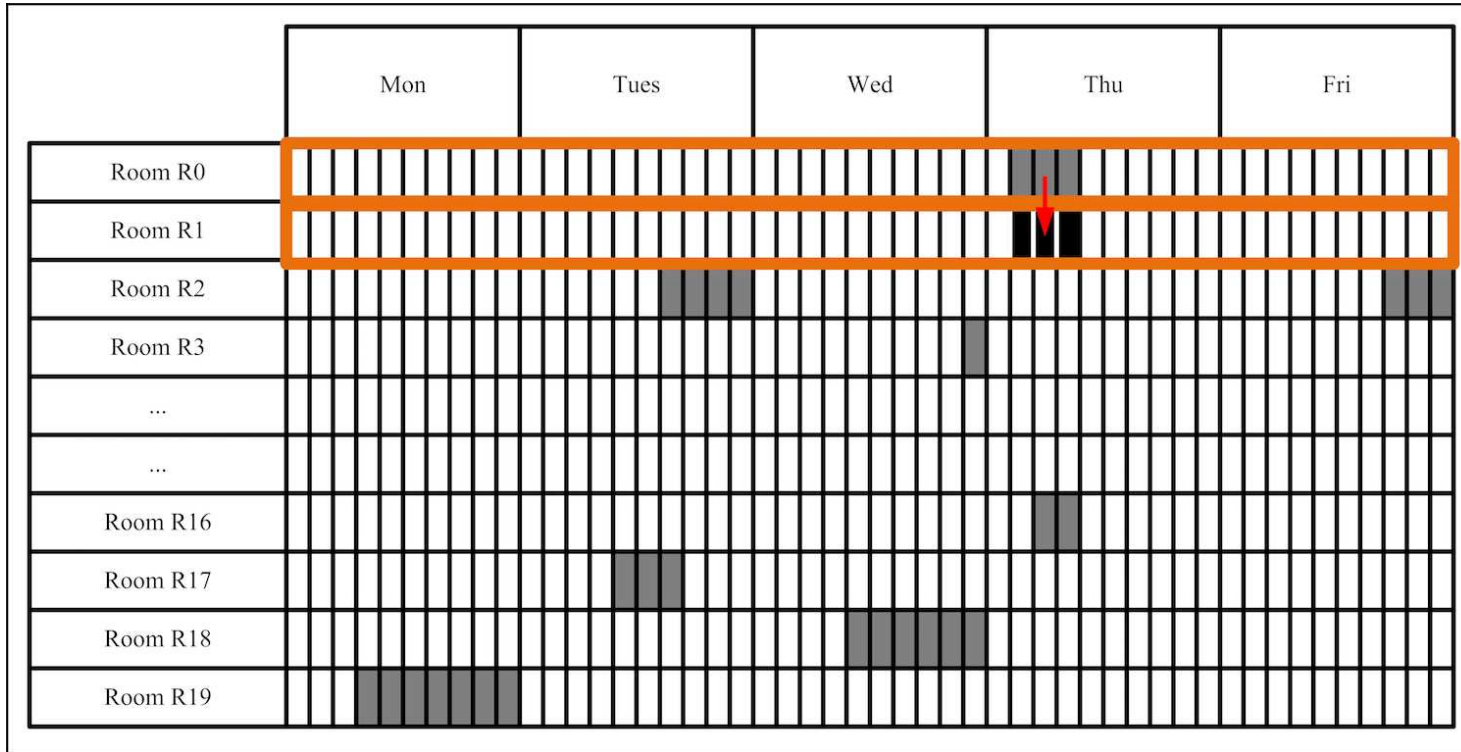
- ◇ Energy-aware scheduling of meetings in a large building or buildings
- ◇ Each meeting has:
 - a set of participants
 - a set of allowable locations (rooms)
 - a set of allowable times
 - a duration (30 minutes, 1 hour, 2 hours, ...)
- ◇ Constraints:
 - All classes are scheduled exactly once
 - All particular constraints on meetings are met
 - No two meetings in the same room at the same time
 - No conflict of meetings for participants
- ◇ Objective function: energy consumption (separately calculated)
- ◇ With 200 meetings in 20 rooms, too hard for DFBB

LNS: Example (2)

	Mon	Tues	Wed	Thu	Fri
Room R0					
Room R1					
Room R2					
Room R3					
...					
...					
Room R16					
Room R17					
Room R18					
Room R19					

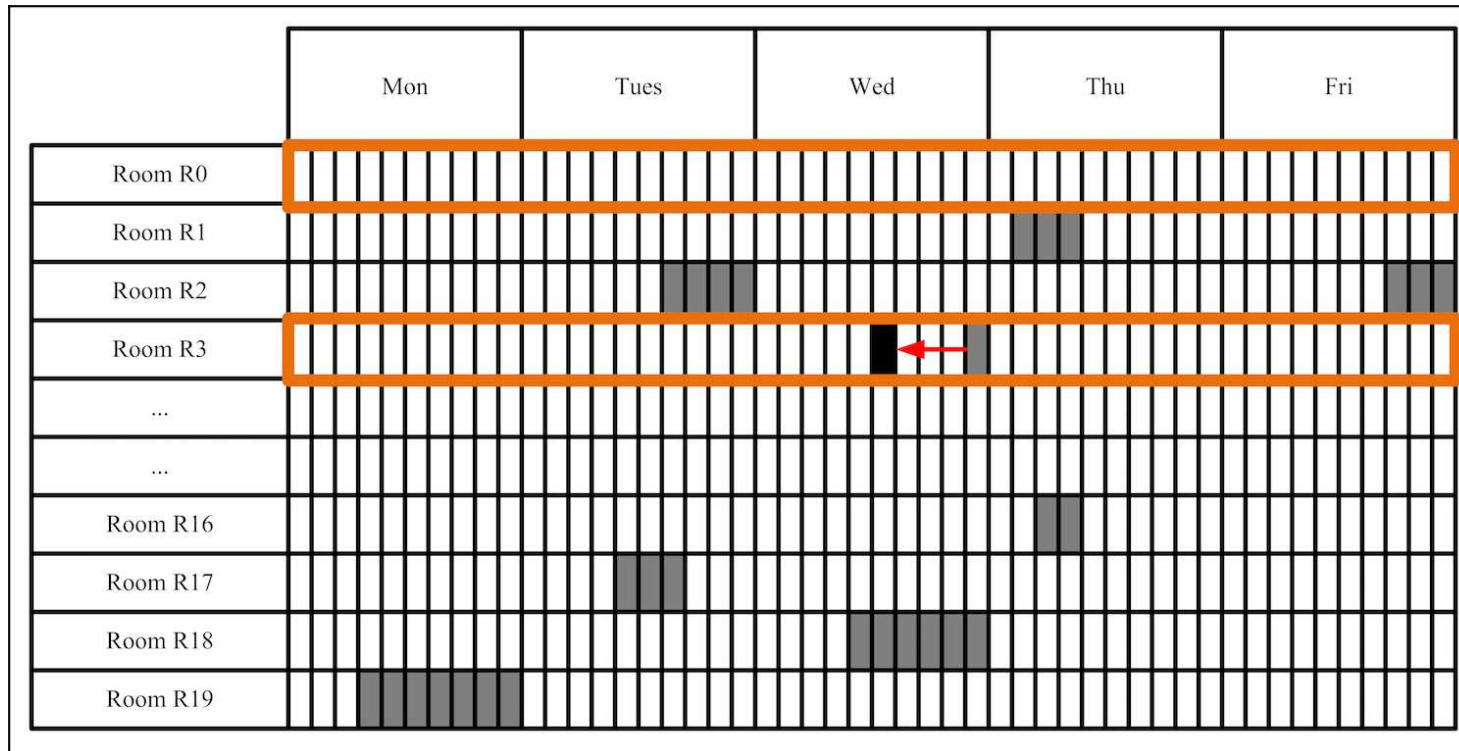
◇ Initial schedule generated by the MIP solver

LNS: Example (2)




◇ Schedule for rooms R0 and R1 destroyed and re-optimised

LNS: Example (2)




◇ Schedule for rooms R0 and R3 destroyed and re-optimised

LNS: Example (2)



◇ Schedule for rooms R0 and R2 destroyed and re-optimised

LNS: Example (3)



- ◇ MIP (complete) solver found fairly good solutions in 2 hours
- ◇ LNS using the same solver found better ones in 15 minutes

LNS: Notes

- ◇ The **initial solution** must come from somewhere.
- ◇ Performance is sensitive to the choice of what to destroy
- ◇ We may choose to **abstract** from the current solution
 - use only some decision variables, for a partial description
 - designed so the rest can be recovered by easy search
 - destroy part of the abstract solution
 - gives the complete search freedom to optimise minor aspects
- ◇ **Local optima** are still a problem, as with all local search
 - random restart is commonly used to escape
- ◇ There is always a **tradeoff** between neighbourhood size and speed
 - Large neighbourhoods increase the chance of improvement
 - but they may create hard problems for the complete search

Summary

- ◇ Local search explores the space of total assignments
- ◇ Usually step to a neighbour, looking for improvement
 - but sometimes jump further, or even re-start entirely
- ◇ No guarantees (incomplete, sub-optimal, ...)
 - But in many cases scales up better than systematic search
- ◇ Many varieties
 - hill-climbing, random walks, simulated annealing, population methods
- ◇ Large neighbourhood search the best of both worlds (?)
 - DFBB or similar to solve small problems quickly and completely
 - local search for scaling up
 - tradeoff between the two is fundamental
 - requires experiment to tune parameters and define neighbourhoods