

APPLIED STATISTICS

Logistic Regression for Binomial Counts

Dr Tao Zou

Research School of Finance, Actuarial Studies & Statistics
The Australian National University

Last Updated: Mon Oct 9 21:12:17 2017

Overview

- Motivating Example
- Logistic Regression for Binomial Counts
- Model Diagnostics
 1. Logit proportion versus explanatory variable plot.
 2. Pearson residual plot.
 3. Deviance goodness-of-fit test.

References

1. **F.L. Ramsey and D.W. Schafer** (2012)
Chapter 21 of *The Statistical Sleuth*
2. ANU STAT3015 Lecture Notes
3. The slides are made by **R Markdown**.
<http://rmarkdown.rstudio.com>

Example: Anaesthetic Data (Con'd)

(Taken from STAT3015 notes.)

Thirty patients were given a particular anaesthetic at various predetermined concentrations for 15 minutes before a stimulus was applied.

The response variable was simply an indication as to whether the patient responded to the stimulus in any way.

“Response” is 1 if the patient responded to the stimulus.

Standard Data Form

```
setwd('~\\Desktop\\Research\\AppliedStat2017\\L12')  
a=read.csv('anaesthetic.csv');a
```

##	Concentration	Response
## 1	0.8	1
## 2	0.8	1
## 3	0.8	1
## 4	0.8	1
## 5	0.8	1
## 6	0.8	1
## 7	0.8	0
## 8	1.0	1
## 9	1.0	1
## 10	1.0	1
## 11	1.0	1
## 12	1.0	0
## 13	1.2	1
## 14	1.2	1
## 15	1.2	0
## 16	1.2	0
## 17	1.2	0
## 18	1.2	0
## 19	1.4	1
## 20	1.4	1
## 21	1.4	0
## 22	1.4	0
## 23	1.4	0
## 24	1.4	0
## 25	1.6	0
## 26	1.6	0
## 27	1.6	0
## 28	1.6	0
## 29	2.5	0
## 30	2.5	0

Some Terminologies to Interpret Standard Data

“Trial”: Thirty patients correspond to thirty trials for this anaesthetic experiment.

Different trials are independent.

“Success”: one trial is called “success” if in this trial, the patient responds to the stimulus, or equivalently, $\text{Response} = 1$.

“Failure”: one trial is called “failure” if in this trial, the patient does not respond to the stimulus, or equivalently, $\text{Response} = 0$.

Probability of “Success”: the probability that one trial is a “success”, **or equivalently**, the probability that the patient responds to the stimulus in the trial, **or equivalently**, the probability that $\text{Response} = 1$.

Aggregated Data Form

Note that in the data, we have repeated concentration (X) values. Can we aggregate response Y by X ?

The number of “successes” (i.e. Response = 1 here) given some $X = x$:

```
attach(a)
tapply(Response, Concentration, sum)
```

```
## 0.8  1 1.2 1.4 1.6 2.5
##  6  4  2  2  0  0
```

The number of “trials” (equivalent to patients here) given some $X = x$:

```
tapply(Response, Concentration, length)
```

```
## 0.8  1 1.2 1.4 1.6 2.5
##  7  5  6  6  4  2
```

The aggregated dataset is

```
b=data.frame(Concentration=unique(Concentration), BinomialCount=tapply(Response, Concentration, sum),
             TotalNumber=tapply(Response, Concentration, length));b
```

```
##      Concentration BinomialCount TotalNumber
## 0.8              0.8              6          7
## 1              1.0              4          5
## 1.2             1.2              2          6
## 1.4             1.4              2          6
## 1.6             1.6              0          4
## 2.5             2.5              0          2
```

Binomial Count

In a lot of real data, we only have **the aggregated data form** instead of the standard data form. Specifically in this example, we have

1. at a specific concentration, we have **the number of "successes"** (the number of patients that respond to the stimulus), which is called **binomial count**, denoted by Z ;
2. at a specific concentration, we have **the total number of "trials"** (the total number of patients given this specific level of anaesthetic concentration), denoted by M .

Logistic regression for binomial counts deals with the particular case of the aggregated data, and is used to model the binomial count response Z .

However, the total number of "trials" M should also be observed such that we can employ this model.

Some Terminologies to Interpret Aggregated Data

A typical row for an aggregated dataset:

Explanatory	Binomial Count	Total Number
X	Z	M

Z is the number of “successes” given some specific X .

M is the number of “trials” given some specific X .

Given X , the probability of one “success” in a trial is denoted by $\pi \in [0, 1]$.

The binomial count Z is modelled by $\text{Binomial}(M, \pi)$ distribution.

Specifically,

$$P(Z = z) = \binom{M}{z} \pi^z (1 - \pi)^{M-z}, \quad \binom{M}{z} = \frac{M(M-1) \cdots (M-z+1)}{z(z-1) \cdots 1},$$

for $z = 0, \dots, M$.

Given X , the proportion of “successes” is Z/M , which can be used to estimate π when M is large.

Overview of This Course

	Continuous X + Categorical X
Continuous Y	MLR + Indicator Variables
Two-Category Y	Binary Logistic Regression + Indicator Variables
Multicategory Y - Nominal	Nominal Response Regression + Indicator Variables
Multicategory Y - Ordinal	Ordinal Response Regression + Indicator Variables
Binomial Count Z	Binomial Logistic Regression + Indicator Variables

Binomial Logistic Regression Model Assumptions

- 1. Binomial distribution:** There is a Binomial distributed (sub)population of responses Z for given values of the explanatory variables $(X_1 = x_1, \dots, X_k = x_k)$. That means if we let $X = (X_1, \dots, X_k)$, the probability that $Z = z$ given X is

$$P(Z = z) = \binom{M}{z} \pi^z (1 - \pi)^{M-z}.$$

- 2. Generalised Linearity:** The transformation of the probability of “success” π falls on a linear function of the explanatory variables

$$g(\pi) = \beta_0 + \beta_1 X_1 + \dots + \beta_k X_k, \text{ for } X = (X_1, \dots, X_k),$$

where $g(u) = \log\{u/(1 - u)\}$, which is called logit link function.

Binary Logistic Regression Model Assumptions (Con'd)

Remark: the inverse function of the logit link function is

$$g^{-1}(v) = \frac{e^v}{1 + e^v} \in [0, 1].$$

Then

$$\pi = g^{-1}(\beta_0 + \beta_1 X_1 + \cdots + \beta_k X_k) \in [0, 1],$$

which is consistent with the range $\pi \in [0, 1]$.

3. Independence: Observations

$$(X_{1,1}, \cdots X_{k,1}, Z_1, M_1),$$

$$\vdots$$

$$(X_{1,m}, \cdots X_{k,m}, Z_m, M_m),$$

are independent, where m is the sample size.

The interpretation of the binomial logistic regression model stays largely the same as for the binary logistic regression model.

Estimation, z-Test and CI

The likelihood function for the observations and the MLE can be obtained. The inferential tools for generalised linear model can be used.

```
attach(b)
Y=cbind(BinomialCount,TotalNumber-BinomialCount)
fit=glm(Y~Concentration,family=binomial(link=logit))
summary(fit)

##
## Call:
## glm(formula = Y ~ Concentration, family = binomial(link = logit))
##
## Deviance Residuals:
##      0.8      1      1.2      1.4      1.6      2.5
## -0.20147  0.45367 -0.56890  0.70000 -0.81838 -0.04826
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      6.469      2.419   2.675  0.00748 **
## Concentration    -5.567      2.044  -2.724  0.00645 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 15.4334  on 5  degrees of freedom
## Residual deviance:  1.7321  on 4  degrees of freedom
## AIC: 13.811
##
## Number of Fisher Scoring iterations: 5
```

Drop-in-Deviance χ^2 -Test

$$H_0 : \beta_1 = 0 \leftrightarrow H_a : \text{otherwise.}$$

```
#drop in deviance test  
#reduced model  
fitr=glm(Y~1,family=binomial(link=logit))  
anova(fitr,fit,test='Chisq')
```

```
## Analysis of Deviance Table
```

```
##
```

```
## Model 1: Y ~ 1
```

```
## Model 2: Y ~ Concentration
```

```
##   Resid. Df Resid. Dev Df Deviance  Pr(>Chi)
```

```
## 1         5      15.4334
```

```
## 2         4       1.7321  1    13.701 0.0002143 ***
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Fitted Probabilities of “Success”

Using MLE $\hat{\beta}_0, \dots, \hat{\beta}_k$, the estimated (fitted) probability of “success” is given by:

$$\begin{aligned}\hat{\pi} &= g^{-1}(\hat{\beta}_0 + \hat{\beta}_1 X_1 + \dots + \hat{\beta}_k X_k) \\ &= \frac{e^{\hat{\beta}_0 + \hat{\beta}_1 X_1 + \dots + \hat{\beta}_k X_k}}{1 + e^{\hat{\beta}_0 + \hat{\beta}_1 X_1 + \dots + \hat{\beta}_k X_k}}.\end{aligned}$$

It is worth noting that

$$g(\hat{\pi}) = \log\left(\frac{\hat{\pi}}{1 - \hat{\pi}}\right) = \hat{\beta}_0 + \hat{\beta}_1 X_1 + \dots + \hat{\beta}_k X_k,$$

where function g is the logit link function. Hence the logit transformation of fitted probabilities should be on a straight line.

When we talk about “fitted”, X is usually from the training dataset (see Lecture Notes 8).

When X_{new} is from the new dataset or the test dataset, we actually talk about prediction.


Predicted Probabilities of “Success”

The predicted probability of “success” is given by

$$\begin{aligned}\hat{\pi} &= g^{-1}(\hat{\beta}_0 + \hat{\beta}_1 X_{1,\text{new}} + \cdots + \hat{\beta}_k X_{k,\text{new}}) \\ &= \frac{e^{\hat{\beta}_0 + \hat{\beta}_1 X_{1,\text{new}} + \cdots + \hat{\beta}_k X_{k,\text{new}}}}{1 + e^{\hat{\beta}_0 + \hat{\beta}_1 X_{1,\text{new}} + \cdots + \hat{\beta}_k X_{k,\text{new}}}}.\end{aligned}$$

Example: Anaesthetic Data (Con'd)

```
plot(Concentration, log(fit$fitted.values/(1-fit$fitted.values)),  
     xlab="concentration", ylab="fitted probabilities")
```



The logit transformations of fitted probabilities are on a straight line.

Example: Anaesthetic Data (Con'd)

By using this example, we might be interested in predicting the probability of “success” if an anaesthetic at a new concentration of 1.5 is given. The forecast of probability is:

```
Xnew=data.frame(Concentration=1.5)  
predict(fit,Xnew,type='response')
```

```
##           1  
## 0.1322204
```

Fitted Values of Response and Prediction of Response

A typical row for an aggregated data:

Explanatory	Binomial Count	Total Number
X	Z	M

The binomial count Z is modelled by Binomial(M, π) distribution. Specifically,

$$P(Z = z) = \binom{M}{z} \pi^z (1 - \pi)^{M-z}, \text{ for } z = 0, \dots, M.$$

Based on the properties of the Binomial distribution, the mean of response Z is given by

$$\mu\{Z|X\} = M\pi.$$

Using MLE $\hat{\beta}_0, \dots, \hat{\beta}_k$, the fitted values of response Z are given by:

$$\hat{Z} = \hat{\mu}\{Z|X\} = M\hat{\pi}.$$

Example: Anaesthetic Data (Con'd)

By using this example, we might be interested in predicting the number of “successes” if an anaesthetic at a new concentration of 1.5 is given.

Suppose the number of “trials” is 5 in this case. The forecast should be

```
Mnew=5  
Mnew*predict(fit,Xnew,type='response')
```

```
##          1  
## 0.6611018
```

However, it is not an integer. So we consider

```
round(Mnew*predict(fit,Xnew,type='response'),0)
```

```
## 1  
## 1
```

roughly as the prediction of response Z .

1. Logit Proportion versus Explanatory Variable Plot

Recall that in the Binomial logistic regression model assumptions, the **logit** transformation of the probability of “success” π falls on a linear function of the explanatory variables

$$g(\pi) = \beta_0 + \beta_1 X_1 + \cdots + \beta_k X_k, \text{ for } X = (X_1, \cdots, X_k),$$

where $g(u) = \log\{u/(1-u)\}$ is the **logit** link function. Here the probability of “success” π depends on unknown parameters and hence is also unknown.

For a typical row for an aggregated data:

Explanatory	Binomial Count	Total Number
X	Z	M

The proportion of “successes” is Z/M given Z , which can be used to estimate π when M is large.


Hence it can be useful to plot the logit proportion $g(Z/M)$ versus explanatory variable (e.g., X_1).

1. Logit Proportion versus Explanatory Variable Plot (Con'd)

The plot should show a straight line. Otherwise the model assumption is violated. See Lecture Notes 3 for solutions to this problem.

Note: $g(Z/M)$ is undefined for Z/M equal to 0 or 1. We need to add (subtract) a small quantity to values of 0 (1).

```
proportion=BinomialCount/TotalNumber
proportion=ifelse(proportion==0, 0.01, proportion)
proportion=ifelse(proportion==1, 1-0.01, proportion)
logitprop=log(proportion/(1-proportion))
plot(Concentration,logitprop)
```



Review: Studentized Residuals for Continuous Response Y

Recall residual: $\text{res}_i = \hat{\mathcal{E}}_i = Y_i - \hat{Y}_i$ for observation i , where \hat{Y}_i is the fitted value of response.

One can obtain

$$\text{SD}(\text{res}_i) = \sigma \sqrt{1 - h_i}, \text{ and}$$

$$\text{SE}(\text{res}_i) = \hat{\sigma} \sqrt{1 - h_i}, \text{ where } \hat{\sigma} = \sqrt{\frac{\sum_{i=1}^n \text{res}_i^2}{n - k - 1}}.$$

A studentized residual is a residual divided by its standard error, namely

$$\text{studres}_i = \frac{\text{res}_i}{\text{SE}(\text{res}_i)}$$

Using studentized residuals allows the residuals to be viewed on the same scale.

These studentized residuals are roughly standard normally distributed $[N(0, 1)]$, if the observation is from the MLR model with the all the assumptions satisfied.

Pearson Residuals for Binomial Count Z

The residual for Binomial count Z is defined as: $\text{res}_i = Z_i - \hat{Z}_i$ for observation i , where \hat{Z}_i is the fitted value of response.

One can obtain

$$\text{SD}(\text{res}_i) = \sqrt{M_i \pi_i (1 - \pi_i)}, \text{ and}$$

$$\text{SE}(\text{res}_i) = \sqrt{M_i \hat{\pi}_i (1 - \hat{\pi}_i)}, \text{ where } \hat{\pi}_i = \frac{e^{\hat{\beta}_0 + \hat{\beta}_1 X_{1,i} + \dots + \hat{\beta}_k X_{k,i}}}{1 + e^{\hat{\beta}_0 + \hat{\beta}_1 X_{1,i} + \dots + \hat{\beta}_k X_{k,i}}}.$$

A Pearson residual is a residual divided by its standard error, namely

$$\text{Peares}_i = \frac{\text{res}_i}{\text{SE}(\text{res}_i)}$$

Using Pearson residuals allows the residuals to be viewed on the same scale.

For large values of M_i (>5), these Pearson residuals are roughly standard normally distributed $[N(0, 1)]$, if the observation is from the Binomial logistic model with the all the assumptions satisfied.

2. Pearson Residual Plot

Due to the nature of the $N(0, 1)$ distribution, most values of the $N(0, 1)$ distribution concentrate in the middle around 0.

Hence, if Peares_i falls into the two tails of the $N(0, 1)$ distribution, namely $|\text{Peares}_i|$ is too large, then it is unlikely that observation i is from the Binomial logistic model with the all the assumptions satisfied.

$\text{Peares}_i > 1.96$ (97.5 % quantile of $N(0, 1)$), or
 $\text{Peares}_i < -1.96$ (2.5 % quantile of $N(0, 1)$).

\Rightarrow

Peares_i falls into the two tails of the $N(0, 1)$ distribution.

\Rightarrow

$|\text{Peares}_i|$ is too large.


\Rightarrow

Observation i is **an outlier**.

Sometimes we use 2 instead of 1.96 for simplicity.

2. Pearson Residual Plot (Con'd)

```
Peares=residuals(fit,type="pearson")  
plot(Peares)  
abline(h=1.96,col='red')  
abline(h=-1.96,col='red')
```



The Pearson residuals do not indicate any problems - none are outside the range of $[-1.96, 1.96]$ and no pattern is evident.

3. Deviance Goodness-of-Fit Test

H_0 : Binomial model $g(\pi) = \beta_0 + \beta_1 X_1 + \cdots + \beta_k X_k$ is appropriate \leftrightarrow

H_a : The Binomial logistic model is not appropriate.

The test statistic is

$$TS = \text{deviance of model } g(\pi) = \beta_0 + \beta_1 X_1 + \cdots + \beta_k X_k.$$

This test statistic should be compared to approximately χ^2_{m-k-1} distribution, where k is the number of explanatory variables.

The p -value is

$$p\text{-value} = P(S > TS), \text{ where } S \sim \chi^2_{m-k-1}.$$

If $p\text{-value} < \alpha \Rightarrow$ reject H_0 ; $p\text{-value} \geq \alpha \Rightarrow$ not reject H_0 .

3. Deviance Goodness-of-Fit Test (Con'd)

The output from the `summary()` output can be used to perform the goodness-of-fit test.

```
summaryfit=(summary(fit))
summaryfit
```

```
##
## Call:
## glm(formula = Y ~ Concentration, family = binomial(link = logit))
##
## Deviance Residuals:
##      0.8      1      1.2      1.4      1.6      2.5
## -0.20147  0.45367 -0.56890  0.70000 -0.81838 -0.04826
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      6.469      2.419   2.675  0.00748 **
## Concentration    -5.567      2.044  -2.724  0.00645 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 15.4334  on 5  degrees of freedom
## Residual deviance:  1.7321  on 4  degrees of freedom
## AIC: 13.811
##
## Number of Fisher Scoring iterations: 5
```

3. Deviance Goodness-of-Fit Test (Con'd)

The test statistic and the number of degrees of freedom are

```
summaryfit$deviance
```

```
## [1] 1.732136
```

```
summaryfit$df.residual
```

```
## [1] 4
```

The p -value for the test is given by

```
1 - pchisq(summaryfit$deviance, summaryfit$df.residual)
```

```
## [1] 0.7848721
```

```
detach(b)
```

We cannot reject the null. There is no evidence that the model is inappropriate.