

## STAT7017 Final Project

Rui Qiu

2018-10-08

## 0. Setup

First, we load some required libraries for this project. Basically, they are libraries of `tidyverse` which are responsible for data frame pipelines and the rest are mainly focusing on some particular calculations. We will mention those functions later.

```
library(mvtnorm)
library(ggplot2)
library(ggfortify)
library(forecast)
library(readr)
library(dplyr)
library(reshape2)
library(heplots)
library(ks)
library(vars)
library(portes)

set.seed(7017)
```

## 1. Part 1 Testing Covariance Matrices

## 1.1 Q1 Reproduce Box's Test for Equality of Covariance Matrices

The followings are some customized helper functions in order to apply Box's M test and print out some intermediate values.

**We have also configured some arguments as flags so that it can be reused in the next few questions, but the output details might vary accordingly.**

```
box.u <- function(n.vec,p,g) {
  u <- (sum(1/(n.vec-1))-1/sum(n.vec-1))*((2*p**2+3*p-1)/(6*(p+1)*(g-1)))
  return(u)
}

log.det <- function(S) {
  return(unlist(determinant(S,logarithm=T))[[1]])
}

box.M <- function(p,g,n.vec,S.list,print.log=F) {
  S.pooled <- matrix(rep(0,p**2),ncol=p,nrow=p)

  for (i in 1:g) {
    S.pooled <- S.pooled + (n.vec[i]-1)*S.list[[i]]/sum(n.vec-1)
  }

  # for this Q1 only
  if (print.log) {
    cat("log determinant of S.pooled:",log.det(S.pooled),"\n")
  }

  M <- sum(n.vec-1)*log.det(S.pooled)
```

```

    for (i in 1:g) {
      M <- M - (n.vec[i]-1)*log.det(S.list[[i]])
    }
    return(M)
  }

box.C <- function(u, M) {
  return((1-u)*M)
}

M.test <- function(p,g,n.vec,S.list,test=F,summary=F) {
  u <- box.u(n.vec,p,g)
  if (test) {
    M <- box.M(p,g,n.vec,S.list,print.log=T)
  } else {
    M <- box.M(p,g,n.vec,S.list)
  }
  C <- box.C(u,M)
  nu <- 0.5*p*(p+1)*(g-1)
  if (summary) {
    cat("u =",u,"\n",
        "M =",M,"\n",
        "C =",C,"\n",
        "chi-sq with", nu, "degrees of freedom is", qchisq(.05,nu), "\n",
        "reject null hypothesis:", C > qchisq(.05,nu), "\n")
  }
  return(c(C,qchisq(.05,nu)))
}

```

the

When all the helper functions are settled, we are safe to reproduce example 6.12 with following Wisconsin Nursing Homes data. Here, as expected, we printed out an intermediate value in the middle of calculation, which should be the cause of differences in final values. Further explanation can be found in the following paragraph

```

M.test(4,3,c(271,138,107),
  list(matrix(c(.291,0,0,0,-.001,.011,0,0,.002,0,.001,0,.010,.003,0,.01),
    nrow=4,byrow=T),
    matrix(c(.561,0,0,0,.011,.025,0,0,.001,.004,.005,0,.037,.007,.002,.019),
    nrow=4,byrow=T),
    matrix(c(.261,0,0,0,.030,.017,0,0,.003,-0.0,.004,0,.018,.006,.001,.013),
    nrow=4,byrow=T)),
  test=T,summary=T)

```

```

## log determinant of S.pooled: -15.42671
## u = 0.01324946
## M = 218.8517
## C = 215.9521
## chi-sq with 20 degrees of freedom is 10.85081
## reject null hypothesis: TRUE
## [1] 215.95207 10.85081

```

Hence we reject the null hypothesis that  $\Sigma_1 = \Sigma_2 = \Sigma_3$ . We also notice that the reproduced valued of test statistic  $C = 215.9521$  differs from the original value which is 285.5. To investigate this, we include an intermediate value, the logarithm of determinant of  $\mathbf{S}_{\text{pooled}}$ . Our reproduced value is  $-15.42671$ , while the text has  $-15.564$ . The standalone computational error should be trivial, but  $\ln|\mathbf{S}_{\text{pooled}}|$  is iterated and subtracted 3 times when calculating the value of  $M$  (correspondingly,  $C$ ). Therefore, the difference of our test statistic is quite “large”, but of course, does not affect our conclusion anyways.

textbook

## 1.2 Q2 Trials on Classic Iris Data


### 1.2.1 EDA Explanatory Data Analysis

For the famous Fisher's Iris data, we still would like to have some exploratory data analysis before starting our covariance homogeneity test. For simplicity, we just take a quick look at the two scatterplots.


```
str(iris)
```

```
## 'data.frame': 150 obs. of 5 variables:
## $ Sepal.Length: num 5.1 4.9 4.7 4.6 5 5.4 4.6 5 4.4 4.9 ...
## $ Sepal.Width : num 3.5 3 3.2 3.1 3.6 3.9 3.4 3.4 2.9 3.1 ...
## $ Petal.Length: num 1.4 1.4 1.3 1.5 1.4 1.7 1.4 1.5 1.4 1.5 ...
## $ Petal.Width : num 0.2 0.2 0.2 0.2 0.2 0.4 0.3 0.2 0.2 0.1 ...
## $ Species : Factor w/ 3 levels "setosa","versicolor",...: 1 1 1 1 1 1 1 1 1 1 ...
```

```
ggplot(data=iris,aes(x=Sepal.Width, y=Sepal.Length)) +
  geom_point(aes(color=Species, shape=Species)) +
  ggtitle("Sepal Width vs Sepal Length") +
  theme_minimal()
```



```
ggplot(data=iris,aes(x=Petal.Width, y=Petal.Length)) +
  geom_point(aes(color=Species, shape=Species)) +
  ggtitle("Petal Width vs Petal Length") +
  theme_minimal()
```



In fact, the variance-covariance matrix between variables are really hard to evaluate graphically. So no direct conclusion can be generated at this point.

### 1.2.2 Method 1: Recycle Pre-defined Functions in Q1

The major reason that we implement Box's M Test with functional programming is the ability to recycle and reuse them! Here we only need to separate the data frame and treat them as inputs of the function `M.test()`.

```
versicolor <- iris %>%
  filter(Species=="versicolor") %>%
  dplyr::select(-"Species")
setosa <- iris %>%
  filter(Species=="setosa") %>%
  dplyr::select(-"Species")
virginica <- iris %>%
  filter(Species=="virginica") %>%
  dplyr::select(-"Species")
n1 <- nrow(versicolor)
n2 <- nrow(setosa)
n3 <- nrow(virginica)
n <- nrow(iris)

S1 <- cov(versicolor)
S2 <- cov(setosa)
S3 <- cov(virginica)
```

In this question, we are dealing with two null hypotheses. We do nothing but plug in the values respectively **as requested**

```
M.test(4,2,c(n1,n2),list(S1,S2),summary=T)
```

Firstly, we test  $H_0: \Sigma_1 = \Sigma_2$

```
## u = 0.04387755
## M = 69.87649
## C = 66.81048
## chi-sq with 10 degrees of freedom is 3.940299
```

```
## reject null hypothesis: TRUE
```

The result of this comparison is obvious, we reject  $H_0$ .

```
## [1] 66.810481 3.940299
```

Reject the null hypothesis.

For the second part, we test  $H_0: S_1=S_2=S_3$ .

```
M.test(4,3,c(n1,n2,n3),list(S1,S2,S3),summary=T)
```

```
## u = 0.03900227
```

```
## M = 146.6632
```

```
## C = 140.943
```

```
## chi-sq with 20 degrees of freedom is 10.85081
```

```
## reject null hypothesis: TRUE
```

```
## [1] 140.94305 10.85081
```

Again, reject the null hypothesis.

we reject

### 1.2.3 Method 2: Use heplots Package

Alternatively, for this particular question,

[insert link of package]

Or here, we could directly use `boxM()` function in Michael Friendly's `heplots` package, which is a package mainly focuses on visualizing hypothesis tests in multivariate linear models.

What we need to do here is to firstly eliminate `virginica` species and drop the level along our data pipeline, and do a hypothesis test for  $\Sigma_1 = \Sigma_2$ . After that, we would repeat for testing  $\Sigma_1 = \Sigma_2 = \Sigma_3$ .

The results are pretty neat:

```
#####  
# Use Michael Friendly's 'heplots' package directly.  
#####
```

```
iris2 <- iris %>%  
  filter(Species!="virginica") %>%  
  droplevels()  
res2 <- boxM(iris2[, 1:4], iris2[, "Species"])  
res2
```

```
##  
## Box's M-test for Homogeneity of Covariance Matrices  
##  
## data: iris2[, 1:4]  
## Chi-Sq (approx.) = 66.81, df = 10, p-value = 1.823e-10
```

```
summary(res2)
```

```
## Summary for Box's M-test of Equality of Covariance Matrices  
##  
## Chi-Sq: 66.81048  
## df: 10  
## p-value: 1.823e-10  
##  
## log of Covariance determinants:  
## setosa versicolor pooled  
## -13.06736 -10.87433 -11.25782  
##  
## Eigenvalues:  
## setosa versicolor pooled  
## 1 0.236455690 0.487873944 0.32938906  
## 2 0.036918732 0.072384096 0.07721833  
## 3 0.026796399 0.054776085 0.05032393  
## 4 0.009033261 0.009790365 0.01008296
```

```
##
## Statistics based on eigenvalues:
##           setosa  versicolor  pooled
## product  2.113088e-06 1.893828e-05 1.290601e-05
## sum      3.092041e-01 6.248245e-01 4.670143e-01
## precision 5.576122e-03 7.338788e-03 7.405502e-03
## max      2.364557e-01 4.878739e-01 3.293891e-01
res3 <- boxM(iris[, 1:4], iris[, "Species"])
res3

##
## Box's M-test for Homogeneity of Covariance Matrices
##
## data:  iris[, 1:4]
## Chi-Sq (approx.) = 140.94, df = 20, p-value < 2.2e-16
summary(res3)

## Summary for Box's M-test of Equality of Covariance Matrices
##
## Chi-Sq: 140.943
## df: 20
## p-value: < 2.2e-16
##
## log of Covariance determinants:
##      setosa versicolor virginica  pooled
## -13.067360 -10.874325  -8.927058  -9.958539
##
## Eigenvalues:
##      setosa versicolor virginica  pooled
## 1 0.236455690 0.487873944 0.69525484 0.44356592
## 2 0.036918732 0.072384096 0.10655123 0.08618331
## 3 0.026796399 0.054776085 0.05229543 0.05535235
## 4 0.009033261 0.009790365 0.03426585 0.02236372
##
## Statistics based on eigenvalues:
##           setosa  versicolor  virginica  pooled
## product  2.113088e-06 1.893828e-05 0.0001327479 4.732183e-05
## sum      3.092041e-01 6.248245e-01 0.8883673469 6.074653e-01
## precision 5.576122e-03 7.338788e-03 0.0169121236 1.304819e-02
## max      2.364557e-01 4.878739e-01 0.6952548382 4.435659e-01
```

The results are identical to our “home-brewed” version.

there is nothing extra we need to emphasize, the results are ...

### 1.3 Q3 what could go wrong with box's chisel approximation?

In this problem, we fix the parameter  $g$  to be 2 and  $p$  changes from 5 to 20 as requested. Since the problem itself does not mention how  $n_i$ 's behave, we also fix them to be  $n_1 = n_2 = 25 > 20$ . For each value of  $p$ , we iterate to resample data for  $sim = 1000$  times.

In addition, my personal preference for monitoring complex for-loops is to add a timer keeping track of how long it takes. This will be repeated several times in the next few questions.

```
sim <- 1000
fixed.g <- 2
p.vec <- c(5:15)
n1 <- 25
n2 <- 25

table <- data.frame()

timer <- Sys.time()
for (iter in 1:length(p.vec)) {
  p <- p.vec[iter]

  for (s in 1:sim) {
    X1 <- matrix(rnorm(p*n1), n1, p)
    X2 <- matrix(rnorm(p*n2, mean=0, sd=1), n2, p)
    X2.prime <- matrix(rnorm(p*n2, mean=0, sd=1.5), n2, p)
    S1 <- cov(X1)
    S2 <- cov(X2)
    S2.prime <- cov(X2.prime)

    table <- rbind(table, c(p, M.test(p, fixed.g, c(n1, n2), list(S1, S2)), 1))
    table <- rbind(table, c(p, M.test(p, fixed.g, c(n1, n2), list(S1, S2.prime)), 2))
  }
}

Sys.time()-timer

## Time difference of 19.42585 secs
names(table) <- c("p", "TS", "chisq", "type")

large.summary <- data.frame()
for (pv in p.vec) {
  nu <- 0.5*pv*(pv+1)*(fixed.g-1)

  table.size <- table %>%
    filter(p==pv, type==1)
  table.power <- table %>%
    filter(p==pv, type==2)

  pv.size <- mean(table.size$TS > qchisq(.05, nu, lower.tail=F))
  pv.power <- mean(table.power$TS > qchisq(.05, nu, lower.tail=F))
  large.summary <- rbind(large.summary, c(n1, n2, pv, fixed.g, pv.size, pv.power))
}
names(large.summary) <- c("n1", "n2", "p", "g", "size", "power")
```

Here we take a look at the table.

```
large.summary


##      n1 n2  p  g  size power
```

```
## 1  25 25  5 2 0.051 0.701
## 2  25 25  6 2 0.057 0.706
## 3  25 25  7 2 0.058 0.668
## 4  25 25  8 2 0.055 0.711
## 5  25 25  9 2 0.058 0.701
## 6  25 25 10 2 0.061 0.681
## 7  25 25 11 2 0.078 0.664
## 8  25 25 12 2 0.065 0.631
## 9  25 25 13 2 0.072 0.638
## 10 25 25 14 2 0.089 0.597
## 11 25 25 15 2 0.093 0.588
```

To understand the performance of simulations, it is better to visualize the size and power changes as a line plot.

Better with a visualization of the line plot.

```
ggplot(large.summary, aes(x=p)) +
  geom_line(aes(y=size,color="size")) +
  geom_line(aes(y=power,color="power")) +
  scale_colour_manual(values=c("#C83E45", "#5289B1")) +
  theme_minimal()
```



One thing noticable is that, since we have set  $n_1 = n_2 = 20$  which are definitely not large enough to be called “infinitely large”, our number of parameters  $p$  soon cataches up with  $n_1$  and  $n_2$ .

But how badly this performs? The power has a weak start with a value around 0.7 and gradually decreases to a more unpleasant range. On the other hand, the size is generally stable but not always significant ( $< .05$ ). Overall, we would suggest that the Box’s  $\chi^2$  approximation performs poorly when  $p > 5$  while holding  $g$  fixed below 5.



## 1.4 Q4 Testing Covariance Matrices WITHOUT adjustment

Besides the required dimension  $p = 5, 10, 50, 100, 300$ , we add an additional 375 for comparison.

```
n <- 500
p.vec <- c(5, 10, 50, 100, 300, 375)
N <- n - 1
sim <- 1000

table <- data.frame()

timer <- Sys.time()

for (i in 1:length(p.vec)) {
  p <- p.vec[i]
  mu <- rep(0, p)
  Sigma <- diag(p)
  Sigma2 <- diag(.05, p)
  Sigma2[1,1] <- 1

  for (s in 1:sim) {
    X <- rmvnorm(n, mu, Sigma)
    S <- cov(X)

    rho <- 1 - (2*p**2 + 3*p - 1) / (6*(n-1)*(p+1))
    ts <- -2*rho*(0.5*p*N + (0.5*N) *
               (unlist(determinant(S, logarithm=T))[[1]] - sum(diag(S))))

    X2 <- rmvnorm(n, mu, Sigma2)
    S2 <- cov(X2)
    ts2 <- -2*rho*(0.5*p*N + (0.5*N) *
              (unlist(determinant(S2, logarithm=T))[[1]] - sum(diag(S2))))

    table <- rbind(table, c(p, ts, 1))
    table <- rbind(table, c(p, ts2, 2))
  }
}

Sys.time() - timer
```

## Time difference of 12.51302 mins

```
names(table) <- c("p", "ts", "type")

large.summary <- data.frame()
for (pv in p.vec) {
  table.size <- table %>%
    filter(p==pv, type==1)
  table.power <- table %>%
    filter(p==pv, type==2)

  f <- 0.5*pv*(pv+1)
  gamma2 <- pv*(2*pv**4 + 6*pv**3 + pv**2 - 12*pv - 13) / (288*(pv+1))
  rho <- 1 - (2*pv**2 + 3*pv - 1) / (6*(n-1)*(pv+1))

  pv.size <- mean(table.size$ts > (
    qchisq(.95, f) + gamma2/rho**2/N**2*(qchisq(.95, f+4) - qchisq(.95, f))))
  pv.power <- mean(table.power$ts > (
    qchisq(.95, f) + gamma2/rho**2/N**2*(qchisq(.95, f+4) - qchisq(.95, f))))
```

```

large.summary <- rbind(large.summary, c(pv, pv.size, pv.power))
}
names(large.summary) <- c("p", "size", "power")

```

Summary:

```
large.summary
```

```

##      p  size power
## 1    5 0.060     1
## 2   10 0.050     1
## 3   50 0.044     1
## 4  100 0.055     1
## 5  300 0.243     1
## 6 375 0.964     1


```

Viz:

```

ggplot(large.summary, aes(x=p)) +
  geom_line(aes(y=size,color="size")) +
  geom_line(aes(y=power,color="power")) +
  scale_colour_manual(values=c("#C83E45", "#5289B1")) +
  theme_minimal()

```



The power is good. And the size is good at the beginning, but when  $n$  gets really large, size tends to 1.

Note that our sample size is 500 which is not a small number, so when dimension  $p$  gets large, modern RMT indicates that the likelihood ratio statistic drifts to infinity almost surely. Therefore, the classical  $\chi^2$  approximation may lead us to many false rejections of  $H_0$  in case of high-dimensional data. (TODO: add footnotes)

add Bai Jiang Yao Zheng 2009

## 1.5 Q5 A Proof

TODO:

**Proof:**

$$\begin{aligned}
 T_1 &= \text{tr} \mathbb{S} - \log |\mathbb{S}| - p \\
 &= p \cdot F^{\mathbb{S}}(f) \\
 &= p \cdot \int (x - \log x - 1) dF^{y_N}(x) \\
 &= p \cdot \int f(x) d(F^{\mathbb{S}}(x) - F^{y_N}(x)) + p \cdot F^{y_N}(f)
 \end{aligned}$$

Since  $N = n - 1, y_N = p/N$ :

$$\begin{aligned}
 F^{y_N}(f) &= 1 - \frac{y_N - 1}{y_N} \log(1 - y_N) \\
 &\doteq 1 - \frac{y - 1}{y} \log(1 - y) \\
 &= 1 + \frac{1 - y}{y} \log(y - 1) \\
 &= d_1(y)
 \end{aligned}$$

By Theorem [lecture](#) in notes,

$$T_1 - p \cdot F^{y_N}(f) \doteq T_1 - p \cdot d_1(y_N)$$

should weakly converges to a Gaussian vector with the mean

$$\mu_1 = -\frac{1}{2} \log(1 - y),$$

and variance

$$\sigma_1^2 = -2 \log(1 - y) - 2y.$$

## 1.6 Q6

This question is a refined version of Q4, an improvement has been made based on Q5 theoretically.

Note that we include a case with  $p = 375$  for consistency as well.

```
n <- 500
p.vec <- c(5, 10, 50, 100, 300, 375)
N <- n - 1
sim <- 1000

table <- data.frame()

timer <- Sys.time()

for (i in 1:length(p.vec)) {
  p <- p.vec[i]
  mu <- rep(0, p)
  Sigma <- diag(p)
  Sigma2 <- diag(.05,p)
  Sigma2[1,1] <- 1

  for (s in 1:sim) {
    y <- p/n
    YN <- p/N
    d1 <- 1+(1-YN)/YN*log(1-YN)

    X <- rmvnorm(n, mu, Sigma)
    S <- cov(X)
    T1 <- sum(diag(S))-unlist(determinant(S,logarithm=T))[[1]]-p

    ts <- T1-p*d1

    X2 <- rmvnorm(n, mu, Sigma2)
    S2 <- cov(X2)
    T1.2 <- sum(diag(S2))-unlist(determinant(S2,logarithm=T))[[1]]-p

    ts.2 <- T1.2-p*d1

    table <- rbind(table,c(p,ts,1))
    table <- rbind(table,c(p,ts.2,2))
  }
}

Sys.time()-timer

## Time difference of 12.62086 mins

names(table) <- c("p","ts","type")

large.summary <- data.frame()

for (pv in p.vec) {
  y <- pv/n
  mu1 <- -0.5*log(1-y)
  sigma1 <- sqrt(-2*log(1-y)-2*y)

  table.size <- table %>%
    filter(p==pv,type==1)
  table.power <- table %>%
    filter(p==pv,type==2)
```

```

pv.size <- mean(table.size$ts>qnorm(.95,mu1,sigma1))
pv.power <- mean(table.power$ts>qnorm(.95,mu1,sigma1))
large.summary <- rbind(large.summary, c(pv, pv.size, pv.power))
}
names(large.summary) <- c("p", "size", "power")

```

Check out the table again:

```
large.summary
```

```


##      p  size power
## 1    5 0.084     1
## 2   10 0.068     1
## 3   50 0.056     1
## 4  100 0.057     1
## 5  300 0.046     1
## 6 375 0.047     1

```

```

ggplot(large.summary, aes(x=p)) +
  geom_line(aes(y=size,color="size")) +
  geom_line(aes(y=power,color="power")) +
  scale_colour_manual(values=c("#C83E45", "#5289B1")) +
  theme_minimal()

```



This time, the size converges to 0.05 as expected, which in comparison, is of course an improvement from Q4.

## 2. Part 2 Multivariate Time Series

### Q7 MTS Basics

2.1

#### 7.(a)(b)(c)(d) MTS Generation and Visualization

2.1.1

For this question, my multivariate first order VAR/VMA model generating function might seem a little bit complex due to the introduction of multiple “flag” arguments. To be specific,

- `viz` controls if a time series plot is needed in the output.
- `lag` controls the upper limit of  $\tau$ .
- `add.S0` controls if a lag-0 autocovariance matrix is added to the “correlation matrix list”, since it is also required for the portmanteau test.
- `burnin` is the number of discarded observations. By default, we include the first 100 observations as burn-ins. For example, if the required observation number is 300, we generate 300+100 such observations and discard the first 100 accordingly.
- `type` is set to be `lag` by default for part (a), (b), (c) and (d). In this case, we are storing observations into the data frame `df`. But if it is set to be `qm`, i.e. when we conduct a portmanteau test, we are operating on residuals of observations, so that the data frame `df` will be assigned with residuals instead.

```
mts.gen.1 <- function(matrix.A, matrix.Sig, obs, method, viz=T, lag=0,
                      add.S0=F, burnin=100, type="lag") {
  df <- data.frame()
  if (method=="var") {
    epsilons <- rmvnorm((obs+burnin), c(0,0), matrix.Sig)
    Xt <- c(0,0)
    for (t in 1:(obs+burnin)) {
      Xt <- matrix.A%*%Xt + epsilons[t,]
      df <- rbind(df, c(Xt))
    }
    df <- df[(1+burnin):(obs+burnin),]
  } else if (method=="vma" || method=="ma") {
    epsilons <- rmvnorm(obs+burnin+1,c(0,0),matrix.Sig)
    for (t in 1:(obs+burnin)) {
      Xt <- epsilons[t+1,] + matrix.A%*%epsilons[t,]
      df <- rbind(df, c(Xt))
    }
    df <- df[(1+burnin):(obs+burnin),]
  }
  names(df) <- c("x1", "x2")
  df$id <- (1+burnin):(obs+burnin)
  df2 <- melt(df, id.var = "id", variable.name = "x")
  if (viz) {
    print(ggplot(df2,aes(x=id,y=value)) +
          geom_line(aes(color=x)) +
          facet_grid(x ~ .) +
          scale_colour_manual(values=c("#C83E45", "#5289B1")) +
          theme_minimal())
  }
  if (lag>0) {
    if (type=="qm") {
      # if we conduct a portmanteau test on residuals,
      # reassign the residual df to current variable df
      mod <- VAR(df,p=1,lag.max=20)
      df <- as.data.frame(cbind(mod$varresult$x1$residuals,mod$varresult$x2$residuals))
      names(df) <- c("x1", "x2")
    }
  }
}
```

```

rho.tau <- list()
X.bar <- c(mean(df$x1), mean(df$x2))
S0 <- cov(df[,1:2])
D <- diag(1/sqrt(diag(S0)))

# also record lag-0 autocovariance matrix, need to be used when doing portmanteau test
if (add.S0) {
  rho.tau[[lag+1]] <- cov2cor(S0) # the same as D %*% S0 %*% D
}


for (tau in 1:lag) {
  Stau <- 0
  for (t in (tau+1):nrow(df)) {
    Stau <- Stau + (c(df$x1[t], df$x2[t]) - X.bar) %*%
      t(c(df$x1[t-tau], df$x2[t-tau]) - X.bar)
  }
  Stau <- Stau / (nrow(df)-1)
  rho.tau[[tau]] <- D %*% Stau %*% D
}
return(rho.tau)
}

## (a) & (b) & (c) & (d)
A <- matrix(c(0.8,0.4,-0.3,0.6),byrow = T,ncol=2)
Sigma <- matrix(c(2,0.5,0.5,1),byrow = T,ncol=2)

```

Part (a) and (b) with time series visualization and its lag-5 correlation matrices.


```
mts.gen.1(A, Sigma, 300, "var", viz=T, lag=5, add.S0=F, burnin=100, type="lag")
```



```
## [[1]]
##           [,1]      [,2]
## [1,]  0.7859249 0.1526886
## [2,] -0.4931342 0.6202365
##
## [[2]]
##           [,1]      [,2]
## [1,]  0.5184852 0.2744502
## [2,] -0.6265335 0.3241870
##
## [[3]]
##           [,1]      [,2]
## [1,]  0.2670795 0.3394800
## [2,] -0.6097292 0.1038583
##
## [[4]]
##           [,1]      [,2]
## [1,]  0.02212521 0.34494099
## [2,] -0.51048627 -0.06452215
##
## [[5]]
##           [,1]      [,2]
## [1,] -0.1600059  0.2918388
## [2,] -0.3385234 -0.2013264
```

Part (c) and (d) with time series visualization and its lag-2 correlation matrices.

```
mts.gen.1(A, Sigma, 300, "vma", viz=T, lag=2, add.S0=F, burnin=100, type="lag")
```



```
## [[1]]
##           [,1]      [,2]
## [1,]  0.4739894 0.2453092
```



```
## [2,] -0.2099934 0.3844573
##
## [[2]]
##           [,1]      [,2]
## [1,] -0.07149522 -0.09460774
## [2,] -0.04072408  0.01342928
```

## 2.1.2 7.(e) Q7(e) Replication add reference link.

In this question, all parameters are set to be the values used in the paper (TODO)

```
A <- matrix(c(-.2, .3, -.6, 1.1), byrow = T, ncol=2)
B <- matrix(c(.4, .1, -1, .5), byrow = T, ncol=2)
C <- matrix(c(-1.5, 1.2, -.9, .5), byrow = T, ncol=2)
k <- 2 # 2 by 2 matrices
m <- 20
n <- 200
sim <- 1000

alpha.vec <- c(.25, -.25, .5, -.5, .75, -.75)

Delta.matrix <- function(alpha.value) {
  return(matrix(c(1, alpha.value, alpha.value, 1), byrow = T, ncol=2))
}

# kronecker(), vec()

inside.Qm <- function(rho1,rho0) {
  temp <- t(vec(t(rho1))) %*% kronecker(solve(rho0), solve(rho0)) %*% vec(t(rho1))
  return(unlist(temp)[1])
}

summary.table <- data.frame()

timer <- Sys.time()

for (i in 1:sim) {
  for (alpha in alpha.vec) {
    delta.matrix <- Delta.matrix(alpha)
    rho1 <- mts.gen.1(A, delta.matrix, n, "var", viz=F, lag=m, add.S0=T, type="qm")
    rho2 <- mts.gen.1(B, delta.matrix, n, "var", viz=F, lag=m, add.S0=T, type="qm")
    rho3 <- mts.gen.1(C, delta.matrix, n, "var", viz=F, lag=m, add.S0=T, type="qm")
    QA <- 0
    QB <- 0
    QC <- 0
    for (j in 1:m) {
      QA <- QA + inside.Qm(rho1[[j]], rho1[[m+1]])
      QB <- QB + inside.Qm(rho2[[j]], rho2[[m+1]])
      QC <- QC + inside.Qm(rho3[[j]], rho3[[m+1]])
    }
    QA <- n * QA
    QAs <- QA + k**2*m*(m+1)/(2*n)
    QB <- n * QB
    QBs <- QB + k**2*m*(m+1)/(2*n)
    QC <- n * QC
    QCs <- QC + k**2*m*(m+1)/(2*n)
    summary.table <- rbind(summary.table,
```

```

      c(alpha, QA, QAs, QB, QBs, QC, QCs))
    }
  }
  Sys.time() - timer

## Time difference of 35.70703 mins

names(summary.table) <- c("alpha", "QA", "QAs", "QB", "QBs", "QC", "QCs")

critical <- qchisq(.05,76,lower.tail = F) #  $k^2(m-p-q) = 4 * (20 - 1 - 0) = 76$ 

summary.table %>%
  mutate(QA = QA>critical,
         QAs = QAs>critical,
         QB = QB>critical,
         QBs = QBs>critical,
         QC = QC>critical,
         QCs = QCs>critical) %>%
  group_by(alpha) %>%
  summarize(
    A.Q20 = sum(QA),
    A.Q20.star = sum(QAs),
    B.Q20 = sum(QB),
    B.Q20.star = sum(QBs),
    C.Q20 = sum(QC),
    C.Q20.star = sum(QCs)
  )

```

```

## # A tibble: 6 x 7
##   alpha A.Q20 A.Q20.star B.Q20 B.Q20.star C.Q20 C.Q20.star
##   <dbl> <int>    <int> <int>    <int> <int>    <int>
## 1 -0.75     30         47     23         41     29         49
## 2 -0.5      26         59     30         54     28         58
## 3 -0.25     26         45     26         54     28         49
## 4  0.25     28         58     19         41     27         49
## 5  0.5      31         52     21         42     33         55
## 6  0.75     23         53     23         46     32         60

```

NOTE 1: The original table in paper (TODO: 1981) might have a serious typo, where “per cent” should actually be “count”. Consider a test with typo 1 error about 50%. Since the simulation run has a value of 1000 instead of 100, I guess this in fact is the “count”, which also explains why  $Q_m^*$  is better in the perspective of “closer to 0.05”.

NOTE 2: The difference between our simulation and the original table might originate from the fact that we discard 100 observations in the implemented function as burnin, while the paper didn’t specify their approach or their particular parameter.

**Note 3: 30 minutes runtime, still have space for optimization.**

### 2.1.3 7.(f) Q7(f) An Experiment on Federal Debts Data.

The two authors of package `portes` are... also... function oo good.

```

debt <- read.table("q-fdebt.txt",header = T) %>%      transformation, just
  dplyr::select(hbfin, hbfrbn, hbpun) %>%            taking the difference
  mutate(hbfin = log(hbfin), hbfrbn = log(hbfrbn), hbpun = log(hbpun))
debt <- as.matrix(debt)
DiffData <- matrix(numeric(ncol(debt) * (nrow(debt)-1)), ncol = ncol(debt))
for (i in 1:ncol(debt)) DiffData[, i] <- diff(debt[, i], lag = 1)
Fit <- ar.ols(DiffData, aic=F, order.max = 1)

# theoretical                                insert there are two versions,
LiMcLeod(Fit, lags=1:10)                    theoretical with... or monte carlo
                                           simulation with portes

```

```
## lags statistic df p-value
## 1 7.161042 0 0.000000e+00
## 2 18.339335 9 3.143462e-02
## 3 36.555260 18 5.985399e-03
## 4 111.935815 27 2.609801e-12
## 5 137.487311 36 8.859580e-14
## 6 144.074991 45 2.663314e-12
## 7 153.214982 54 1.954914e-11
## 8 201.506096 63 2.220446e-16
## 9 213.347314 72 6.661338e-16
## 10 220.915081 81 6.550316e-15
```

*# Monte Carlo version*

```
portest(Fit, lags=1:10, test="LiMcLeod", ncores=4)
```

```
## lags statistic p-value
## 1 7.161042 0.003996004
## 2 18.339335 0.038961039
## 3 36.555260 0.003996004
## 4 111.935815 0.000999001
## 5 137.487311 0.000999001
## 6 144.074991 0.000999001
## 7 153.214982 0.000999001
## 8 201.506096 0.000999001
## 9 213.347314 0.000999001
## 10 220.915081 0.000999001
```

## 2.2 Q8 VAR model order selection on CPI data


aka  
plotting  
CPI  
data

2.2.1 (a) unlike Q7 part f, here we take more complex transformation on CPI index, so that the CPI rate is  
=====

```
dat <- read.table("m-cpi tb3m.txt", header = T) %>%
  dplyr::select(tb=tb3m, cpi=cpi aucs1) %>%
  mutate(cpi=100*log(cpi))

dat <- diff(ts(dat, frequency = 12, start = c(1947, 1)))

autoplot(dat) +
  scale_colour_manual(values=c("#C83E45", "#5289B1")) +
  theme_minimal()
```



2.2.2 (b) The selecti on algorithm inconsistency in notation

```
# library(MTS)
# MTS::VARorder(dat, maxp=10)
```

modified the VARorder function in MTS library. Actually I simplified it by removing the order selection criteria like AIC, BIC, only keeping the Barlet test statistic and calculated p-value. And implement my own version of m selection.

```
VARorder2 <- function (x, P = 10) {
  x1 <- as.matrix(x)
  T <- nrow(x1)
  k <- ncol(x1)
  if (P < 1) {P = 1}
  obs = T - P
  y = x1[(P + 1):T, , drop = FALSE]
  ist = P + 1
  xmtx = cbind(rep(1, obs), x1[P:(T - 1), ])
```

```

if (P > 1) {
  for (i in 2:P) {
    xmtx = cbind(xmtx, x1[(ist - i):(T - i), ])
  }
}
chidet = rep(0, (P + 1))
s = cov(y) * (obs - 1)/obs
chidet[1] = log(det(s))
y = as.matrix(y)
for (l in 1:P) {
  idm = k * l + 1
  xm = xmtx[, 1:idm]
  xm = as.matrix(xm)
  xpx <- crossprod(xm, xm)
  xpy <- crossprod(xm, y)
  beta <- solve(xpx, xpy)
  yhat <- xm %*% beta
  resi <- y - yhat
  sse <- crossprod(resi, resi)/obs
  d1 = log(det(sse))
  chidet[l + 1] = d1
}
Mstat = rep(0, P)
pv = rep(0, P)
for (j in 1:P) {
  Mstat[j] = (T - P - k * j - 1.5) * (chidet[j] - chidet[j + 1])
  pv[j] = 1 - pchisq(Mstat[j], k ** 2)
}

output = cbind(c(0:P), c(0, Mstat), c(0, pv))
colnames(output) <- c("l", "M(l)", "p-value")

# which m? i.e. starting from this l=m, the p-value of following models
# should all be greater than 0.05

m.output <- 1
for (m in 1:nrow(output)) {
  if (all(output[m:nrow(output),3] >= 0.05)) {
    m.output <- m
    break()
  }
}
print(round(output, 3))
cat("Select VAR (", m.output-1, ")")
}

insert here

VARorder2(dat, 50)

```

```

##      1      M(1) p-value
## [1,] 0      0.000 0.000
## [2,] 1 441.733 0.000
## [3,] 2  66.496 0.000
## [4,] 3   9.399 0.052
## [5,] 4  25.067 0.000
## [6,] 5  19.022 0.001
## [7,] 6  55.022 0.000
## [8,] 7  19.353 0.001

```

For this part of the question, we only set P=50

```
## [9,] 8 21.604 0.000
## [10,] 9 23.523 0.000
## [11,] 10 15.876 0.003
## [12,] 11 5.698 0.223
## [13,] 12 23.276 0.000
## [14,] 13 13.020 0.011
## [15,] 14 8.099 0.088
## [16,] 15 19.250 0.001
## [17,] 16 13.831 0.008
## [18,] 17 5.467 0.243
## [19,] 18 2.474 0.649
## [20,] 19 13.430 0.009
## [21,] 20 11.943 0.018
## [22,] 21 2.940 0.568
## [23,] 22 1.783 0.776
## [24,] 23 3.090 0.543
## [25,] 24 12.153 0.016
## [26,] 25 8.388 0.078
## [27,] 26 4.946 0.293
## [28,] 27 1.894 0.755
## [29,] 28 1.726 0.786
## [30,] 29 8.431 0.077
## [31,] 30 0.724 0.948
## [32,] 31 10.173 0.038
## [33,] 32 1.822 0.768
## [34,] 33 0.308 0.989
## [35,] 34 0.291 0.990
## [36,] 35 5.653 0.227
## [37,] 36 2.444 0.655
## [38,] 37 0.444 0.979
## [39,] 38 0.866 0.929
## [40,] 39 6.145 0.189
## [41,] 40 6.934 0.139
## [42,] 41 0.881 0.927
## [43,] 42 0.892 0.926
## [44,] 43 4.728 0.316
## [45,] 44 0.774 0.942
## [46,] 45 8.692 0.069
## [47,] 46 2.158 0.707
## [48,] 47 2.106 0.716
## [49,] 48 3.445 0.486
## [50,] 49 3.054 0.549
## [51,] 50 2.979 0.561
## Select VAR ( 32 )
```

the curse of dimensionality

(c)

THE INTUITION OF THIS PART OF  
THE QUESTION IS

```
library(tsdyn)
increase.dim <- function(data, p) {
  dat.p <- dat
  # dim 3 to p
  for (dim.p in 3:p) {
    dat.p <- cbind(dat.p, sample(dat[,sample(1:2, 1, prob=c(0.5,0.5))],
                                nrow(dat), replace=T))
  }
  dim(dat.p)
```

didn't run the algorithm with a bunch of different P as a simulation.

I totally understand the result model selection highly depends on the P you select. But seriously, for small P, say we have P around 10, the selected  $m$  is pretty much close to P. But if P becomes large enough,  $m$  would tend to a stable value.

```

mod <- lineVar(dat.p, lag=1)
new.mod <- VAR.boot(mod, "resample")
VARorder2(new.mod, P=20)
}

```

```
increase.dim(dat, 10)
```

```

##          1      M(1) p-value
## [1,]  0      0.000  0.000
## [2,]  1 608.674  0.000
## [3,]  2  89.173  0.773
## [4,]  3  98.026  0.537
## [5,]  4  80.760  0.921
## [6,]  5  78.809  0.942
## [7,]  6  93.331  0.668
## [8,]  7  89.448  0.766
## [9,]  8 110.501  0.222
## [10,] 9 109.152  0.250
## [11,] 10 89.639  0.762
## [12,] 11 105.577  0.332
## [13,] 12  89.135  0.773
## [14,] 13 124.502  0.049
## [15,] 14 124.377  0.050
## [16,] 15 109.658  0.239
## [17,] 16 101.651  0.435
## [18,] 17  92.212  0.698
## [19,] 18  89.250  0.771
## [20,] 19 114.667  0.150
## [21,] 20  86.146  0.837
## Select VAR ( 15 )

```

```
increase.dim(dat, 20)
```

```

##          1      M(1) p-value
## [1,]  0      0.000  0.000
## [2,]  1 1149.768  0.000
## [3,]  2  414.855  0.294
## [4,]  3  371.294  0.845
## [5,]  4  436.758  0.099
## [6,]  5  399.094  0.503
## [7,]  6  398.796  0.508
## [8,]  7  377.260  0.787
## [9,]  8  422.052  0.215
## [10,] 9  412.137  0.327
## [11,] 10  377.474  0.785
## [12,] 11  393.017  0.589
## [13,] 12  390.518  0.623
## [14,] 13  406.624  0.399
## [15,] 14  448.810  0.046
## [16,] 15  443.344  0.066
## [17,] 16  389.596  0.636
## [18,] 17  425.331  0.184
## [19,] 18  433.851  0.117
## [20,] 19  385.575  0.689
## [21,] 20  385.094  0.695
## Select VAR ( 15 )

```

```
increase.dim(dat, 30)
```

```
##          1      M(1) p-value
## [1,] 0      0.000 0.000
## [2,] 1 2117.922 0.000
## [3,] 2  888.507 0.601
## [4,] 3  993.378 0.016
## [5,] 4  818.134 0.976
## [6,] 5  872.855 0.736
## [7,] 6  881.030 0.668
## [8,] 7  944.988 0.145
## [9,] 8  813.550 0.982
## [10,] 9 1053.468 0.000
## [11,] 10  948.118 0.129
## [12,] 11  963.036 0.071
## [13,] 12  914.459 0.361
## [14,] 13  986.043 0.024
## [15,] 14  957.778 0.089
## [16,] 15  965.372 0.064
## [17,] 16  878.385 0.691
## [18,] 17  895.413 0.537
## [19,] 18  911.608 0.387
## [20,] 19  951.776 0.112
## [21,] 20  976.452 0.038
## Select VAR ( 0 )
```

```
increase.dim(dat, 40) # computationally singular
```

```
## Error in solve.default(xpx, xpy): system is computationally singular: reciprocal condition number = 2.206
```

```
increase.dim(dat, 50) # computationally singular
```

```
## Error in solve.default(xpx, xpy): system is computationally singular: reciprocal condition number = 1.531
```

Curse of dimensionality

reference



## 2.3 Q9

### how we separate data?

- plot 1
  - $p = 20, n = 40, \tau = 0$
  - $p = 50, n = 100, \tau = 0$
- plot 2
  - $p = 20, n = 20, \tau = 0$
  - $p = 50, n = 50, \tau = 0$
- plot 3
  - $p = 20, n = 40, \tau = 1$
  - $p = 50, n = 100, \tau = 1$
- plot 4
  - $p = 20, n = 20, \tau = 1$
  - $p = 50, n = 50, \tau = 1$
- plot 5
  - $p = 20, n = 40, \tau = 2$
  - $p = 50, n = 100, \tau = 2$
- plot 6
  - $p = 20, n = 20, \tau = 2$
  - $p = 50, n = 50, \tau = 2$

```
ps <- c(20,50)
ratios <- c(.5, 1)
taus <- c(0,1,2)

lag.sample.acf <- function(n, tau, X) {
  if (tau==0) {
    return(cov(X))
  } else {
    Ctau <- 0
    for (t in 1:(n-tau)) {
      Ctau <- Ctau + X[t,] %*% t(X[t+tau,]) + X[t+tau,] %*% t(X[t,])
    }
    Ctau <- Ctau/(2*n)
    return(Ctau)
  }
}

df <- data.frame()
for (ratio in ratios) {
  for (tau in taus){
    for (p in ps) {
      n <- p / ratio

      # draw the ESDs for iid
      Z <- rmvnorm(n,mean=rep(0,p),sigma=diag(p))
      S <- lag.sample.acf(n, tau, Z)
      ev <- eigen(S)$values

      F_ <- function(x) {
        total <- 0.
        for (i in 1:p) {
          if (ev[i] <= x){
            total <- total + 1
          }
        }
        return(total/p)
      }
    }
  }
}
```

```

empirical.cdf <- Vectorize(F_)

df <- rbind(df, data.frame(p=rep(p,n),
                           ratio=rep(ratio,n),
                           tau=rep(tau, n),
                           x=ev,
                           Fvalue=empirical.cdf(ev),
                           type=rep(paste("ESD.iid p =",p),n)))

# draw the ESDs for MA(1)
A1 <- diag(ncol(Z))
Z.copy <- Z
Z.copy <- cbind(rep(1,ncol(Z.copy)), Z.copy[,1:(ncol(Z.copy)-1)])
X <- Z+Z.copy
ev2 <- eigen(lag.sample.acf(n, tau, X))$values
df <- rbind(df, data.frame(p=rep(p,n),
                           ratio=rep(ratio,n),
                           tau=rep(tau, n),
                           x=ev2,
                           Fvalue=empirical.cdf(ev2),
                           type=rep(paste("ESD.MA(1) p =",p),n)))

# draw the LSD for iid
pp <- 500
nn <- pp / ratio
Zn <- rmvnorm(nn,mean=rep(0,pp),sigma=diag(pp))
Sn <- lag.sample.acf(nn,tau,Zn)
evn <- eigen(Sn)$values

df <- rbind(df, data.frame(p=rep(pp,length(evn)),
                           ratio=rep(ratio,length(evn)),
                           tau=rep(tau,length(evn)),
                           x=evn,
                           Fvalue=empirical.cdf(evn),
                           type=rep("LSD.iid",length(evn))))

# draw the LSD for MA(1)
A1n <- diag(ncol(Zn))
Zn.copy <- Zn
Zn.copy <- cbind(rep(1,ncol(Zn.copy)), Zn.copy[,1:(ncol(Zn.copy)-1)])
Xn <- Zn+Zn.copy
ev2n <- eigen(lag.sample.acf(length(evn), tau, Xn))$values
df <- rbind(df, data.frame(p=rep(pp,length(ev2n)),
                           ratio=rep(ratio,length(ev2n)),
                           tau=rep(tau, length(ev2n)),
                           x=ev2n,
                           Fvalue=empirical.cdf(ev2n),
                           type=rep("LSD.MA(1)",length(ev2n))))
}
}
}

# plotting

df$para <- paste("tau = ", df$tau, "p/n = ", df$ratio)
ggplot(df, aes(x=x, y=Fvalue, color=type)) +
  stat_ecdf(geom = "step", size=.5, alpha=.8) +

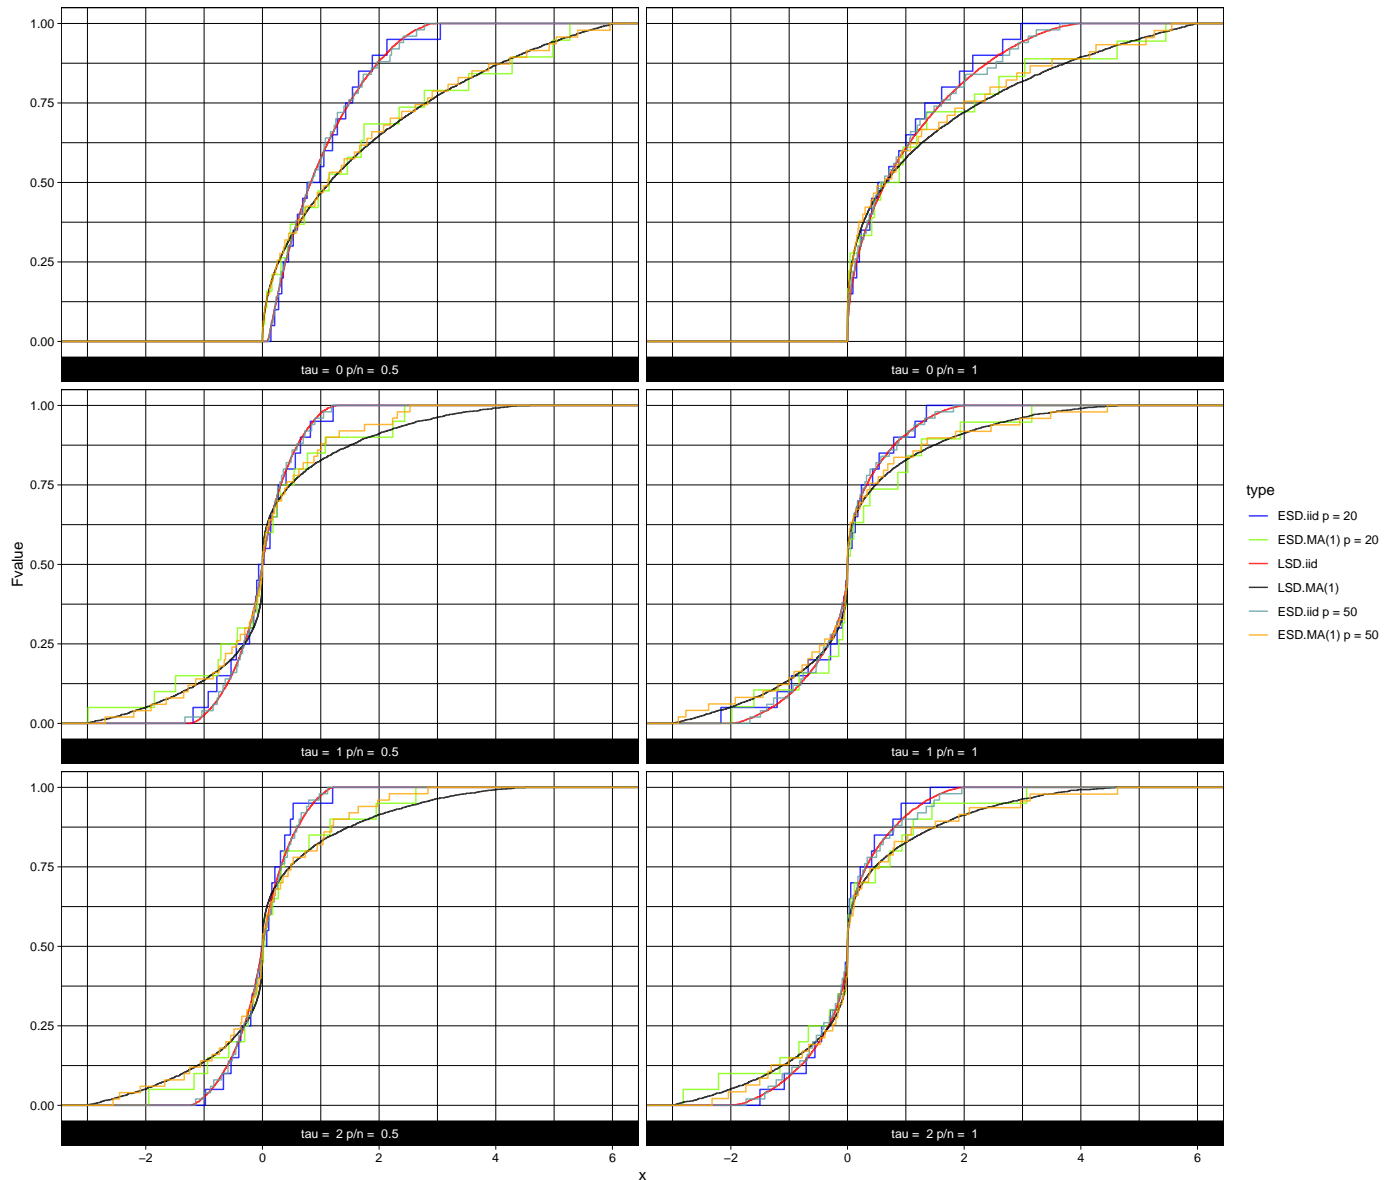
```

```

facet_wrap(para ~ ., nrow=3, strip.position = "bottom") +
scale_x_continuous(limits = c(-3,6)) +
scale_y_continuous(limits = c(0,1)) +
scale_color_manual(values=c("LSD.iid" = "red",
                             "LSD.MA(1)" = "black",
                             "ESD.iid p = 20" = "blue",
                             "ESD.MA(1) p = 20" = "chartreuse",
                             "ESD.iid p = 50" = "cadetblue",
                             "ESD.MA(1) p = 50" = "orange")) +

theme_linedraw()

```



We are pretty sure that we had some mistakes in the MA(1) model since they look different comparing with the ones in the paper (TODO:Li) but the three of iid models are pretty much the same.

## References

- Visualizing Tests for Equality of Covariance Matrices, Michael Friendly and Matthew Siga. <https://arxiv.org/pdf/1805.05756.pdf>
- <https://stats.stackexchange.com/questions/40868/experiment-or-simulation-to-understand-type-i-and-type-ii-errors/40874#40874>
- TODO: and many more with decent format of citations
- 1. all covariance matrix use `cov()` for better accuracy
- 2. all `log(det())` use `determinant()` with argument `log=TRUE`
- 3. for type 1 error and power calculation, use statistic vs critical value instead p-value vs significance.

portmanteau test

<https://faculty.chicagobooth.edu/ruey.tsay/teaching/mts/sp2009/lec1-09.pdf> <https://faculty.chicagobooth.edu/ruey.tsay/teaching/mts/sp2009/lec2-09.pdf>

- <https://rdrr.io/cran/MTS/man/VARorder.html>