

Workshop 1

- 1 Crash course
 - 1.1 Multidimensional data basics
 - 1.1.1 Constructing vectors
 - 1.1.2 Constructing matrices
 - 1.1.3 Basic operations
 - 1.1.4 Other operations
 - 1.2 Plotting multidimensional data
 - 1.2.1 Box plot
 - 1.2.2 Scatter plot
 - 1.2.3 Bubble plots
 - 1.2.4 Star plots
 - 1.2.5 Parallel coordinates plot
 - 1.3 Random vectors and matrices
 - 1.3.1 Random matrix
 - 1.3.2 Variance and standard deviation of a vector
 - 1.3.3 Sample variance-covariance
- 2 Salient features of Big Data
 - 2.1 Outliers in higher dimensions are not obvious
 - 2.1.1 Outliers in univariate case
 - 2.1.2 Multivariate outliers
 - 2.1.3 More dimensions
 - 2.2 Heterogeneity
 - 2.3 Noise accumulation

1 Crash course

This section is inspired by code from Everitt's book *An Introduction to Applied Multivariate Analysis with R*. This is potentially a good read if you like the applied side of multivariate analysis.

1.1 Multidimensional data basics

1.1.1 Constructing vectors

Create a vector, using the command `c()` (for concatenate)

```
x <- c(1,2,3,4)
x
```

```
## [1] 1 2 3 4
```

Sum all elements of x:

```
sum(x)
```

```
## [1] 10
```

Square all elements of x:

```
x^2
```

```
## [1] 1 4 9 16
```

Get the third element of x:

```
x[3]
```

```
## [1] 3
```

Add an extra element:

```
x <- c(x,10)
x
```

```
## [1] 1 2 3 4 10
```

Create a vector using the command `seq()` (for sequence)

```
x <- seq(from=1, to=10, by=.1)
x
```

```
## [1] 1.0 1.1 1.2 1.3 1.4 1.5 1.6 1.7 1.8 1.9 2.0 2.1 2.2 2.3
## [15] 2.4 2.5 2.6 2.7 2.8 2.9 3.0 3.1 3.2 3.3 3.4 3.5 3.6 3.7
## [29] 3.8 3.9 4.0 4.1 4.2 4.3 4.4 4.5 4.6 4.7 4.8 4.9 5.0 5.1
## [43] 5.2 5.3 5.4 5.5 5.6 5.7 5.8 5.9 6.0 6.1 6.2 6.3 6.4 6.5
## [57] 6.6 6.7 6.8 6.9 7.0 7.1 7.2 7.3 7.4 7.5 7.6 7.7 7.8 7.9
## [71] 8.0 8.1 8.2 8.3 8.4 8.5 8.6 8.7 8.8 8.9 9.0 9.1 9.2 9.3
## [85] 9.4 9.5 9.6 9.7 9.8 9.9 10.0
```

Obtain the length of x:

```
length(x)
```

```
## [1] 91
```

Create a vector using the command `'rep()'` (for repeat)

```
x <- rep(0,times=10)
x
```

```
## [1] 0 0 0 0 0 0 0 0 0 0
```

1.1.2 Constructing matrices

Library for the matrix commands.

```
library(Matrix)
```

Construct a matrix.

```
A<-matrix(c(1, 2, 3, 4, 5, 6), byrow=T, ncol=3)  
print(A)
```

```
##      [,1] [,2] [,3]  
## [1,]    1    2    3  
## [2,]    4    5    6
```

Access elements.

```
A[1,1]
```

```
## [1] 1
```

Access columns.

```
A[1,]
```

```
## [1] 1 2 3
```

Access rows.

```
A[,1]
```

```
## [1] 1 4
```

Construct by column first.

```
B<-matrix(c(1, 2, 3, 4, 5, 6), byrow=F, ncol=3)  
print(B)
```

```
##      [,1] [,2] [,3]  
## [1,]    1    3    5  
## [2,]    2    4    6
```

Construct a diagonal matrix.

```
D<-diag(c(1,2,3))
```

Construct an identity matrix.

```
I<-diag(c(1,1,1))
```

Construct a matrix of all ones.

```
ONES<-matrix(rep(1,9),ncol=3)
```

1.1.3 Basic operations

Create some vectors and matrices.

```
x<-c(1, 2, 3)
y<-c(4, 5, 6)
z<-seq(1,10,by=1)
```

To make sure that R respects dimensions, turn them into matrices.

```
x<-as.matrix(x)
y<-as.matrix(y)
```

1.1.3.1 Basic operations

Transpose operations.

```
t(A)
```

```
##      [,1] [,2]
## [1,]    1    4
## [2,]    2    5
## [3,]    3    6
```

```
t(B)
```

```
##      [,1] [,2]
## [1,]    1    2
## [2,]    3    4
## [3,]    5    6
```

```
t(D)
```

```
##      [,1] [,2] [,3]
## [1,]    1    0    0
## [2,]    0    2    0
## [3,]    0    0    3
```

Element-wise operations on matrices.

```
A+B
```

```
##      [,1] [,2] [,3]
## [1,]    2    5    8
```

```
## [2,]    6    9   12
```

A-B

```
##      [,1] [,2] [,3]
## [1,]    0  -1  -2
## [2,]    2   1   0
```

A*B

```
##      [,1] [,2] [,3]
## [1,]    1   6  15
## [2,]    8  20  36
```

A/B

```
##      [,1]      [,2] [,3]
## [1,]    1 0.6666667 0.6
## [2,]    2 1.2500000 1.0
```

A^B

```
##      [,1] [,2] [,3]
## [1,]    1   8 243
## [2,]   16 625 46656
```

Element-wise operations on vectors.

x+y

```
##      [,1]
## [1,]    5
## [2,]    7
## [3,]    9
```

x-y

```
##      [,1]
## [1,]   -3
## [2,]   -3
## [3,]   -3
```

x*y

```
##      [,1]
## [1,]    4
## [2,]   10
```

```
## [3,] 18
```

x/y

```
##      [,1]
## [1,] 0.25
## [2,] 0.40
## [3,] 0.50
```

y^x

```
##      [,1]
## [1,] 4
## [2,] 25
## [3,] 216
```

1.1.3.2 Matrix and vector operations

This would give an error message: non-conformable.

```
# A %*% B
```

Check the matrix dimension.

```
dim(A)
```

```
## [1] 2 3
```

```
dim(B)
```

```
## [1] 2 3
```

A correct calculation

```
A %*% t(B)
```

```
##      [,1] [,2]
## [1,] 22  28
## [2,] 49  64
```

or some alternatives

```
t(A) %*% B
```

```
##      [,1] [,2] [,3]
## [1,] 9   19  29
## [2,] 12  26  40
```

```
## [3,] 15 33 51
```

```
t(B) %*% A
```

```
##      [,1] [,2] [,3]
## [1,] 9 12 15
## [2,] 19 26 33
## [3,] 29 40 51
```

```
B %*% t(A)
```

```
##      [,1] [,2]
## [1,] 22 49
## [2,] 28 64
```

```
x %*% t(y)
```

```
##      [,1] [,2] [,3]
## [1,] 4 5 6
## [2,] 8 10 12
## [3,] 12 15 18
```

```
t(x) %*% y
```

```
##      [,1]
## [1,] 32
```

```
t(x) %*% t(A)
```

```
##      [,1] [,2]
## [1,] 14 32
```

Multiplies each column of B by a number

```
B %*% D
```

```
##      [,1] [,2] [,3]
## [1,] 1 6 15
## [2,] 2 8 18
```

Multiplies each row of B by a number

```
diag(c(3,4)) %*% B
```

```
##      [,1] [,2] [,3]
## [1,] 3 9 15
## [2,] 8 16 24
```

1.1.4 Other operations

Determinant of a matrix

```
det(D)
```

```
## [1] 6
```

```
det(ONES)
```

```
## [1] 0
```

Inverse of a matrix

```
Di<-solve(D)
```

```
D %*% Di
```

```
##      [,1] [,2] [,3]
## [1,]    1    0    0
## [2,]    0    1    0
## [3,]    0    0    1
```

```
Di %*% D
```

```
##      [,1] [,2] [,3]
## [1,]    1    0    0
## [2,]    0    1    0
## [3,]    0    0    1
```

You can create an almost-singular matrix $(I+N)$ by choosing small variance for the noise matrix N and see what happens with the inverse.

```
N<-matrix(rnorm(9,sd=10^-1),3,3)
```

```
Ii<-solve(ONES+N)
```

```
(ONES+N)%*%Ii
```

```
##      [,1]      [,2]      [,3]
## [1,] 1.000000e+00 1.970949e-17 8.881784e-16
## [2,] 4.957927e-16 1.000000e+00 0.000000e+00
## [3,] 0.000000e+00 -2.220446e-16 1.000000e+00
```

```
Ii%*%(I+N)
```

```
##      [,1]      [,2]      [,3]
## [1,] -4.910364 -3.256198 8.405292
## [2,] -1.549238 4.504265 -2.211232
## [3,] 6.238835 -1.057717 -5.201819
```

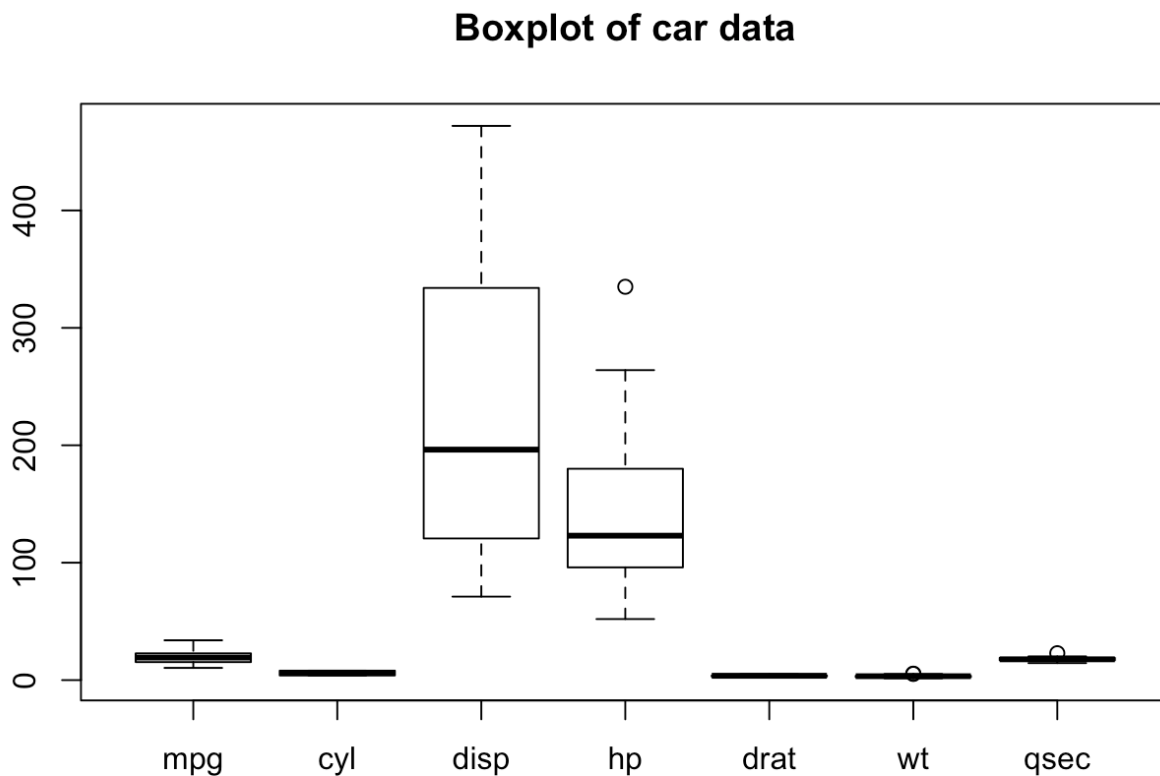

1.2 Plotting multidimensional data

Visualisation is very important when exploring a new data set. Here are some useful ways to look at multi-dimensional data.

1.2.1 Box plot

Standard boxplot.


```
boxplot(mtcars[,1:7], main="Boxplot of car data")
```



1.2.2 Scatter plot

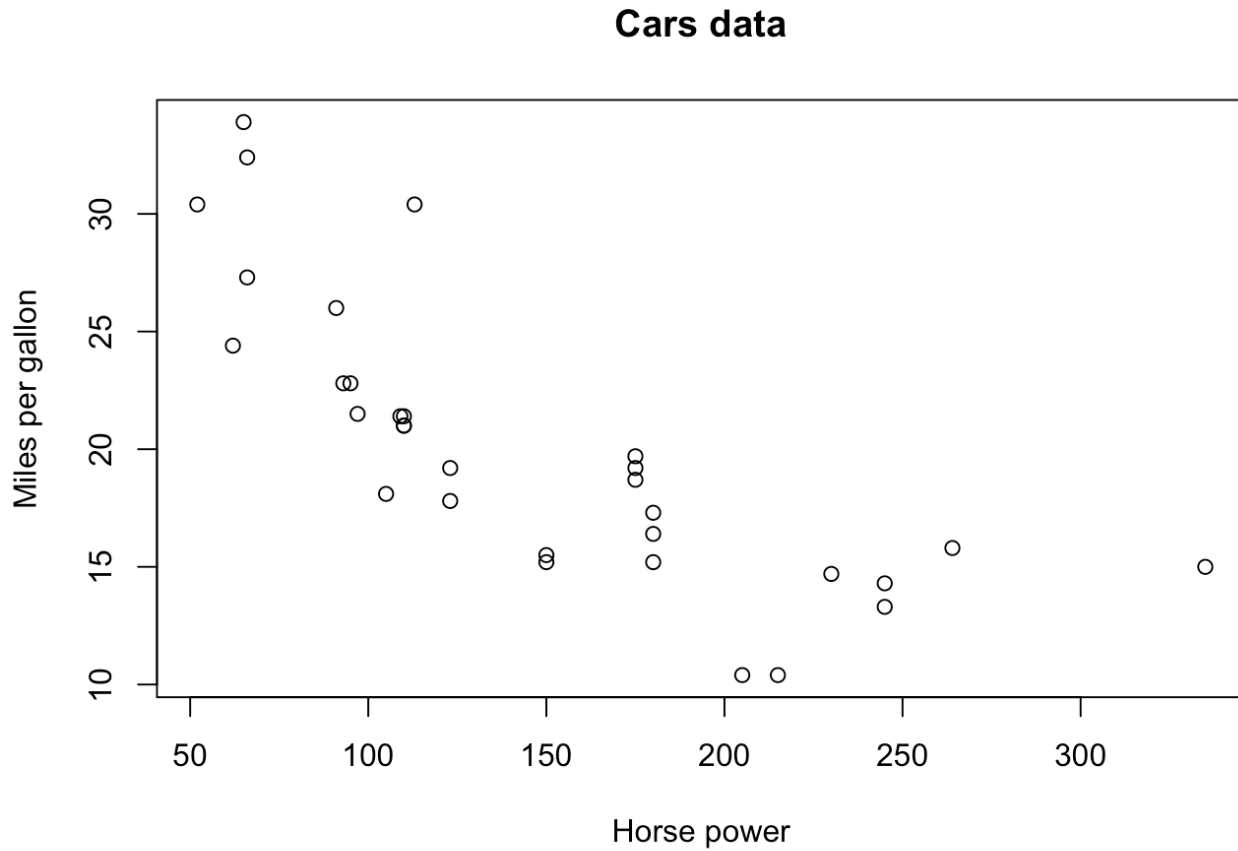
Bivariate scatter plot.

```
plot(mtcars$hp, mtcars$mpg)
```



Add labels along the axes

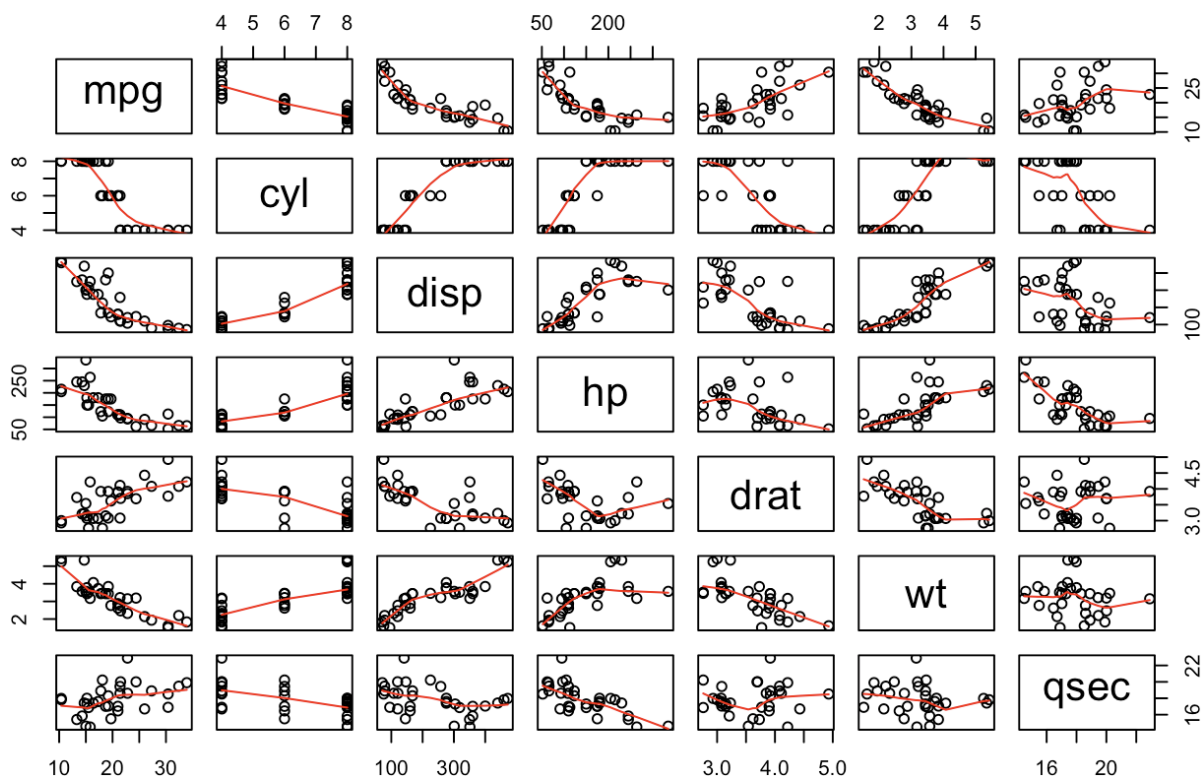
```
plot(mtcars$hp, mtcars$mpg, xlab="Horse power",  
      ylab="Miles per gallon", main="Cars data")
```



All possible bivariate scatter plots.

```
pairs(mtcars[,1:7], panel=panel.smooth,  
      main="Scatterplot matrix of car data")
```

Scatterplot matrix of car data



We can make a nicer version but we need to create two helper functions first. This one puts histograms on the diagonal.

```
panel.hist <- function(x, ...) {
  usr <- par("usr"); on.exit(par(usr))
  par(usr = c(usr[1:2], 0, 1.5) )
  h <- hist(x, plot = FALSE)
  breaks <- h$breaks; nB <- length(breaks)
  y <- h$counts; y <- y/max(y)
  rect(breaks[-nB], 0, breaks[-1], y,
      col = "cyan", ...)
}
```

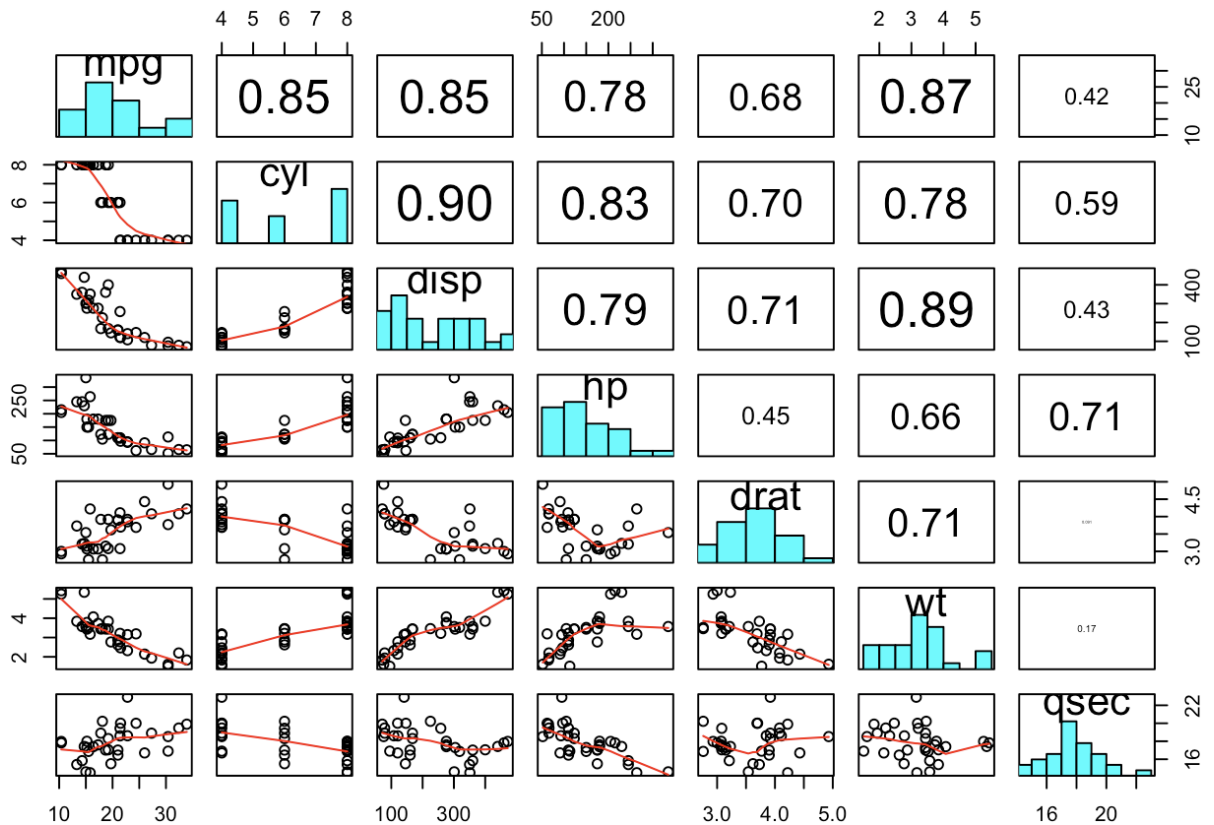
This one puts (absolute) correlations on the upper panels, with size proportional to the correlations.

```
panel.cor <- function(x, y, digits = 2, prefix = "", cex.cor, ...) {
  usr <- par("usr"); on.exit(par(usr))
  par(usr = c(0, 1, 0, 1))
  r <- abs(cor(x, y))
  txt <- format(c(r, 0.123456789), digits = digits)[1]
  txt <- paste0(prefix, txt)
  if(missing(cex.cor)) cex.cor <- 0.8/strwidth(txt)
```

```
text(0.5, 0.5, txt, cex = cex.cor * r)
}
```


Now let's do it.

```
pairs(mtcars[,1:7], panel=panel.smooth,
      upper.panel=panel.cor, diag.panel=panel.hist)
```



1.2.3 Bubble plots

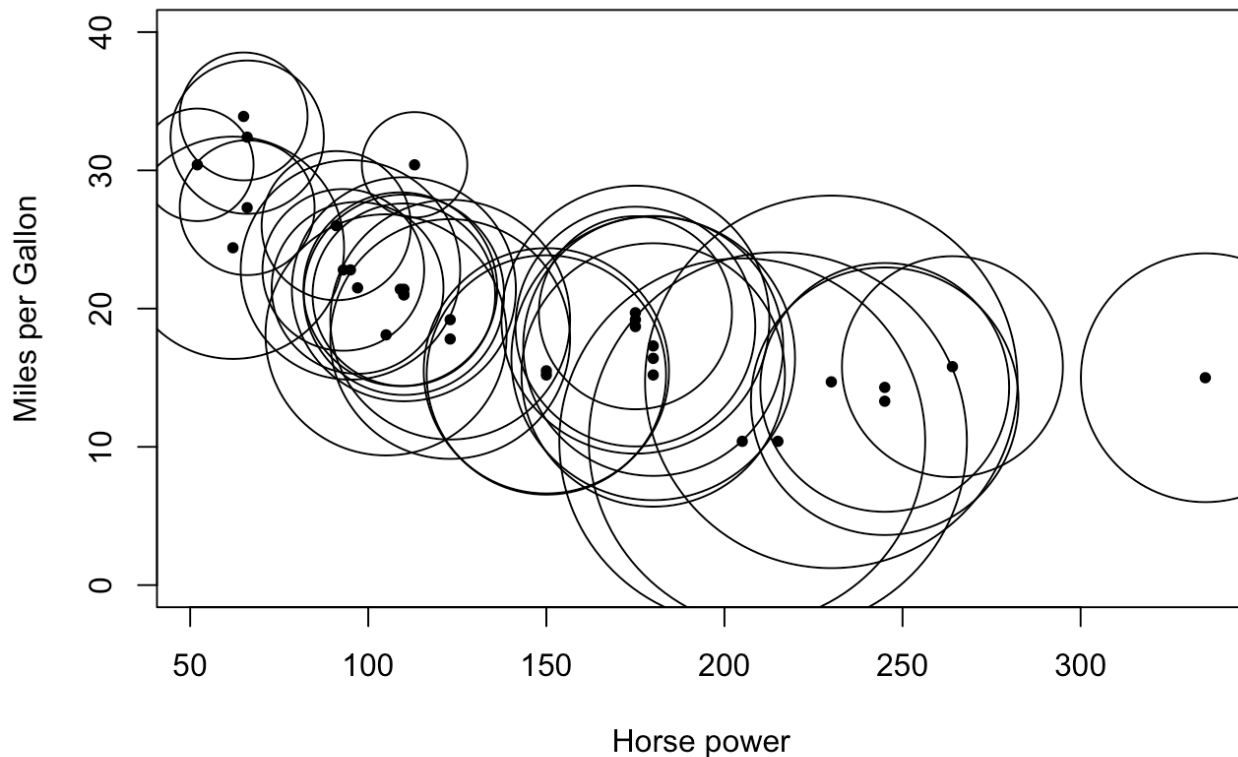
```
plot(mtcars$hp, mtcars$mpg, pch = 20,
      xlab = "Horse power", ylab = "Miles per Gallon",
      ylim=c(0,40))
```



We can add a third variable by using circles.

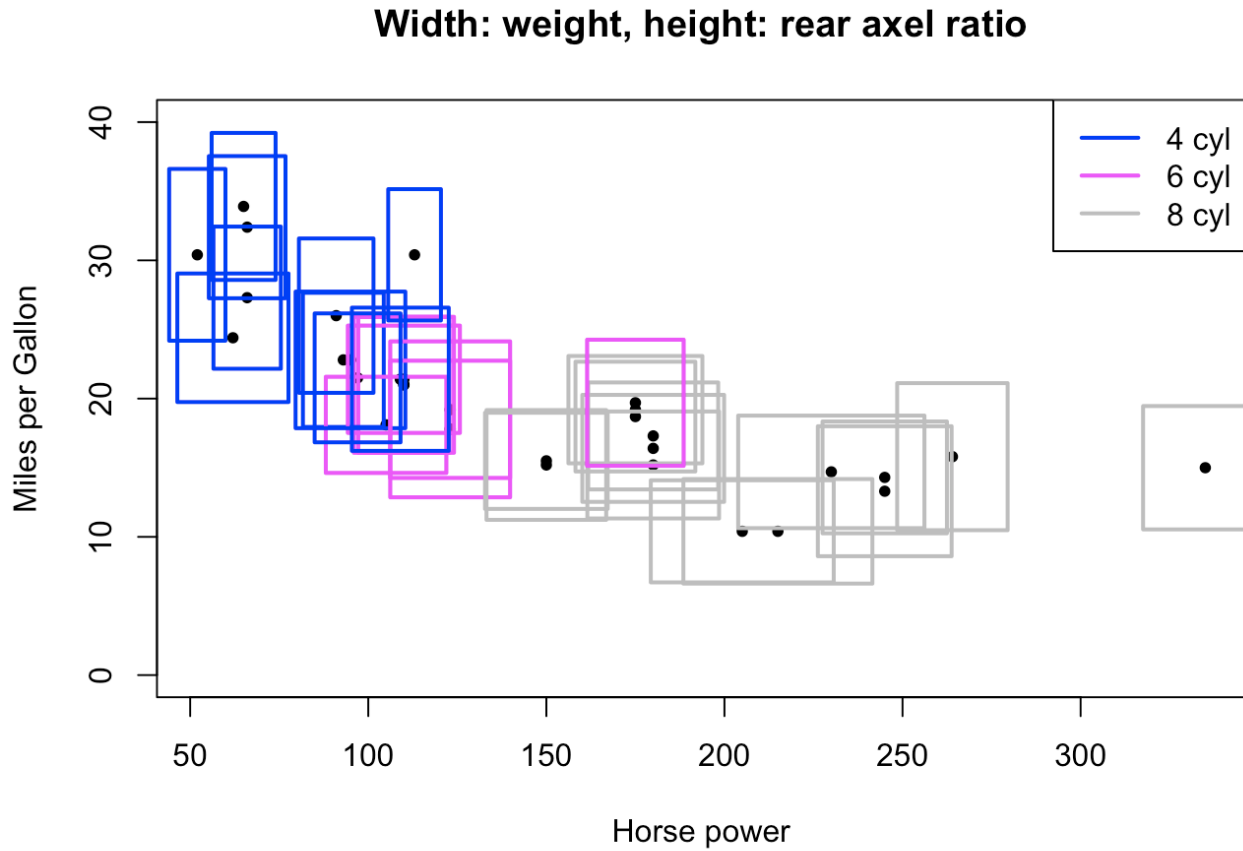
```
plot(mtcars$hp, mtcars$mpg, pch = 20,  
     xlab = "Horse power", ylab = "Miles per Gallon",  
     ylim=c(0,40))  
symbols(mtcars$hp, mtcars$mpg, circles = mtcars$wt,  
        add = TRUE)  
title("Bubble plot: The radius of the circle  
      indicates weight")
```

Bubble plot: The radius of the circle indicates weight



We can add even more variables with rectangles and colors:

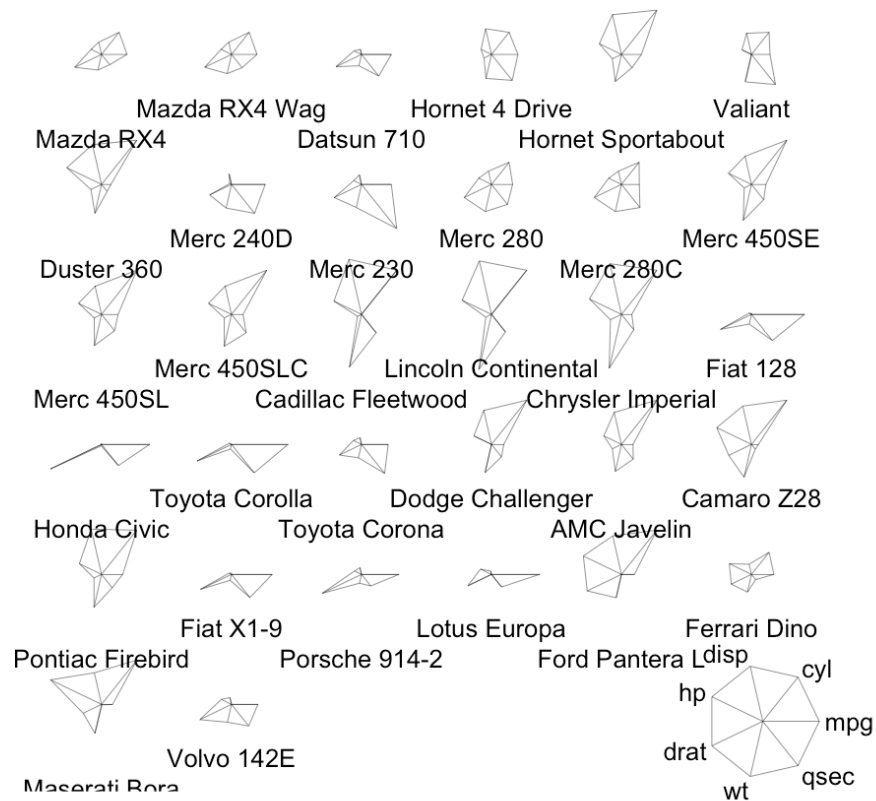
```
plot(mtcars$hp, mtcars$mpg, pch = 20,
     xlab = "Horse power", ylab = "Miles per Gallon",
     ylim=c(0,40))
symbols(mtcars$hp, mtcars$mpg,
        rectangles = cbind(mtcars$wt,mtcars$drat),
        fg=mtcars$cyl, lwd=2, add = TRUE)
title("Width: weight, height: rear axel ratio")
legend("topright", c("4 cyl", "6 cyl", "8 cyl"),
      lty=1, col=c(4,6,8), lwd=2)
```



1.2.4 Star plots

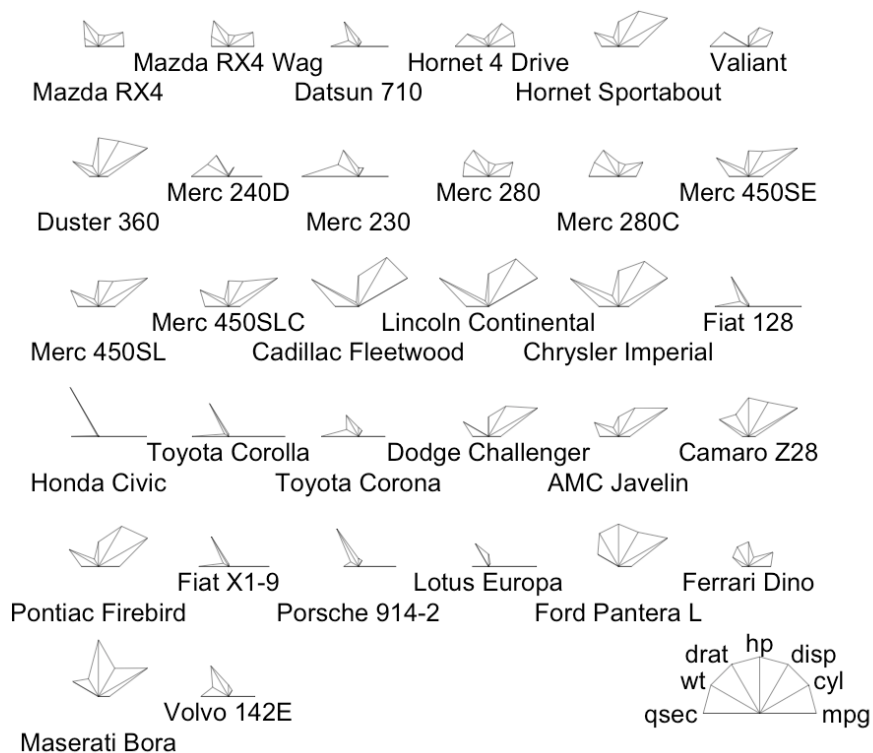
```
stars(mtcars[, 1:7], key.loc = c(14, 2), scale=T,
      main = "Motor Trend Cars")
```


Motor Trend Cars



```
stars(mtcars[, 1:7], key.loc = c(14, 2), scale=T,
      main = "Motor Trend Cars", full = FALSE)
```

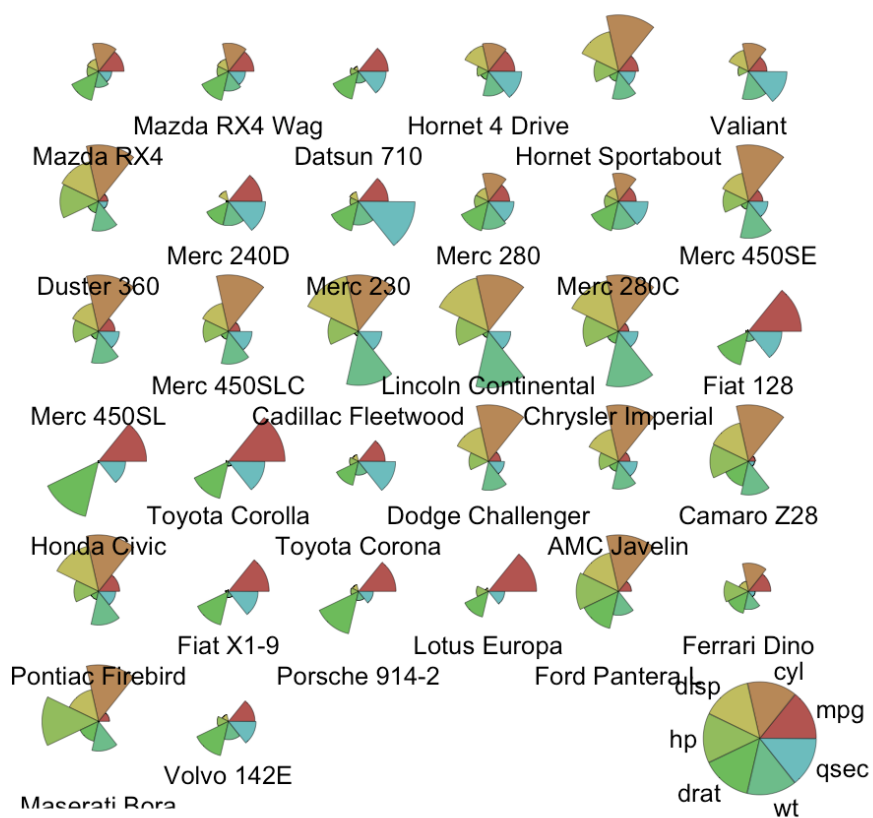
Motor Trend Cars



Better approach, segment plot with colors.

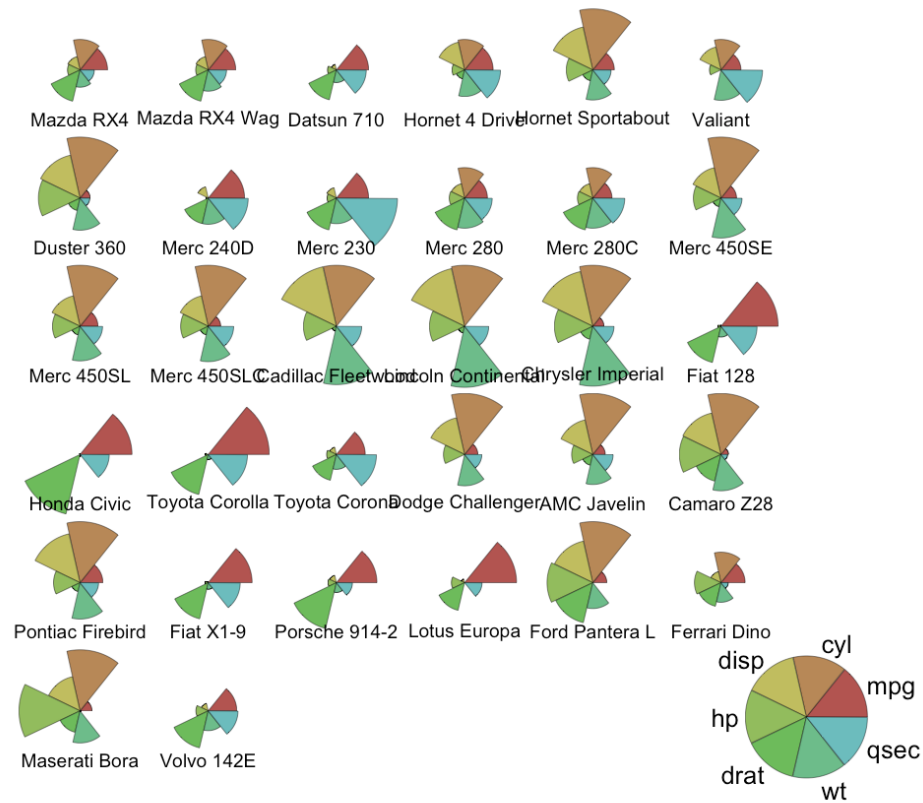
```
palette(rainbow(12, s = 0.6, v = 0.75))
stars(mtcars[, 1:7], key.loc = c(14,2), scale=T,
      main = "Motor Trend Cars",
      draw.segment=TRUE)
# with more control over position of labels:
loc <- stars(mtcars[, 1:7], key.loc = c(14,2), scale=T,
             main = "Motor Trend Cars",
             draw.segment=TRUE)
```

Motor Trend Cars



```
loc <- stars(mtcars[, 1:7], key.loc = c(14,2), scale=T,
  labels=NULL, main = "Motor Trend Cars",
  draw.segment=TRUE)
# loc contains the centers of the segment plots
# write name of car in the middle of each segment plot
text(loc[,1],loc[,2]-.8,
  row.names(mtcars),
  col = "black", cex = 0.6, xpd = TRUE)
```

Motor Trend Cars

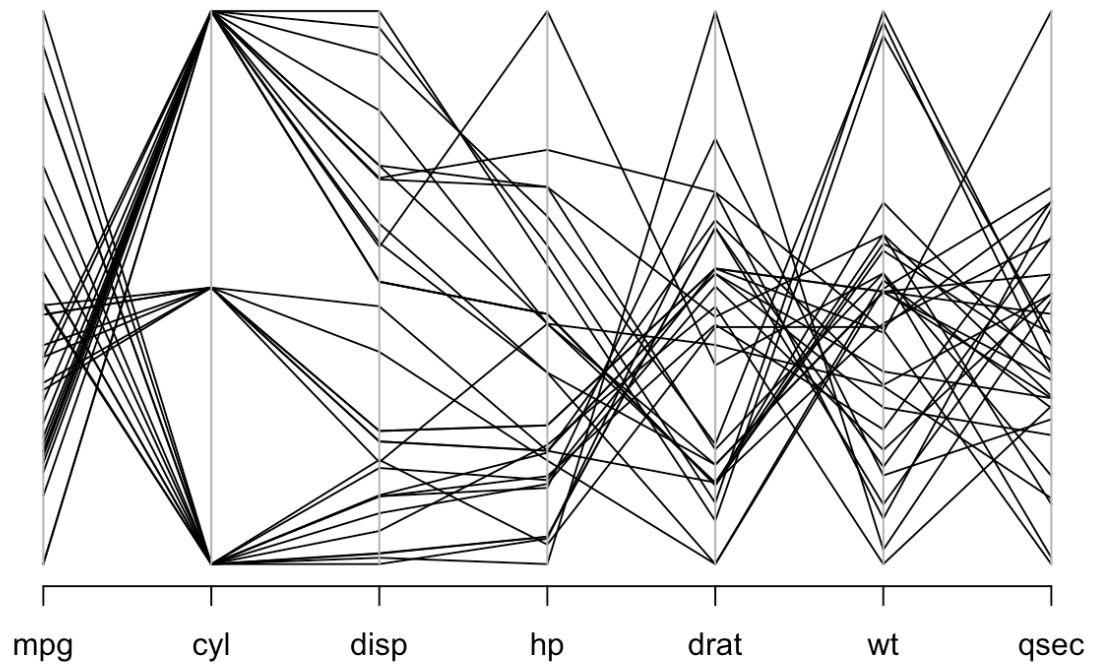


```
palette("default") # set colors back to default
```

1.2.5 Parallel coordinates plot

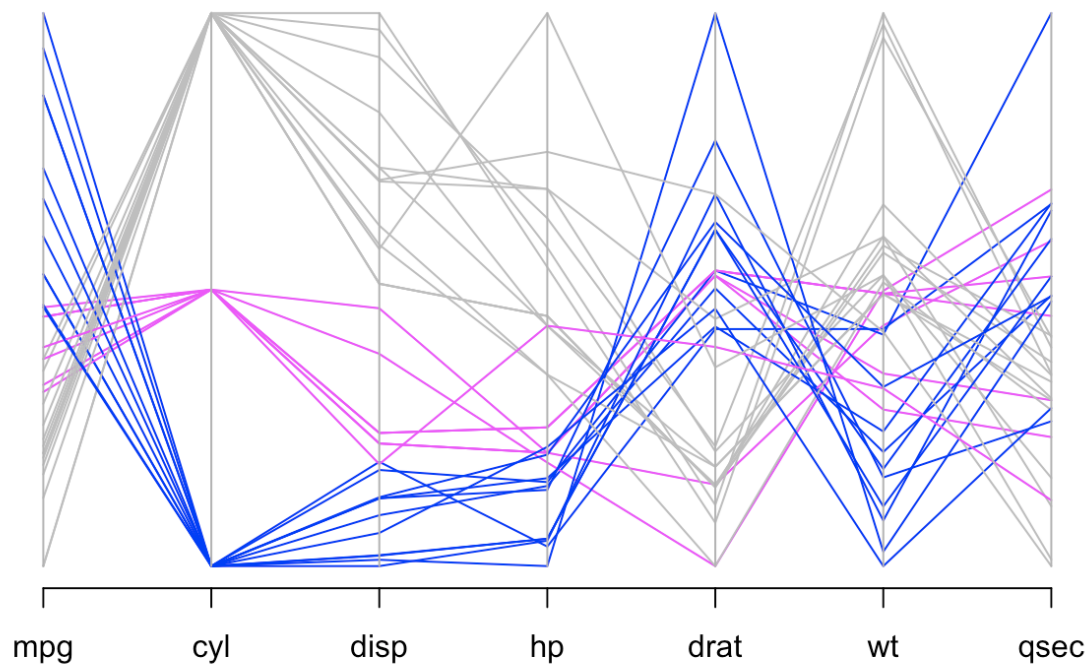
This is one of my favourites for high-dimensional data.

```
library(MASS)
parcoord(mtcars[,1:7])
```



Color can help.

```
parcoord(mtcars[,1:7], col=mtcars$cyl)
```



1.3 Random vectors and matrices

1.3.1 Random matrix

```
x<-matrix(rnorm(6), ncol=2)
x
```

```
##           [,1]      [,2]
## [1,]  2.4868938 -0.5839087
## [2,] -0.7161627  0.1765573
## [3,]  0.1950095 -0.4694244
```

Notice that `mean(x)` DOES NOT produce what we want.

```
mean(x)
```

```
## [1] 0.1814941
```

Empirical mean.

```
n<-dim(x)[1]
ones<-matrix(rep(1,n),ncol=1)
mu<-t(x) %*% ones / n
print(mu)
```

```
##           [,1]
## [1,]  0.6552469
## [2,] -0.2922586
```

1.3.2 Variance and standard deviation of a vector

```
x
```

```
##           [,1]      [,2]
## [1,]  2.4868938 -0.5839087
## [2,] -0.7161627  0.1765573
## [3,]  0.1950095 -0.4694244
```

```
var(x[,1])
```

```
## [1] 2.723757
```

```
var(x[,2])
```

```
## [1] 0.1681179
```

```
sd(x[,1])
```

```
## [1] 1.650381
```

```
sd(x[,2])
```

```
## [1] 0.4100219
```

```
# covariance
```

```
var(x[,1], x[,2])
```

```
## [1] -0.5478001
```

Variance-covariance matrix.

```
var(x)
```

```
##           [,1]      [,2]
## [1,]  2.7237566 -0.5478001
## [2,] -0.5478001  0.1681179
```

Correlation matrix.

```
cor(x)
```

```
##           [,1]      [,2]
## [1,]  1.0000000 -0.8095263
## [2,] -0.8095263  1.0000000
```

1.3.3 Sample variance-covariance

3x3 matrix of 1s.

```
ones %*% t(ones)
```

```
##           [,1] [,2] [,3]
## [1,]      1      1      1
## [2,]      1      1      1
## [3,]      1      1      1
```

Identity matrix.

```
diag(3)
```

```
##           [,1] [,2] [,3]
## [1,]      1      0      0
## [2,]      0      1      0
## [3,]      0      0      1
```

Matrix computation of S (unbiased)

```
(1/(n-1)) * t(x) %*% (diag(3)-(1/n)*ones %*% t(ones)) %*% x
```

```
##           [,1]      [,2]
## [1,]  2.7237566 -0.5478001
## [2,] -0.5478001  0.1681179
```

Produces the same result

```
var(x)
```

```
##           [,1]      [,2]
## [1,]  2.7237566 -0.5478001
## [2,] -0.5478001  0.1681179
```

2 Salient features of Big Data

2.1 Outliers in higher dimensions are not obvious

2.1.1 Outliers in univariate case

```
set.seed(123)
dat <- matrix(rnorm(5*100),100,5)
summary(dat)
```

##	V1	V2	V3
## Min.	:-2.30917	Min. :-2.0532	Min. :-1.75653
## 1st Qu.:	-0.49385	1st Qu.:-0.8011	1st Qu.:-0.53131
## Median :	0.06176	Median :-0.2258	Median : 0.03591
## Mean :	0.09041	Mean :-0.1075	Mean : 0.12047
## 3rd Qu.:	0.69182	3rd Qu.: 0.4678	3rd Qu.: 0.76363
## Max. :	2.18733	Max. : 3.2410	Max. : 2.29308


##	V4	V5
## Min.	:-2.465898	Min. :-2.6609
## 1st Qu.:	-0.729376	1st Qu.:-0.3964
## Median :	-0.003509	Median : 0.1651
## Mean :	-0.036223	Mean : 0.1059
## 3rd Qu.:	0.688690	3rd Qu.: 0.7216
## Max. :	2.571458	Max. : 2.3975


```
dat[23,4] <- dat[23,4] * 10
summary(dat)
```

##	V1	V2	V3
## Min.	:-2.30917	Min. :-2.0532	Min. :-1.75653
## 1st Qu.:	-0.49385	1st Qu.:-0.8011	1st Qu.:-0.53131
## Median :	0.06176	Median :-0.2258	Median : 0.03591
## Mean :	0.09041	Mean :-0.1075	Mean : 0.12047
## 3rd Qu.:	0.69182	3rd Qu.: 0.4678	3rd Qu.: 0.76363
## Max. :	2.18733	Max. : 3.2410	Max. : 2.29308

##	V4	V5
## Min.	:-13.387743	Min. :-2.6609
## 1st Qu.:	-0.729376	1st Qu.:-0.3964
## Median :	-0.003509	Median : 0.1651
## Mean :	-0.156713	Mean : 0.1059
## 3rd Qu.:	0.688690	3rd Qu.: 0.7216
## Max. :	2.571458	Max. : 2.3975

```
pairs(dat)
```



Extract its location

```
which.min(dat[,4])
```


```
## [1] 23
```

2.1.2 Multivariate outliers

```
load(file = "simpleExample.rda")
```

There are no clear univariate outliers here


```
par(mfrow=c(1,4))
plot(dat, xlim = c(-5,5), ylim = c(-5,5))
barplot(dat[,1], main="x values")
barplot(dat[,2], main="y values")
```



Using Mahalanobis distance

```
d <- mahalanobis(dat, colMeans(dat), cov(dat))
barplot(d, main="Mahalanobis")
```

Mahalanobis




```
which.max(d)
```

```
## [1] 301
```


Create chi-squared QQ-plot.

```
par(mfrow=c(1,1))
qqplot(qchisq(ppoints(301), df = 2), d,
       main = expression("Q-Q plot for" ~- {chi^2}[nu == 2]))
```



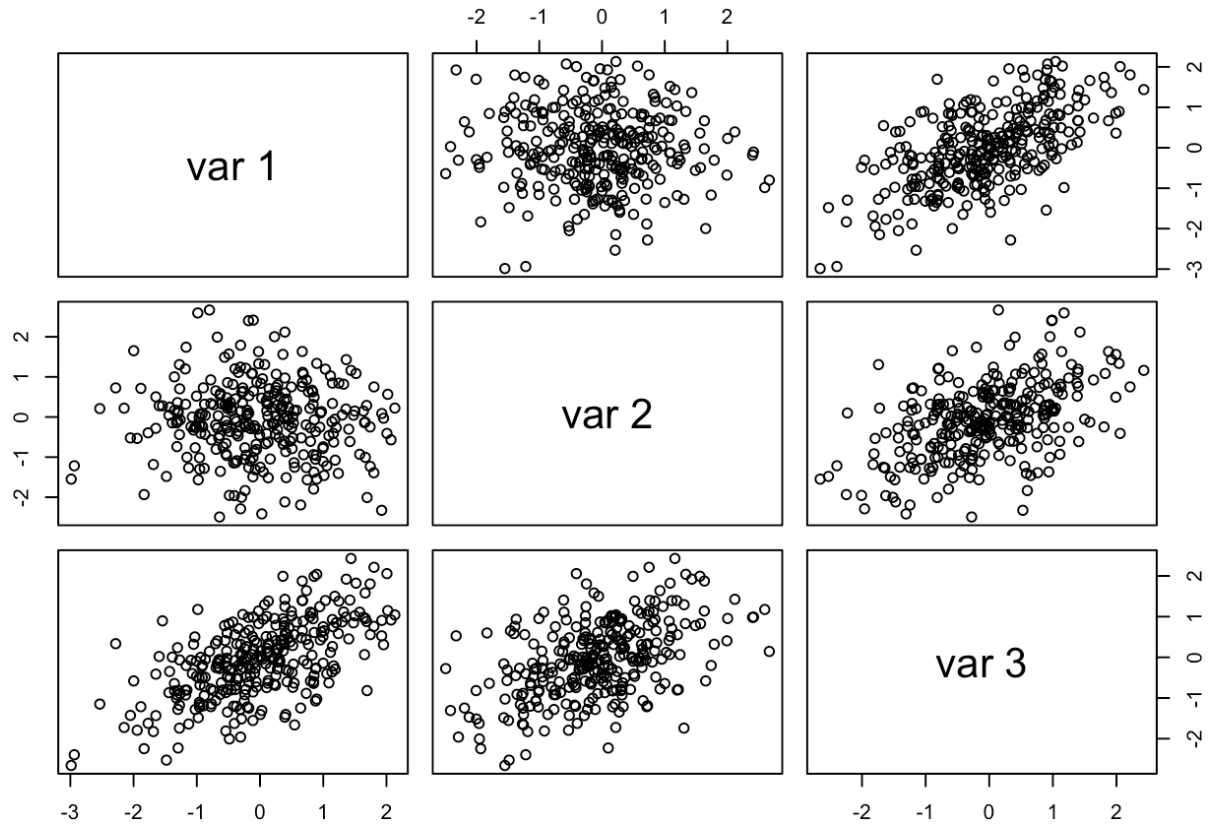
Now the outlier is clearly visible.

```
par(mfrow = c(1,1))
plot(dat, xlim=c(-5,5), ylim=c(-5,5))
points(dat[301,1],dat[301,2],col="red")
```



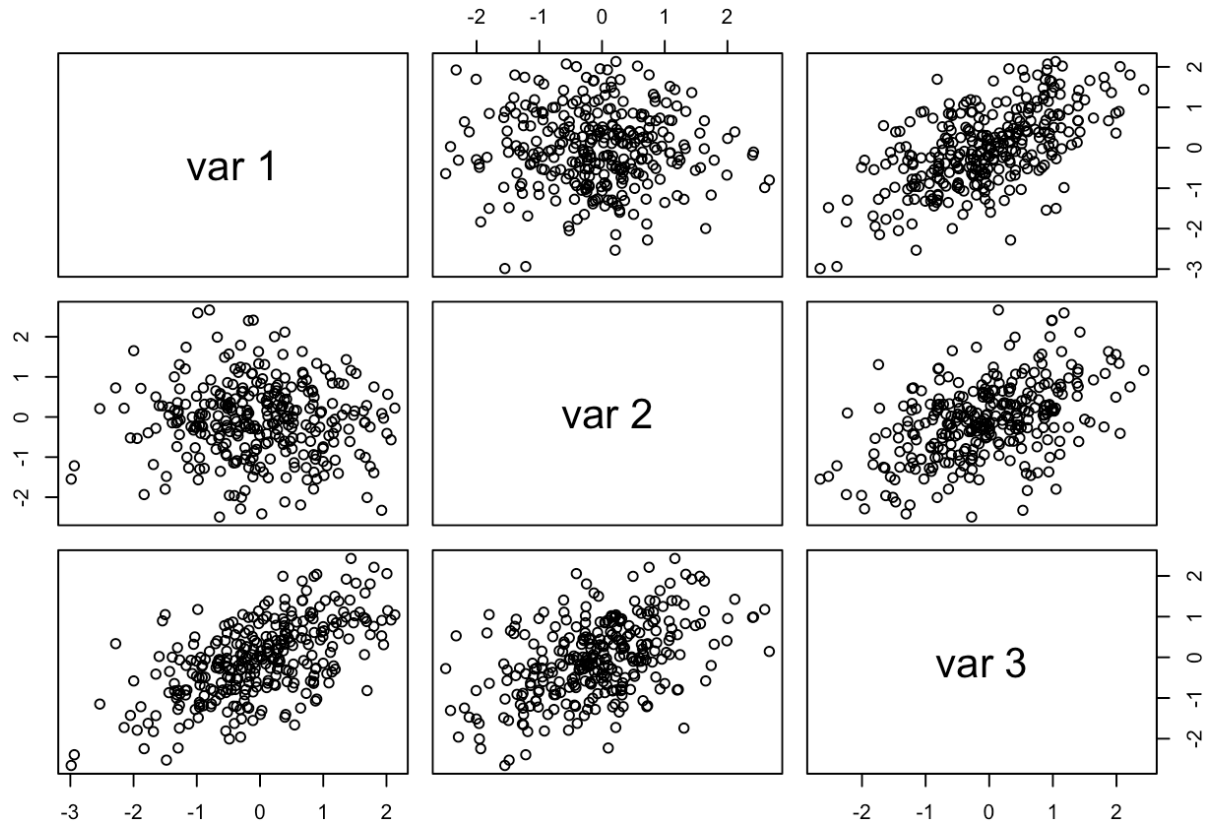
2.1.3 More dimensions

```
load(file = "3dExample.rda")  
pairs(dat)
```

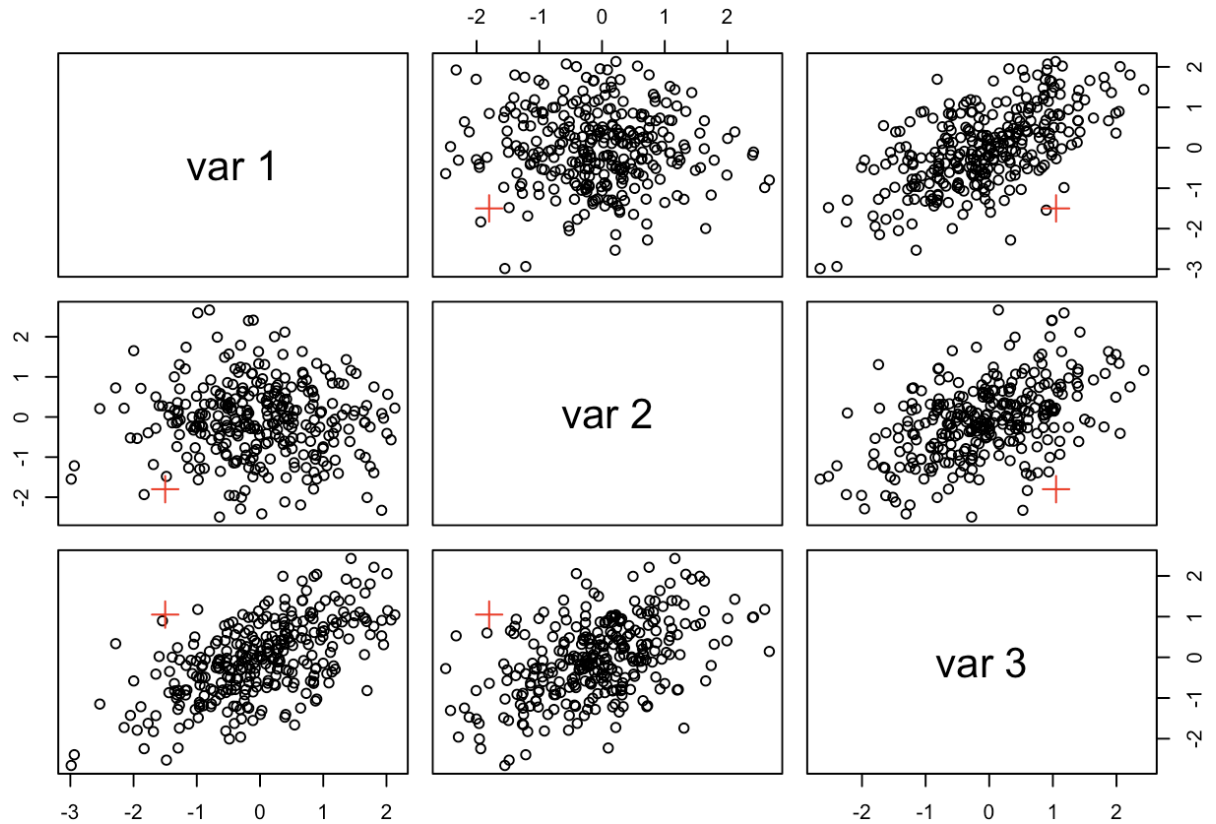


Introduce an outlier.

```
outFactor <- 1.5  
dat <- rbind(dat, outFactor*c(-1,-1.2,0.7))  
pairs(dat)
```




```
pairs(dat, col = c(rep(1,300), 2), pch = c(rep(1,300), 3), cex = c(rep(1,300), 2))
```

In none of the plots, the point is an outlier.

```
d <- mahalanobis(dat, colMeans(dat), cov(dat))
barplot(d, main="Mahalanobis")
```

Mahalanobis

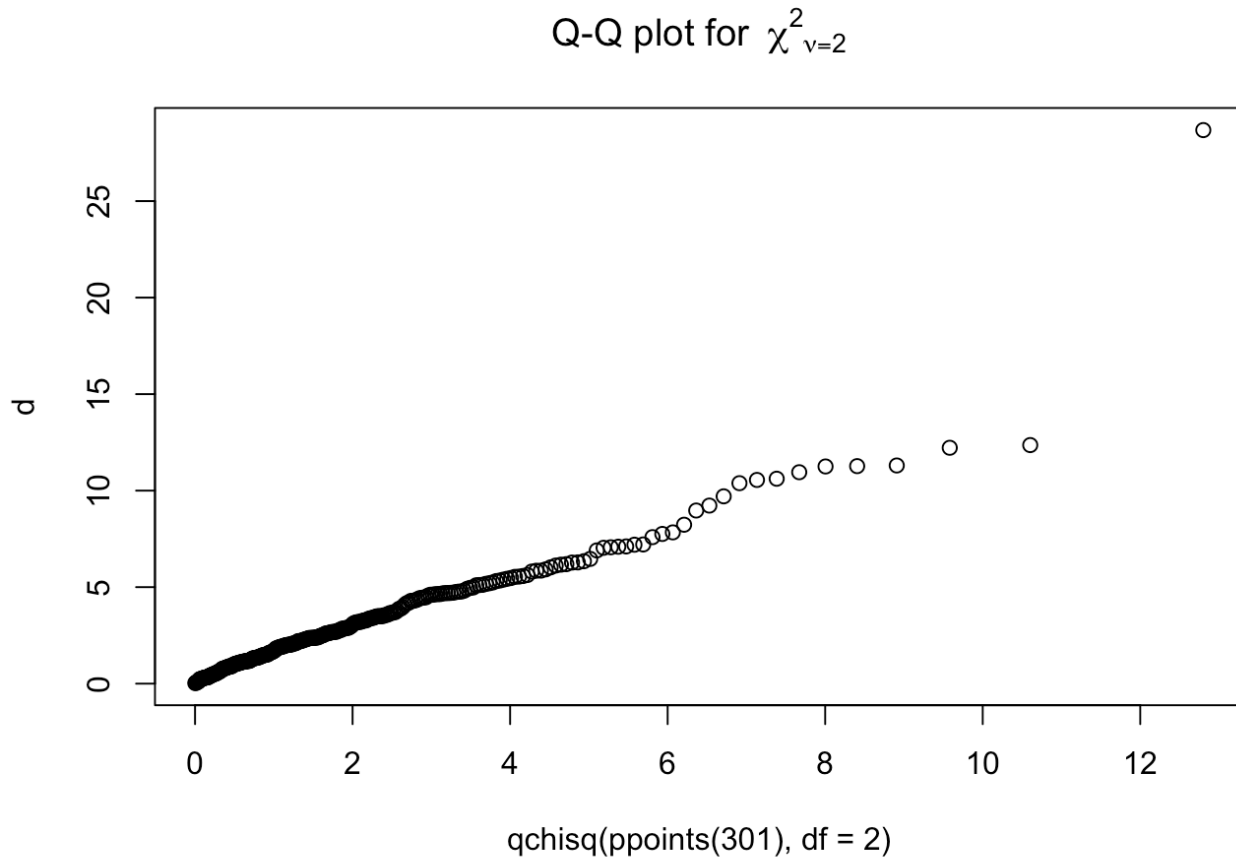


```
which.max(d)
```

```
## [1] 301
```

Create chi-squared QQ-plot, which will help show that it is a multivariate outlier.

```
par(mfrow=c(1,1))
qqplot(qchisq(ppoints(301), df = 2), d,
       main = expression("Q-Q plot for" ~~ {chi^2}[nu == 2]))
```




Let's do a fancy 3D plot. First run `install.packages('rgl')` to install the `rgl` package.

```
library(rgl)
plot3d(dat, col = c(rep(1,300), 2))
```

2.2 Heterogeneity

```
set.seed(123)
mus <- rnorm(5, mean=0, sd=10)
sds <- rchisq(5, 1)

library(MASS)
Sigma <- diag(sds)
X <- as.vector(mvrnorm(n=3000, mus, Sigma))
hist(X, breaks=100)
```



```
summary(X)
```

```
##      Min.  1st Qu.  Median    Mean  3rd Qu.    Max.
## -13.7700  -2.3920   0.6767   1.9200   1.6330   17.2600
```


```
mu <- mean(X)
```

```
sd <- sd(X)
```

```
hist(X, breaks=100, freq=FALSE, main="Bad model that doesn't capture heterogeneity")
```

```
curve(dnorm(x, mean=mu, sd=sd), col=2, lty=2, lwd=2, add=TRUE)
```

Bad model that doesn't capture heterogeneity



2.3 Noise accumulation

Try to replicate the example (Fig. 1) in the “Noise Accumulation” section of the paper *Challenges of Big Data analysis* by Fan, Han, and Liu. Useful functions are `prcomp` and `mvrnorm` (in the `MASS` package).

We setup the number of observations n and the number of dimensions p .

```
n <- 100
p <- 1000
```

The first class has a zero mean.

```
mu <- rep(0, p)
```

The second class has mean η . It is a sparse mean with 0 entries everywhere except for the first 10 dimensions where the mean is 3.


```
eta <- rep(0, p)
eta[1:10] <- 3
```

Sample from each class.

```
I <- diag(1,p,p)
X <- mvrnorm(n=n, mu, I)
Y <- mvrnorm(n=n, eta, I)
```

Plot the first two dimensions against each other for the two classes.

```
plot(X[,1], X[,2], bg=1, pch=19, xlim=c(-5,5), ylim=c(-5,5))
points(Y[,1], Y[,2], col=2, bg=2, pch=24)
```



Stack the data into one matrix.

```
Z <- rbind(X,Y)

#m <- 2
#pcaZ <- prcomp(t(Z), scale=TRUE)
#features <- pcaZ$rotation[1:m,]
#Zc <- t(features) %*% t(Z)
#Za <- t(features) %*% Zc
```

```
#dim(Za)

# pca <- prcomp(Z[,1:500], center=FALSE, scale=TRUE, retx=TRUE)
# pX <- pca$x[1:n,1:2]
# pY <- pca$x[(n+1):nrow(pZ),1:2]
# plot(pX[,1], pX[,2], bg=1, pch=19, xlim=c(-5,5), ylim=c(-5,5))
# points(pY[,1], pY[,2], col=2, bg=2, pch=24)
```