

Tutorial 5 Solutions

STAT 3013/8027

- Rice Chapter 8, Questions 13, 17 (a, b, c). See handwritten solutions.
- Question 25:
 - A priori I believe the probability of landing up (θ) is between 0.10 and 0.40. A specific value would be 0.25.
 - While you may not have found a thumbtack to throw, there is a data set in R of 100 tosses! Use the first 20 observations (1 = “Up”, 0 = “Down”)

```
library("isdals")
data(thumbtack)

y <- thumbtack[1:20]
y
```


```
## [1] 1 1 0 0 1 1 0 1 0 0 1 0 1 1 0 0 1 1 1 0
```

The likelihood of the data:

$$\begin{aligned} L(\theta|\mathbf{y}) &= \prod_{i=1}^{n=20} \theta^{y_i} (1-\theta)^{1-y_i} \\ &= \theta^{\sum_{i=1}^{20} y_i} (1-\theta)^{20-\sum_{i=1}^{20} y_i} \\ \ell(\theta|\mathbf{y}) &= \left(\sum_{i=1}^{20} y_i \right) \log(\theta) + \left(20 - \sum_{i=1}^{20} y_i \right) \log(1-\theta) \end{aligned}$$

```
theta <- seq(0.1, 0.99, by=0.01)
log.lik <- NULL
for(i in 1:length(theta)){
  log.lik <- c(log.lik, sum(dbinom(y, 1, theta[i], log=TRUE)))
}

plot(theta, log.lik, type="l", lwd=3)
```



- Now let's run the experiment a bit differently . . . let's flip till we get 5 Ups:

```
check <- 0
c <- 1

while(check!=5){
  z <- thumbtack[21:(21+c)]
  check <- sum(z)
  c <- c+1
}

z

## [1] 1 0 0 1 1 0 0 0 0 1 1
```

Based on the experiment, the likelihood is based on a negative binomial (flip until we get 5 = r successes):

$$L(\theta|z) = \binom{11-1}{5-1} \theta^5 (1-\theta)^{14-5}$$

$$\ell(\theta|z) = \log \binom{11-1}{5-1} + 5 \log(\theta) + (11-5) \log(1-\theta)$$

Notice that this is the same likelihood as above except for a constant in front, which won't change the maximization!


```
theta <- seq(0.1, 0.99, by=0.01)
log.lik <- NULL
for(i in 1:length(theta)){
```

```

log.lik.i <- log(choose(11-1, 5-1)) + 5*log(theta[i]) + (11-5)*log(1-theta[i])
log.lik <- c(log.lik, log.lik.i)
}

plot(theta, log.lik, type="l", lwd=3)

```



- Now let's determine the distribution for θ under a uniform prior for θ .

$$\begin{aligned}
 p(\theta|\mathbf{y}) &= \frac{p(\mathbf{y}|\theta)p(\theta)}{p(\mathbf{y})} \\
 &\propto p(\mathbf{y}|\theta)p(\theta) \\
 &= \theta^{\sum_{i=1}^n y_i} (1-\theta)^{n-\sum_{i=1}^n y_i} \times 1 \\
 &= \theta^{\sum_{i=1}^n y_i} (1-\theta)^{n-\sum_{i=1}^n y_i}
 \end{aligned}$$

As θ is the random variable, we can see that this is a kernel for a $\text{beta}(a, b)$ distribution:

$$\theta^{(\sum_{i=1}^n y_i + 1) - 1} (1 - \theta)^{(n - \sum_{i=1}^n y_i + 1) - 1}$$

Where $a = (\sum_{i=1}^n y_i + 1)$ and $b = (n - \sum_{i=1}^n y_i + 1)$. Based on beta distribution the mean and variance are:

$$\begin{aligned}
 E[\theta|\mathbf{y}] &= \frac{\sum_{i=1}^n y_i + 1}{\sum_{i=1}^n y_i + 1 + n - \sum_{i=1}^n y_i + 1} \\
 &= \frac{\sum_{i=1}^n y_i + 1}{n + 2}
 \end{aligned}$$

```

n <- length(y)
a <- sum(y)+1
b <- n - sum(y)+1

m <- a/(a+b)
v <- (a*b)/((a+b)^2 * (a + b + 1))

m

```

```
## [1] 0.5454545
```

```
v
```

```
## [1] 0.01077973
```

- Let's plot the posterior based on observing \mathbf{y} along with a normal approximation based on the mean and variance above:

```


theta <- seq(0.1, 0.99, by=0.01)
post <- NULL
norm.approx <- NULL

for(i in 1:length(theta)){
  post.i <- dbeta(theta[i], sum(y)+1, 20-sum(y)+1)
  post <- c(post, post.i)

  norm.approx.i <- dnorm(theta[i], m, sqrt(v))
  norm.approx <- c(norm.approx, norm.approx.i)
}

plot(theta, post, type="l", lwd=3)
lines(theta, norm.approx, lwd=3, col="blue")

```



The normal approximation (blue) is very similar to the posterior (black).

- Now let's throw the tack 20 more times (label these x) and examine the two posteriors.

```
n.y <- length(y)


x <- thumbtack[40:59]
w <- c(y, x)
n.w <- length(w)

theta <- seq(0.1, 0.99, by=0.01)
post.y <- NULL
post.w <- NULL

for(i in 1:length(theta)){
  post.y.i <- dbeta(theta[i], sum(y)+1, n.y-sum(y)+1)
  post.y <- c(post.y, post.y.i)

  post.w.i <- dbeta(theta[i], sum(w)+1, n.w-sum(w)+1)
  post.w <- c(post.w, post.w.i)
}

plot(theta, post.y, type="l", lwd=3, ylim=c(0, 6))
lines(theta, post.w, lwd=3, col="orange")
```




From the figure, with the full 40 data points, we see that are beliefs have changed (orange posterior). Also the variability is smaller.

- A question for you all: What would the posterior look like if I used the first 20 data points (y) and calculated a posterior. Now used the posterior as the prior for θ and observed the next 20 data points (x). What would that posterior look like?
- Question 43:
 - a. Let's load in the data and examine a histogram of the data.

```
data <- read.table("gamma-arrivals.txt")
y <- data$V1
hist(y, col="azure2")
```

Histogram of y



Based on the histogram, an $\text{gamma}(a,b)$ distribution does not seem like an unreasonable model for the data.

$$f(y) = \frac{1}{\gamma(a)b^a} y^{a-1} \exp(-y/b)$$

b. Let's first determine the **Method of Moments** for a and b . Our system of equations is:

$$\begin{aligned} E[Y] = ab &= \frac{1}{n} \sum_{i=1}^n y_i = m_1 \\ E[Y^2] = ab^2 + a^2b^2 &= \frac{1}{n} \sum_{i=1}^n y_i^2 = m_2 \end{aligned}$$

Solving this system of equations we have:

$$\tilde{a} = \frac{m_1^2}{m_2 - m_1^2} \quad \tilde{b} = \frac{m_2 - m_1^2}{m_1}$$

```
n <- length(y)
m1 <- mean(y)
m2 <- sum(y^2)/n

a.mom <- m1^2/(m2-m1^2)
```

```

b.mom <- (m2-m1^2)/m1

a.mom

## [1] 1.012352

b.mom

## [1] 78.95989

```

Now let's consider the **Maximum Likelihood** estimators.

$$L(a, b) = \prod_{i=1}^n \frac{1}{\Gamma(a)} \frac{1}{b^a} y_i^{a-1} \exp(-y_i/b)$$

Here we have the log-likelihood:

$$\ell(a, b) = -n \log(\Gamma(a)) - n a \log(b) + (a-1) \sum_{i=1}^n \log(y_i) - \sum_{i=1}^n y_i/b$$

$$\frac{\partial \ell(a, b)}{\partial a} = -n\psi(a) - n\log(b) + \sum_{i=1}^n \log(y_i) \quad (1)$$

$$\frac{\partial^2 \ell(a, b)}{\partial a^2} = -n\psi'(a) \quad (2)$$

In Eqn (1), let's substitute in for $b = \sum_{i=1}^n y_i / (na)$. So we have an equation which only has a :

$$\frac{\partial \ell(a, b)}{\partial a} = -n\psi(a) - n\log\left(\sum_{i=1}^n y_i / (na)\right) + \sum_{i=1}^n \log(y_i) \quad (3)$$

$$\frac{\partial^2 \ell(a, b)}{\partial a^2} = -n\psi'(a) \quad (4)$$

- Where $\psi(a) = \text{digamma}(a)$ and $\psi'(a) = \text{trigamma}(a)$.

```

## Let's find the MLE of a using the N-R Approach.
## Then we can solve for b analytically
## Write some functions for U and H
U <- function(a){
  n <- length(y)
  out <- -n* digamma(a) - n*log(sum(y)/(n*a)) + sum(log(y))
  return(out)
}

H <- function(a){
  n <- length(y)

```



```

    out <- -n*trigamma(a)
    return(out)
  }

## Starting values - use MoM estimator
a <- mean(y)^2/( (n-1)*var(y)/n)

## set a stopping point
eps <- 1e-07
check <- 10

## Save the results.
out <- a

## Run the algorithm
while(check > eps){
  a.new <- a - U(a)/H(a)
  check <- sum(abs(a-a.new))
  a <- a.new
  out <- rbind(out, t(a))
}

a.mle <- a
b.mle <- sum(y)/(n*a.mle)

##
a.mle

## [1] 1.026332
b.mle

## [1] 77.8844

```

The ML and MoM estimates are very similar.

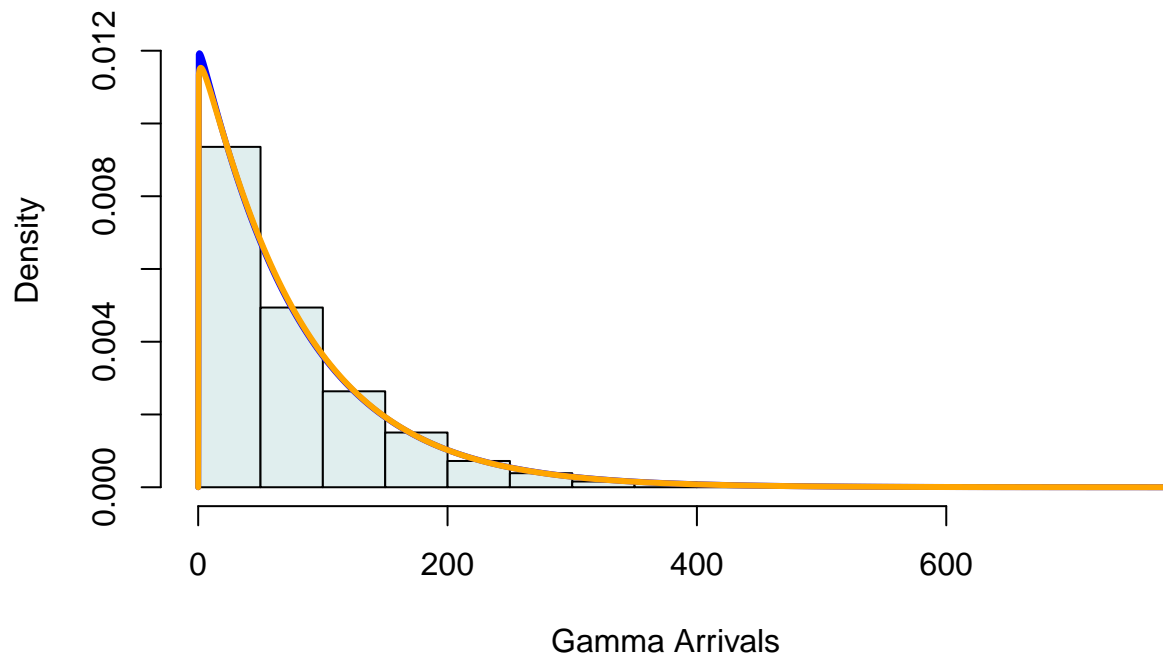
c. Let's plot the two fitted densities on top of the histogram.

```

hist(y, col="azure2", freq=FALSE, ylim=c(0, 0.013), xlab="Gamma Arrivals", main="MoM (bl
x <- seq(0, 800, by=0.5)
lines(x, dgamma(x, shape=a.mom, scale=b.mom), lwd=3, col="blue")
lines(x, dgamma(x, shape=a.mle, scale=b.mle), lwd=3, col="orange")

```

MoM (blue) & MLE (orange)



We can see the results are very similar as the two lines essentiall overlap, except for a slight difference near 0.