



# Course Introduction

CSC343 - Introduction to Databases

Sina Meraji

Thanks to Ryan Johnson, John Mylopoulos, Arnold Rosenbloom  
and Renee Miller for material in these slides



## Course administrivia

- Read the course syllabus online!
- Course Website:
  - <http://www.cdf.toronto.edu/~csc343h/summer/index.shtml>
- Contact:
  - website and discussion board (MYbb) are required reading
  - personal matters: email me (include “343” in the subject)
- Office hours:
  - Wed 5-6 –

## COURSE ADMINISTRIVIA



## Course administrivia(cont.)

- TAs
  - Fatemeh Nargesian(fnargesian at cs.toronto.edu )
  - Aakar Gupta (aakar at cs.toronto.edu)
  - Krish Perumal( krishperumal11 at gmail.com)
- Tutorials
  - Wed 8:00-9:00 pm at BA1240, BA2145, BA3116



## Course prerequisites

- A&S students
  - CSC165/240 or MAT137/157 and
  - CSC207
  - Engineering students
  - Talk to me




## Course Marking Scheme

Work	Weight	Comment
3 Assignments	30%	10% each
1 Midterm Test	15%	-
Final Exam	45%	You must get $\geq 40\%$ to pass the course



## Recommended Resources

- A First Course in Database Systems, by Jeffrey D. Ullman and Jennifer Widom, 2008 (3rd Edition)
- Jennifer Widom's online mini-courses from Stanford (link on the website)



## Assignment Policies

- You may work with a partner on assignments
  - Can change partners between assignments
  - You may not dissolve a partnership in an assignment without permission
- Assignments must be submitted via MarkUs
  - Your code must run on our lab computers ("cdf")
- Late policy
  - 10% per day
  - You can submit up to 2 days late
  - **Submit on time!**



## To-do List

- Anyone new to the CDF labs:
  - You all have a CDF account; see details on the course website
  - Try logging in
- Read the course syllabus
- Bookmark the course website



## Today

- What is a database?
- Relational Model



## Today

- What is a database?
- Relational Model



## What is a database system?

- Database: a large, integrated **collection of data**
- Models relevant aspects of reality
  - **Entities** (teams, players)
  - **Relationships** (Lionel Messi *plays for* Barcelona)
  - **Constraints** (*at least* one **goalkeeper** per team)
  - More recently, active components (“business logic”)
- Database Management System (DBMS): a software system designed to **store, manage, and facilitate access to** databases

## In the beginning...

- There was *The Mainframe*

- Cost: millions
- Watts: millions
- Size: 2000 m<sup>2</sup>
- MIPS: 0.04
- Memory: 2kB
- Storage: 3.5MB (tape)



SAGE (1954)

*Few organizations could afford two!*

## Early computing challenges

- Time sharing

- ~100 terminals per mainframe
- Users share hardware
- Want to share *data*, too



SABRE (1960)

=> "The Database"

- Bare hardware

- No OS
- No device drivers
- No file system



UNIVAC (1951)

=> File Management System

## "The Database"

- Abstract concept dating back to the 1950's

- *Centralized* repository for all the enterprise's data
- *Real-time* updates from many sources
- *Concurrent* access by many users
- *Interactive* (ad-hoc) exploration and reporting

Example System: Semi Automatic Ground Environment (SAGE)

- **Goal:** Produce a single unified image of the airspace over an area
- Computer-aided tracking and interception of aircraft
- Dozens of SAGE installations (big one in North Bay)
- Hundreds of radar stations throughout North America
- Thousands of operators

*Goal: all relevant information at your fingertips*

## File management systems (FMS)

- File management ca. 1935

- **File:** box of punchcards
- **Metadata:** label on the box
- **Ad-hoc report:** no big deal
- **Hardware change:** no big deal



- File management ca. 1955

- **File:** several km of magnetic tape
- **Metadata:** embedded in application logic
- **Ad-hoc report:** hire a couple programmers
- **Hardware change:** hire a dozen programmers...

*Huge need for portability, abstraction*



## Database Management System

- File management systems meet *The Database*
  - Protect users from each other (isolation, consistency)
  - Protect application from data changes (at logical level)
  - Protect data from hardware changes (at physical level)
- Split personality remains to this day
  - Theory/applications (declarative access to changing data)
  - Systems (make it run fast on ever-changing hardware)

*This semester: the theory/application side*



## Why study databases?

- Shift from computation to information
  - always true for corporate computing
  - Web made this point for personal computing
  - more and more true for scientific computing
- Need for DBMS has exploded
  - **Corporate**: retail swipe/clickstreams, “customer relationship mgmt”, “supply chain mgmt”, “data warehouses”, etc.
  - **Scientific**: digital libraries, Human Genome project, Sloan Digital Sky Survey, physical sensors, grid physics network
- A practical discipline spanning much of CS
  - OS, languages, theory, AI, multimedia, logic
  - Yet with a focus on real-world apps

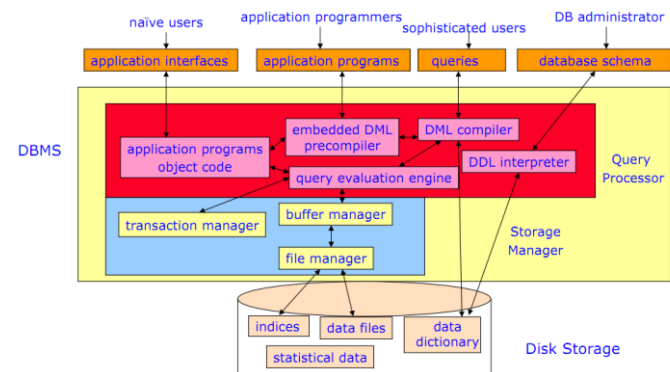


## CSC343 Topics

- The Relational Data Model
- Relational Algebra
- Defining, Querying and Manipulating Databases
  - the Structured Query Language (SQL)
- Application Programming with SQL
- Database Design and Normalization
- Modeling and Querying semi-structured data
- Other Topics (SQL security, NoSQL, ...)



## DBMS High-level Architecture





## Advantages of a DBMS

- Data independence
- Efficient data access
- Data integrity & security
- Data administration
- Concurrent access, crash recovery
- Reduced application development time
- So why not use them always?
  - Expensive/complicated to set up & maintain
  - Cost & complexity must be offset by need
  - General-purpose, not suited for special-purpose tasks (e.g. text search!)



## Summary (part 1)

- DBMS marries two very old concepts
  - The Database (idealistic vision)
  - File management system (imminently practical)
- Benefits
  - Maintain, query large datasets
  - Manipulate data and exploit semantics
  - Recover from system crashes
  - Juggle/balance concurrent access, automatic parallelization
  - Quick application development
  - Preserve data integrity and security
- Powerful abstractions provide data independence
  - Application safe from changes to data organization, hardware



## What comes next?

- If you are heading for industry:
  - Database professionals are in demand and well paid
- If you want to do research:
  - Many interesting problems ahead [The Claremont Report'08]
    - Revisiting Database Engines
    - Declarative Programming for Emerging Platforms
    - The Interplay of Structured and Unstructured Data
    - Cloud Data Services
    - Mobile Applications and Virtual Worlds
- Further studies in databases at DCS:
  - csc443: Database Systems Technology
  - csc2508: Advanced DBMSs (grad course, taken with permission)



## Summary (cont.)

DB administrators, developers  
are the bedrock of the  
information economy



Data management R&D spans a  
broad, fundamental branch of  
the science of computation

*This semester: become an effective DBMS user*

## Today

- What is a database?
- Relational Model

## Data Models

- Data model: a notation for describing data, including
  - the **structure** of the data
  - **constraints** on the content of the data
  - **operations** on the data
- Many possible data models:
  - network data model
  - hierarchical data model
  - relational data model -- the most widely used
  - semistructured model (later in the term)


## Comparing data models

### Student job example


Mary (M) and Xiao (X) both work at Tim Hortons (T)

Jaspreet (J) works at both Bookstore (B) and Wind (W)

### Network (graph)



### Hierarchical (tree)



### Relational (table)

E	S	R
B ...	J ...	M T
T ...	M ...	X T
W ...	X ...	J B
		J W

## Why the relational model?

- Matches how we think about data
- Real reason: **data independence**!
- Earlier models tied to physical data layout
  - Procedural access to data (low-level, explicit access)
  - Relationships stored in data (linked lists, trees, etc.)
  - **Change in data layout => application rewrite**
- Relational model
  - Declarative access to data (system optimizes for you)
  - Relationships specified by queries (schemas help, too)
  - **Develop, maintain apps and data layout separately**

*Similar battle today with languages*



## What is the relational model?

- Logical representation of data
  - Two-dimensional tables (relations)
- Formal system for manipulating relations
  - Relational algebra (coming next)
- Result
  - High-level (logical, declarative) description of data
  - Mechanical rules for rewriting/optimizing low-level access
  - Formal methods to reason about soundness

*Relational algebra is the key*



## Mathematical Relations

- Given sets  $D_1, D_2, \dots, D_n$ , not necessarily distinct, the **Cartesian product**  $D_1 \times D_2 \times \dots \times D_n$  is the set of all (ordered)  $n$ -tuples  $\langle d_1, d_2, \dots, d_n \rangle$  such that  $d_1 \in D_1, d_2 \in D_2, \dots, d_n \in D_n$
- A **mathematical relation** on  $D_1, D_2, \dots, D_n$  is a subset of the Cartesian product  $D_1 \times D_2 \times \dots \times D_n$ .
- $D_1, D_2, \dots, D_n$  are **domains** of the relation, while  $n$  is the **degree** of the relation.
- The number of  $n$ -tuples in a given relation is the **cardinality** of that relation




## The Relational Model

- Proposed by Edgar F. Codd in 1970 (**Turing Award**, 1981) as a data model that strongly supports data independence
- Made available in commercial DBMSs in 1981 -- it is not easy to implement data independence efficiently and reliably!
- It is based on (a variant of) the mathematical notion of **relation**
- Relations are represented as tables



## Relations (tables) and tuples (rows)





## An Example

- Games  $\subseteq$  String x String x Integer x Integer

Juve	Lazio	3	1
Lazio	Milan	2	0
Juve	Roma	1	2
Roma	Milan	0	1

- Note that String and Integer each play two roles, distinguished by means of position
- The structure of a mathematical relation is **positional**

## Attributes

- We can make the structure of a relation **non-positional** by associating a unique name (**attribute**) with each domain that describes its role in the relation
- In the tabular representation, attributes are used as column headings

HomeTeam	VisitingTeam	HomeGoals	VisitorGoals
Juve	Lazio	3	1
Lazio	Milan	2	0
Juve	Roma	1	2
Roma	Milan	0	1

## Notation

- $t[A]$  (or  $t.A$ ) denotes the value on attribute  $A$  for a tuple  $t$
- In our example, if  $t$  is the first tuple in the table  $t[\text{VisitingTeam}] = \text{Lazio}$
- The same notation is extended to sets of attributes, thus denoting tuples:  
 $t[\text{VisitingTeam}, \text{VisitorGoals}]$  is a tuple on two attributes,  $\langle \text{Lazio}, 1 \rangle$
- More generally, if  $X$  is a sequence of attributes  $A_1, \dots, A_n$ ,  $t[X]$  is  $\langle t[A_1], t[A_2], \dots, t[A_n] \rangle$

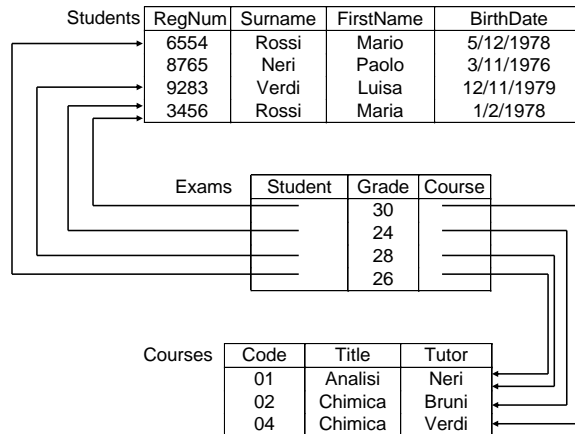
## Value-based References

Students	RegNum	Surname	FirstName	BirthDate
	6554	Rossi	Mario	5/12/1978
	8765	Neri	Paolo	3/11/1976
	9283	Verdi	Luisa	12/11/1979
	3456	Rossi	Maria	1/2/1978

Exams	Student	Grade	Course
	3456	30	04
	3456	24	02
	9283	28	01
	6554	26	01

Courses	Code	Title	Tutor
	01	Analisi	Neri
	02	Chimica	Bruni
	04	Chimica	Verdi

## Value-based References (cont.)



## Advantages of Value-Based References

- Value-based references lead to independence from physical data structures, such as pointers
  - Pointers are implemented differently on different hardware, inhibit portability of a database

## Definitions

- Relation schema:** Relation name R with a set of attributes  $A_1, \dots, A_n$ :

$R(A_1, \dots, A_n)$

- Database schema:** A set of relation schemas with different names

$D = \{R_1(X_1), \dots, R_n(X_n)\}$

- Relation (instance) on a relation schema**

$R(X)$ : Set  $r$  of tuples on  $X$

- Database (instance) on a schema**

$D = \{R_1(X_1), \dots, R_n(X_n)\}$ : Set of relations  $r = \{r_1, \dots, r_n\}$  (where  $r_i$  is a relation on  $R_i$ )

## Example Data

Da Mario			Da Mario			Da Mario		
Receipt No: 1357			Receipt No: 2334			Receipt No: 3007		
Date: 5/5/92			Date: 4/7/92			Date: 4/8/92		
3	covers	3.00	2	covers	2.00	2	covers	3.00
2	hors d'oeuvre	5.00	2	hors d'oeuvre	2.50	2	hors d'oeuvre	6.00
3	first course	9.00	2	first course	6.00	3	first course	8.00
2	steak	12.00	2	bream	15.00	1	bream	7.50
			2	coffee	2.00	1	salad	3.00
						2	coffee	2.00
Total:		29.00	Total:		27.50	Total:		29.50

## Data Representation

### Details

Number	Quantity	Description	Cost
1357	3	Covers	3.00
1357	2	Hors d'oeuvre	5.00
1357	3	First course	9.00
1357	2	Steak	12.00
2334	2	Covers	2.00
2334	2	Hors d'oeuvre	2.50
2334	2	First course	6.00
2334	2	Bream	15.00
2334	2	Coffee	2.00
3007	2	Covers	3.00
3007	2	Hors d'oeuvre	6.00
3007	3	First course	8.00
3007	1	Bream	7.50
3007	1	Salad	3.00
3007	2	Coffee	2.00

### Receipts

Number	Date	Total
1357	5/5/92	29.00
2334	4/7/92	27.50
3007	4/8/92	29.50

## Questions

- Have we represented all details of receipts?
- Well, it depends on what we are interested in:
  - does the order of lines matter?
  - could we have duplicate lines in a receipt?
    - If so, there is a problem ... Why?
- If needed, an alternative representation is possible ...

## More Detailed Representation

### Details

Number	Line	Quantity	Description	Cost
1357	1	3	Covers	3.00
1357	2	2	Hors d'oeuvre	5.00
1357	3	3	First course	9.00
1357	4	2	Steak	12.00
2334	1	2	Covers	2.00
2334	2	2	Hors d'oeuvre	2.50
2334	3	2	First course	6.00
2334	4	2	Bream	15.00
2334	5	2	Coffee	2.00
3007	1	2	Covers	3.00
3007	2	2	Hors d'oeuvre	6.00
3007	3	3	First course	8.00
3007	4	1	Bream	7.50
3007	5	1	Salad	3.00
3007	6	2	Coffee	2.00

### Receipts

Number	Date	Total
1357	5/5/92	29.00
2334	4/7/92	27.50
3007	4/8/92	29.50

## Incomplete Information: Motivation

(County towns have government offices, other towns do not.)

- Florence is a county town; so it has a government office, but we do not know its address
- Tivoli is not a county town; so it has no government office
- Prato has recently become a county town; has the government office been established? We don't know!

City	GovtAddress
Roma	Via IV novembre
Florence	?
Tivoli	??
Prato	???



## Null Value

- A **null value** is a special value (not a value of any domain) which denotes the absence of a value
- Types of Null Values:
  - **unknown value**: there is a domain value, but it is not known (Florence)
  - **non-existent value**: the attribute is not applicable for the tuple (Tivoli)
  - **no-information value**: we don't know if a value exists or not (Prato). (This is the disjunction - logical or - of the other two)
- DBMSs do not distinguish between these types: they implicitly adopt the no-information value



## Integrity Constraints

- An **integrity constraint** is a property that must be satisfied by all meaningful database instances
- A database is **legal** if it satisfies all integrity constraints
- Types of constraints:
  - **Intra-relational constraints**
    - domain constraints
    - tuple constraints
    - keys
  - **Inter-relational constraints**
    - foreign keys



## A Meaningless Database ...

Exams

RegNum	Name	Course	Grade	Honours
6554	Rossi	B01	K	
8765	Neri	B03	C	
3456	Bruni	B04	B	honours
3456	Verdi	B03	A	honours

Courses

Code	Title
B01	Physics
B02	Calculus
B03	Chemistry

Honours are awarded only if grade is A. Can you spot some others?



## Rationale for Integrity Constraints

- Describe the application in greater detail
- Contribute to **data quality**
- An important part of the database design process (we will discuss later **normal forms**)
- Used by the system in choosing a strategy for query processing



64

## Tuple and Domain Constraints

- A **tuple constraint** expresses conditions on the values of each tuple, independently of other tuples
  - `NOT((Honours = 'honours') OR (Grade = 'A'))`
  - `Net = Gross - Deductions`
- A **domain constraint** is a tuple constraint that involves a single attribute
  - `(Grade ≤ 'A') AND (Grade ≥ 'F')`



66

## Keys

- A **key** is a set of attributes that uniquely identifies tuples in a relation
- More formally:
  - A set of attributes  $K$  is a **superkey** for a relation  $r$  if  $r$  can not contain two distinct tuples  $t_1$  and  $t_2$  such that  $t_1[K] = t_2[K]$ ;
  - $K$  is a **key** for  $r$  if  $K$  is a minimal superkey; that is, there exists no other superkey  $K'$  such that  $K' \subset K$



65

## Unique Identification for Tuples

RegNum	Surname	FirstName	BirthDate	DegreeProg
284328	Smith	Luigi	29/04/59	Computing
296328	Smith	John	29/04/59	Computing
587614	Smith	Lucy	01/05/61	Engineering
934856	Black	Lucy	01/05/61	Fine Art
965536	Black	Lucy	05/03/58	Fine Art

- Registration number identifies students
  - no pair of tuples with the same value for **RegNum**
- Personal data could identify students as well
  - E.g. no pair of tuples with the same values for all of **Surname**, **FirstName**, **BirthDate**



67

## An Example

RegNum	Surname	FirstName	BirthDate	DegreeProg
284328	Smith	Luigi	29/04/59	Computing
296328	Smith	John	29/04/59	Computing
587614	Smith	Lucy	01/05/61	Engineering
934856	Black	Lucy	01/05/61	Fine Art
965536	Black	Lucy	05/03/58	Fine Art

- **RegNum** is a key
  - **RegNum** is a superkey and it contains a sole attribute, so it is minimal
- **Surname, Firstname, BirthDate** is a key
  - the three attributes form a superkey and there is no proper subset that is also a superkey

## Beware!

RegNum	Surname	FirstName	BirthDate	DegreeProg
296328	Smith	John	29/04/59	Computing
587614	Smith	Lucy	01/05/61	Engineering
934856	Black	Lucy	01/05/61	Fine Art
965536	Black	Lucy	05/03/58	Engineering

- There is no pair of tuples with the same values on both **Surname** and **DegreeProg**; i.e., in each programme students have different surnames.
- Can we conclude that **Surname** and **DegreeProg** form a key for this relation? No! There **could be** students with the same surname in the same programme

## Keys and Null Values

- If there are nulls, keys do not work well:
  - They do not guarantee unique identification
  - They do not help in establishing correspondences between data in different relations

RegNum	Surname	FirstName	BirthDate	DegreeProg
<b>NULL</b>	Smith	John	<b>NULL</b>	Computing
587614	Smith	Lucy	01/05/61	Engineering
934856	Black	Lucy	<b>NULL</b>	<b>NULL</b>
<b>NULL</b>	Black	Lucy	05/03/58	Engineering

How do we access the first tuple?

Are the third and fourth tuple the same?

## Existence of Keys (Proof Sketch)

- Relations are sets; therefore each relation is composed of distinct tuples
- It follows that the whole set of attributes for a relation defines a superkey
- Therefore each relation has a key, which is the set of all its attributes (or a subset thereof)
- The existence of keys guarantees that each piece of data in the database can be accessed

Keys are a major feature of the Relational Model and allow to say that it is “**value-based**”

## Primary Keys

- The presence of nulls in keys has to be limited
- Each relation must have a **primary key** on which **nulls are not allowed**
- Notation: the attributes of the primary key are **underlined**
- References between relations** are realized through **primary keys**

<u>RegNum</u>	Surname	FirstName	BirthDate	DegreeProg
643976	Smith	John	<b>NULL</b>	Computing
587614	Smith	Lucy	01/05/61	Engineering
934856	Black	Lucy	<b>NULL</b>	<b>NULL</b>
735591	Black	Lucy	05/03/58	Engineering

## References Between Relations

Students	<u>RegNum</u>	Surname	FirstName	BirthDate
	6554	Rossi	Mario	5/12/1978
	8765	Neri	Paolo	3/11/1976
	9283	Verdi	Luisa	12/11/1979
	3456	Rossi	Maria	1/2/1978

Exams	<u>Student</u>	Grade	<u>Course</u>
	3456	30	04
	3456	24	02
	9283	28	01
	6554	26	01

Courses	<u>Code</u>	Title	Tutor
	01	Analisi	Neri
	02	Chimica	Bruni
	04	Chimica	Verdi

## ...Consider...

- Suppose we want a database that maintains information on course offerings at the University of Toronto and use the following relation schema

`Course(name,dept,year,sem)`

- What would it mean if we used each of the following attribute sets as a primary key:

`name?`  
`name,sem,dept?`  
`sem,year?`  
`name,dept?`  
`dept,year?`

## Do we Always Have Primary Keys?

- In most cases **YES**
- In other cases **NO**
  - need to introduce new attributes by identifying **codes**
- Goal:** Unambiguously identify things
  - social insurance number
  - student number
  - area code
  - ...

## Referential Constraints (Foreign Keys)

- Data in different relations are referenced through (primary) key values
- Referential integrity constraints are imposed in order to guarantee that the values refer to existing tuples in the referenced relation
- For example, if a student with id "3456" took an exam for course with id "04", there better be a student with such an id and a course with such an id in the referenced relations
- Also called **inclusion dependencies**



## Example of Referential Constraints

Offences	<u>Code</u>	Date	Officer	Dept	Registration
	143256	25/10/1992	567	75	5694 FR
	987554	26/10/1992	456	75	5694 FR
	987557	26/10/1992	456	75	6544 XY
	630876	15/10/1992	456	47	6544 XY
	539856	12/10/1992	567	47	6544 XY

Officers	<u>RegNum</u>	Surname	FirstName
	567	Brun	Jean
	456	Larue	Henri
	638	Larue	Jacques

Cars	<u>Registration</u>	<u>Dept</u>	Owner	...
	6544 XY	75	Cordon Edouard	...
	7122 HT	75	Cordon Edouard	...
	5694 FR	75	Latour Hortense	...
	6544 XY	47	Mimault Bernard	...

## Referential Constraints

- A **referential constraint** requires that the values on a set X of attributes of a relation R1 must appear as values for the primary key of another relation R2
- In such a situation, we say that X is a **foreign key** of relation R1
- In the previous example, we have referential constraints between the attribute **Officer** of the relation **Offences** and the relation **Officers**; also between the attributes **Registration** and **Department** of relations **Offences** and **Cars**.



## Violation of Referential Constraints

Offences	<u>Code</u>	Date	Officer	Dept	Registration
	987554	26/10/1992	456	75	5694 FR
	630876	15/10/1992	456	47	6544 XY

Officers	<u>RegNum</u>	Surname	FirstName
	567	Brun	Jean
	638	Larue	Jacques

Cars	<u>Registration</u>	<u>Dept</u>	Owner	...
	7122 HT	75	Cordon Edouard	...
	5694 FR	93	Latour Hortense	...
	6544 XY	47	Mimault Bernard	...



## Referential Constraints: Comments

- Referential constraints play an important role in making the relational model **value-based**
- It is possible to have features that support the management of referential constraints (“actions” activated by violations)



## Summary

- The relational model
  - Relations, tuples, attributes
  - Value-based References
  - Incomplete information: The NULL value
  - Integrity Constraints
    - domain constraint
    - tuple constraint
    - unique tuple identification constraint (primary key)
    - referential constraints (foreign Key)
- Next: Relational Algebra

