

# Introduction to SQL (Structured Query Language)

Introduction to Databases  
Sina Meraji

Thanks to Ryan Johnson, John Mylopoulos, Arnold Rosenbloom  
and Renee Miller for material in these slides

## Announcement

- TA forgot to show up for the first hour of class on Wed May 27<sup>th</sup>!!!
- A1 deadline is extended to Friday June 5<sup>th</sup>
- Midterm will on Wednesday July 8, class time

## What is SQL?

- Declarative
  - Say “what to do” rather than “how to do it”
    - Avoid data-manipulation details needed by procedural languages
  - Database engine figures out “best” way to execute query
    - Called “query optimization”
    - Crucial for performance: “best” can be a million times faster than “worst”
- Data independent
  - Decoupled from underlying data organization
    - Views (= precomputed queries) increase decoupling even further
    - Correctness always assured... performance not so much
  - SQL is standard and (nearly) identical among vendors
    - Differences often shallow, syntactical

*Fairly thin wrapper around relational algebra*

## What does SQL look like?

- Query syntax

**SELECT** <desired attributes>  
**FROM** <one or more tables>  
**WHERE** <predicate holds for selected tuple>  
**GROUP BY** <key columns, aggregations>  
**HAVING** <predicate holds for selected group>  
**ORDER BY** <columns to sort>

## Example

Orders

OID	OrderDate	OrderPrice	Customer
1	2008/11/12	1000	Hansen
2	2008/10/23	1600	Nilsen
3	2008/09/02	700	Hansen
4	2008/09/03	300	Hansen
5	2008/08/30	2000	Jensen
6	2008/10/04	100	Nilsen

Find if the customers "Hansen" or "Jensen" have a total order of more than 1500

Query:

**SELECT** Customer, SUM(OrderPrice) AS Total  
**FROM** Orders  
**WHERE** Customer = 'Hansen' OR Customer = 'Jensen'  
**GROUP BY** Customer  
**HAVING** SUM(OrderPrice) > 1500  
**ORDER BY** Customer DESC

Query Result:


Customer	Total
Jensen	2000
Hansen	2000

## Other aspects of SQL

- Data Definition Language ("DDL")
  - Manipulate database schema
  - Specify, alter physical data layout
- Data Manipulation Language ("DML")
  - Manipulate data in databases
  - Insert, delete, update rows
- "Active" Logic
  - Triggers and constraints
  - User-defined functions, stored procedures
  - Transaction management/ Consistency levels

*We'll come back to these later in the course*

## What does SQL \*really\* look like?



*That's not so bad, is it?*

## SELECT-FROM-WHERE QUERIES



## 'SELECT' clause

- Identifies which attribute(s) query returns
  - Comma-separated list
  - => Determines schema of query result
- Optional: extended projection
  - Compute arbitrary expressions
  - Usually based on selected attributes, but not always
- Optional: rename attributes
  - "Prettify" column names for output
  - Disambiguate (E1.name vs. E2.name)
- Optional: specify groupings
  - More on this later
- Optional: duplicate elimination
  - SELECT DISTINCT ...



## 'FROM' clause

- Identifies the tables (relations) to query
  - Comma-separated list
- Optional: specify joins
  - ... but often use WHERE clause instead
- Optional: rename table
  - Using the same table twice (else they're ambiguous)
  - Nested queries (else they're unnamed)



## 'SELECT' clause – examples

- **SELECT** E.name ...
  - => Explicit attribute
- **SELECT** name ...
  - => Implicit attribute (error if R.name and S.name exist)
- **SELECT** E.name **AS** 'Employee name' ...
  - => Prettified for output (like table renaming, 'AS' usually not required)
- **SELECT** sum(S.value) ...
  - => Grouping (compute sum)
- **SELECT** sum(S.value)\*0.13 'HST' ...
  - => Scalar expression based on aggregate
- **SELECT** \* ...
  - => Select all attributes (no projection)
- **SELECT** E.\* ...
  - => Select all attributes from E (no projection)



## 'FROM' clause – examples

- ... **FROM** Employees
  - => Explicit relation
- ... **FROM** Employees **AS** E
  - => Table alias (most systems don't require "AS" keyword)
- ... **FROM** Employees, Sales
  - => Cartesian product
- ... **FROM** Employees E **JOIN** Sales S
  - => Cartesian product (*no join condition given!*)
- ... **FROM** Employees E **JOIN** Sales S **ON** E.EID=S.EID
  - => Equi-join



13

## 'FROM' clause – examples (cont)

- ... **FROM** Employees **NATURAL JOIN** Sales  
=> Natural join (*bug-prone, use equijoin instead*)
- ... **FROM** Employees E  
**LEFT JOIN** Sales S **ON** E.EID=S.EID  
=> Left join
- ... **FROM** Employees E1  
**JOIN** Employees E2 **ON** E1.EID < E2.EID  
=> Theta self-join (*what does it return?*)



14

## Gotcha: natural join in practice

- Uses *\*all\** same-named attributes
  - May be too many or too few
- Implicit nature reduces readability
  - Better to list explicitly all join conditions
- Fragile under schema changes
  - Nasty interaction of above two cases..

*Moral of the story: Avoid using Natural Join*



15

## Gotcha: join selectivity

- Consider tables **R**, **S**, **T** with **T=∅** and this query:

```
SELECT R.x
FROM R, S, T
WHERE R.x=S.x OR R.x=T.x
```

(*what does it return?*)

- Result contains no rows!
  - Selection (WHERE) operates on pre-joined tuples
  - $R \times S \times T = R \times S \times \emptyset = \emptyset$
  - => No tuples for WHERE clause to work with!
- Workaround?
  - Two coming up later

*Moral of the story: WHERE cannot create tuples*



16

## Explicit join ordering

- Use parentheses to group joins
  - e.g. (A join B) join (C join D)
- Special-purpose feature
  - Helps some (inferior) systems optimize better
  - Helps align schemas for natural join
- Recommendation: **avoid**
  - People are notoriously bad at optimizing things
  - Optimizer usually does what it wants anyway
  - ... but sometimes treats explicit ordering as a constraint



17

## 'WHERE' clause

- Conditions which all returned tuples must meet
  - Arbitrary boolean expression
  - Combine multiple expressions with AND/OR/NOT
- Attention to data of interest
  - Specific people, dates, places, quantities
  - Things which do (or do not) correlate with other data
- Often used instead of JOIN
  - FROM tables (Cartesian product, e.g. A, B)
  - Specify join condition in WHERE clause (e.g. A.ID=B.ID)
  - Optimizers (usually) understand and do the right thing



19

## 'WHERE' clause – examples

- ... **WHERE** S.date > '01-Jan-2010'  
=> Simple tuple-literal condition
- ... **WHERE** E.EID = S.EID  
=> Simple tuple-tuple condition (equi-join)
- ... **WHERE** E.EID = S.EID **AND** S.PID = P.PID  
=> Conjunctive tuple-tuple condition (three-way equijoin)
- ... **WHERE** S.value < 10 **OR** S.value > 10000  
=> Disjunctive tuple-literal condition



18

## Scalar expressions in SQL

- Literals, attributes, single-valued relations
- Boolean expressions
  - Boolean T/F coerce to 1/0 in arithmetic expressions
  - Zero/non-zero coerce to F/T in boolean expressions
- Logical connectors: AND, OR, NOT
- Conditionals
  - = != < > <= >= <>
  - BETWEEN, [NOT] LIKE, IS [NOT] NULL, ...
- Operators: + - \* / % & | ^
- Functions: math, string, date/time, etc. (more later)

*Similar to expressions in C, python, etc.*



20

## Pattern matching

- Compare a string to a pattern
  - <attribute> **LIKE** <pattern>
  - <attribute> **NOT LIKE** <pattern>
- Pattern is a quoted string
  - % => "any string"
  - \_ => "any character"
- To escape '%' or '\_':
  - LIKE '%x\_%' ESCAPE 'x' (replace 'x' with character of choice)
  - => matches strings containing '\_' (the underscore character)

*DBMS increasingly allow regular expressions*



## Pattern matching – examples

- ... **WHERE** phone **LIKE** '%268-\_\_\_\_\_'
  - phone numbers with exchange 268
  - WARNING: spaces are wrong, only shown for clarity
- ... **WHERE** last\_name **LIKE** 'Jo%'
  - Jobs, Jones, Johnson, Jorgensen, etc.
- ... **WHERE** Dictionary.entry **NOT LIKE** '%est'
  - Ignore 'biggest', 'tallest', 'fastest', 'rest', ...
- ... **WHERE** sales **LIKE** '%30!%%' **ESCAPE** '!'
  - Sales of 30%



## 'GROUP BY' clause

- Specifies **grouping key** of relational operator **T**
  - Comma-separated list of attributes (names or positions) which identify groups
  - Tuples agreeing in their grouping key are in same "group"
  - SELECT gives attributes to aggregate (and functions to use)
- SQL specifies several **aggregation functions**
  - COUNT, MIN, MAX, SUM, AVG, STD (standard deviation)
  - Some systems allow user-defined aggregates



## MORE COMPLEX QUERIES (GROUP BY-HAVING-ORDER BY)



## 'GROUP BY' clause – gotchas

- WHERE clause cannot reference aggregated values (sum, count, etc.)
  - Aggregates don't "exist yet" when WHERE runs
  - => Use **HAVING** clause instead (coming next)
- GROUP BY **must** list all non-aggregate attributes used in SELECT clause
  - Think projection
  - => Some systems do this implicitly, others throw error
- Grouping often (but not always!) sorts on grouping key
  - Depends on system and/or optimizer decisions
  - => Use **ORDER BY** to be sure (coming next)



25

## 'GROUP BY' clause – examples

- SELECT EID, SUM(value)  
FROM SALES **GROUP BY** EID  
– Show total sales for each employee ID
- SELECT EID, SUM(value), MAX(value)  
FROM Sales **GROUP BY** 1  
– Show total sales and largest sale for each employee ID
- SELECT EID, COUNT(EID)  
FROM Complaints **GROUP BY** EID  
– Show how many complaints each salesperson triggered



27

## Eliminating duplicates in aggregation

- Use **DISTINCT** inside an aggregation

```
SELECT EmpID, COUNT(DISTINCT CustID)
FROM CustomerComplaints
GROUP BY 1
```

=> Number of customers who complained about the employee



26

## 'GROUP BY' clause – examples (cont)

- SELECT EID, SUM(value) FROM Sales  
– Error: non-aggregate attribute (EID) missing from GROUP BY
- SELECT EID, value FROM Sales **GROUP BY** 1,2  
– Not an error – eliminates duplicates
- SELECT SUM(value) FROM Sales **GROUP BY** EID  
– Not an error, but rather useless: report per-employee sales anonymously
- SELECT SUM(value) FROM Sales  
– No GROUP BY => no grouping key => all tuples in same group




28

## 'HAVING' clause

- Allows predicates on aggregate values  
– Groups which do not match the predicate are eliminated  
=> **HAVING** is to **groups** what **WHERE** is to **tuples**
- Order of execution  
– WHERE is before GROUP BY  
=> Aggregates not yet available when WHERE clause runs  
– GROUP BY is before HAVING  
=> Scalar attributes still available

- In tree form:





29

## 'HAVING' clause – examples

- `SELECT EID, SUM(value)`  
`FROM Sales GROUP BY EID`  
**HAVING** `SUM(Sales.value) > 10000`  
– Highlight employees with “impressive” sales
- `SELECT EID, SUM(value)`  
`FROM Sales GROUP BY EID`  
**HAVING** `AVG(value) < (SELECT AVG(GroupAVG)`  
`FROM (SELECT EID, AVG(value) AS GroupAVG`  
`FROM Sales GROUP BY EID ) AS B);`  
– Highlight employees with below-average sales  
– Subquery to find the avg value of average employee sales



31

## 'ORDER BY' clause – examples

- **... ORDER BY** E.name  
=> Defaults to ascending order
- **... ORDER BY** E.name **ASC**  
=> Explicitly ascending order
- **... ORDER BY** E.name **DESC**  
=> Explicitly descending order
- **... ORDER BY** CarCount **DESC**, CarName **ASC**  
=> Matches our car example from previous lecture
- `SELECT E.name ... ORDER BY 1`  
=> Specify attribute's position instead of its name



30

## 'ORDER BY' clause

- Each query can sort by one or more attributes
  - Refer to attributes by name or position in SELECT
  - Ascending (default) or descending (reverse) order
  - Equivalent to relational operator  $\tau$
- Definition of 'sorted' depends on data type
  - Numbers use natural ordering
  - Date/time uses earlier-first ordering
  - NULL values are not comparable, cluster at end or beginning
- Strings are more complicated
  - Intuitively, sort in “alphabetical order”
  - Problem: which alphabet? case sensitive?
  - Answer: user-specified “collation order”
  - Default collation: case-sensitive latin (ASCII) alphabet

*String collation not covered in this class*



32

## What's next?

- Examples





33

## WORKING EXAMPLES

### Example: Simple SQL Query

Employee(FirstName, Surname, Dept, Office, Salary, City)  
Department(DeptName, Address, City)

"Find the salaries of employees named Brown"

```
SELECT Salary AS Remuneration
FROM Employee
WHERE Surname = 'Brown'
```

Result:

Remuneration
45
80



34

## Example Database

Employee(FirstName, Surname, Dept, Office, Salary, City)  
Department(DeptName, Address, City)

EMPLOYEE	FirstName	Surname	Dept	Office	Salary	City
	Mary	Brown	Administration	10	45	London
	Charles	White	Production	20	36	Toulouse
	Gus	Green	Administration	20	40	Oxford
	Jackson	Neri	Distribution	16	45	Dover
	Charles	Brown	Planning	14	80	London
	Laurence	Chen	Planning	7	73	Worthing
	Pauline	Bradshaw	Administration	75	40	Brighton
	Alice	Jackson	Production	20	46	Toulouse

Home city

DEPARTMENT	DeptName	Address	City
	Administration	Bond Street	London
	Production	Rue Victor Hugo	Toulouse
	Distribution	Pond Road	Brighton
	Planning	Bond Street	London
	Research	Sunset Street	San José

City of work



34

## Example: \* in the Target List

Employee(FirstName, Surname, Dept, Office, Salary, City)  
Department(DeptName, Address, City)

"Find all the information relating to employees named Brown" :

```
SELECT *
FROM Employee
WHERE Surname = 'Brown'
```

Result:

FirstName	Surname	Dept	Office	Salary	City
Mary	Brown	Administration	10	45	London
Charles	Brown	Planning	14	80	London



## Example: Attribute Expressions

Employee(FirstName, Surname, Dept, Office, Salary, City)  
Department(DeptName, Address, City)

"Find the monthly salary of employees named White" :

```
SELECT Salary / 12 AS MonthlySalary
FROM Employee
WHERE Surname = 'White'
```

Result:

MonthlySalary
3.00



## Example: Simple (Equi-)Join Query

Employee(FirstName, Surname, Dept, Office, Salary, City)  
Department(DeptName, Address, City)

"Find the names of employees and their cities of work"

```
SELECT Employee.FirstName, Employee.Surname, Department.City
FROM Employee, Department
WHERE Employee.Dept = Department.DeptName
```

Result:

FirstName	Surname	City
Mary	Brown	London
Charles	White	Toulouse
Gus	Green	London
Jackson	Neri	Brighton
Charles	Brown	London
Laurence	Chen	London
Pauline	Bradshaw	London
Alice	Jackson	Toulouse

(alternative?)

Alternative (and more correct):

```
SELECT Employee.FirstName, Employee.Surname, Department.City
FROM Employee E JOIN Department D ON E.Dept = D.DeptName
```



## Example: Table Aliases

Employee(FirstName, Surname, Dept, Office, Salary, City)  
Department(DeptName, Address, City)

"Find the names of employees and their cities of work"  
(using an alias):

```
SELECT FirstName, Surname, D.City
FROM Employee, Department D
WHERE Dept = DeptName
```

Result:

FirstName	Surname	City
Mary	Brown	London
Charles	White	Toulouse
Gus	Green	London
Jackson	Neri	Brighton
Charles	Brown	London
Laurence	Chen	London
Pauline	Bradshaw	London
Alice	Jackson	Toulouse



## Example: Predicate Conjunction

Employee(FirstName, Surname, Dept, Office, Salary, City)  
Department(DeptName, Address, City)

"Find the first names and surnames of employees who  
work in office number 20 of the Administration  
department":

```
SELECT FirstName, Surname
FROM Employee
WHERE Office = '20' AND Dept = 'Administration'
```

Result:

FirstName	Surname
Gus	Green



## Example: Predicate Disjunction

Employee(FirstName, Surname, Dept, Office, Salary, City)  
Department(DepartmentName, Address, City)

"Find the first names and surnames of employees who work in either the Administration or the Production department":

```
SELECT FirstName, Surname
FROM Employee
WHERE Dept = 'Administration' OR Dept = 'Production'
```

Result:

FirstName	Surname
Mary	Brown
Charles	White
Gus	Green
Pauline	Bradshaw
Alice	Jackson



## Example: String Matching Operator LIKE

Employee(FirstName, Surname, Dept, Office, Salary, City)  
Department(DepartmentName, Address, City)

"Find employees with surnames that have 'r' as the second letter and end in 'n'":

```
SELECT *
FROM Employee
WHERE Surname LIKE '_r%n'
```

Result:

FirstName	Surname	Dept	Office	Salary	City
Mary	Brown	Administration	10	45	London
Gus	Green	Administration	20	40	Oxford
Charles	Brown	Planning	14	80	London



## Example: Complex Logical Expressions

Employee(FirstName, Surname, Dept, Office, Salary, City)  
Department(DepartmentName, Address, City)

"Find the first names of employees named Brown who work in the Administration department or the Production department":

```
SELECT FirstName
FROM Employee
WHERE Surname = 'Brown' AND (Dept = 'Administration' OR Dept = 'Production')
```

Result:

FirstName
Mary



## Example: Aggregate Queries: Operator Count

Employee(FirstName, Surname, Dept, Office, Salary, City)  
Department(DepartmentName, Address, City)

"Find the number of employees":

```
SELECT count(*) FROM Employee
```

"Find the number of different values on attribute Salary for all tuples in Employee":

```
SELECT count(DISTINCT Salary) FROM Employee
```

"Find the number of tuples in Employee having non-null values on the attribute Salary":

```
SELECT count(ALL Salary) FROM Employee
```



## Example: Operators **Sum**, **Avg**, **Max** and **Min**

Employee(FirstName, Surname, Dept, Office, Salary, City)  
Department(DeptName, Address, City)

"Find the sum of all salaries for the Administration department":

```
SELECT sum(Salary) AS SumSalary
FROM Employee
WHERE Dept = 'Administration'
```

Result:

SumSalary
125



## Example: Operators **Sum**, **Avg**, **Max** and **Min**

Employee(FirstName, Surname, Dept, Office, Salary, City)  
Department(DeptName, Address, City)

"Find the maximum and minimum salaries among all employees":

```
SELECT max(Salary) AS MaxSal, min(Salary) AS MinSal
FROM Employee
```

Result:

MaxSal	MinSal
80	36



## Example: Aggregate Operators with Join

Employee(FirstName, Surname, Dept, Office, Salary, City)  
Department(DeptName, Address, City)

"Find the maximum salary among the employees who work in a department based in London":

```
SELECT max(Salary) AS MaxLondonSal
FROM Employee, Department
WHERE Dept = DeptName AND Department.City = 'London'
```

Result:

MaxLondonSal
80