

Data Definition Language (DDL) Data Manipulation Language (DML) Views & Indexes

Introduction to Databases

Sina Meraji

Thanks to Ryan Johnson, John Mylopoulos, Arnold Rosenbloom
and Renee Miller for material in these slides

DATA DEFINITION LANGUAGE (DDL)

SQL Main Components

- **Queries**
 - Subset of SQL for read-only access to database
 - SELECT statements
- **Data Definition Language (DDL)**
 - Subset of SQL used to describe database schemas
 - CREATE, ALTER, DROP statements
 - Data types
 - Integrity constraints
- **Data Manipulation Language (DML)**
 - Subset of SQL used to manipulate data in databases
 - INSERT, DELETE, UPDATE statements

Creating (Declaring) a Schema

- A schema is essentially a namespace
 - it contains named objects (tables, data types, functions, etc.)
- The schema name must be distinct from any existing schema name in the current database
- Syntax:
CREATE SCHEMA schemaname [AUTHORIZATION username]
[schema_element [...]]

Examples:

CREATE SCHEMA myschema;

CREATE SCHEMA myschema **AUTHORIZATION** sina;



Creating (Declaring) a Relation/Table

- To create a relation:

```
CREATE TABLE <name> (  
    <list of elements>  
);
```

- To delete a relation:

```
DROP TABLE <name>;
```

- To alter a relation (add/remove column):

```
ALTER TABLE <name> ADD <element>  
ALTER TABLE <name> DROP <element>
```



Elements of Table Declarations

- Elements:
 - attributes and their type
 - constraints (see later)
- The most common types are:
 - INT or INTEGER (synonyms)
 - REAL or FLOAT (synonyms)
 - CHAR(*n*) = fixed-length string of *n* characters
 - VARCHAR(*n*) = variable-length string of up to *n* characters



Examples

- To create a relation:

```
CREATE TABLE employees (  
    id            INTEGER,  
    first_name    CHAR(50),  
    last_name     VARCHAR(100));
```

- To delete a relation:

```
DROP TABLE employees;
```

- To alter a relation (add/remove column):

```
ALTER TABLE employees ADD age INTEGER;  
ALTER TABLE employess DROP last_name;
```



SQL Values

- Integers and reals are represented as you would expect
- Strings are too, except they require single quotes.
 - Two single quotes = real quote, e.g., 'Joe''s Bar'
- Any value can be NULL
 - Unless attribute has NOT NULL constraint
E.g.: price REAL **NOT NULL**



Dates and Times

- DATE and TIME are types in SQL.
 - The form of a **date value** is: DATE 'yyyy-mm-dd'

Example (for Oct. 19, 2011):

DATE '2011-10-19'

- The form of a **time value** is: TIME 'hh:mm:ss' with an optional decimal point and fractions of a second following.
- Example (for two and a half seconds after 6:40PM):

TIME '18:40:02.5'



Running Example

Beers(name, manf)
 Bars(name, addr, license)
 Drinkers(name, addr, phone)
 Likes(drinker, beer)
 Sells(bar, beer, price)
 Frequents(drinker, bar)

Underline = **key** (tuples cannot have the same value in all key attributes)

- Excellent example of a constraint



INTEGRITY CONSTRAINTS

Kinds of Constraints

- **Keys**
- **Foreign-key** or referential-integrity constraints
 - Inter-relation constraints
- **Value-based** constraints
 - Constrain values of a particular attribute
- **Tuple-based** constraints
 - Relationship among components
- **Assertions**





13

Declaring Keys

- An attribute or list of attributes may be declared **PRIMARY KEY** or **UNIQUE**
 - Either says that no two tuples of the relation may agree in all the attribute(s) on the list
 - There are a few distinctions to be mentioned later
 - Place **PRIMARY KEY** or **UNIQUE** after the type in the declaration of the attribute.

```
CREATE TABLE Beers (
    name CHAR(20) PRIMARY KEY,
    manf CHAR(20)
);
```



15

PRIMARY KEY vs. UNIQUE

- There can be only one **PRIMARY KEY** for a relation, but several **UNIQUE** attributes
- No attribute of a **PRIMARY KEY** can ever be NULL in any tuple. But attributes declared **UNIQUE** may have NULL's, and there may be several tuples with NULL.



14

Declaring Multi-attribute Keys

- A key declaration can appear as element in the list of elements of a CREATE TABLE statement
- This form is essential if the key consists of more than one attribute

```
CREATE TABLE Sells (
    bar CHAR(20),
    beer VARCHAR(20),
    price REAL,
    PRIMARY KEY (bar, beer));
```

The bar and beer together are the key for Sells



16

Foreign Keys

- Values appearing in attributes of one relation must appear together in certain attributes of another relation

Example:

We might expect that a value in **Sells.beer** also appears as value in **Beers.name**

Beers(name, manf)

Sells(bar, beer, price)



Expressing Foreign Keys

- Use keyword **REFERENCES**, either:
 - After an attribute (for one-attribute keys)
REFERENCES <relation> (<attributes>)
 - As an element of the schema:
FOREIGN KEY (<list of attributes>)
REFERENCES <relation> (<attributes>)
- Referenced attributes must be declared **PRIMARY KEY** or **UNIQUE**



Example: As Schema Element

```
CREATE TABLE Beers (
  name      CHAR(20) PRIMARY KEY,
  manf      CHAR(20) );
```

```
CREATE TABLE Sells (
  bar       CHAR(20),
  beer      CHAR(20),
  price     REAL,
  FOREIGN KEY(beer) REFERENCES Beers(name));
```



Example: With Attribute

```
CREATE TABLE Beers (
  name      CHAR(20) PRIMARY KEY,
  manf      CHAR(20) );
```

```
CREATE TABLE Sells (
  bar       CHAR(20),
  beer      CHAR(20) REFERENCES Beers(name),
  price     REAL );
```



Enforcing Foreign-Key Constraints

- If there is a foreign-key constraint from relation **R** to relation **S**, two violations are possible:
 - An insert or update to **R** introduces values not found in **S**
 - A deletion or update to **S** causes some tuples of **R** to “dangle”

Example: suppose **R = Sells**, **S = Beers**

- An insert or update to **Sells** that introduces a non-existent beer must be rejected
- A deletion or update to **Beers** that removes a beer value found in some tuples of **Sells** can be handled in three ways (next slide).



21

Actions Taken

- **DEFAULT:** Reject the modification
 - Deleted beer in **Beer**: reject modifications in **Sells** tuples
 - Updated beer in **Beer**: reject modifications in **Sells** tuples
- **CASCADE:** Make the same changes in **Sells**
 - Deleted beer in **Beer**: delete **Sells** tuple
 - Updated beer in **Beer**: change value in **Sells**
- **SET NULL:** Change the beer to NULL
 - Deleted beer in **Beer**: set NULL values in **Sells** tuples
 - Updated beer in **Beer**: set NULL values in **Sells** tuples



24

Choosing a Policy

- When we declare a foreign key, we may choose policies **SET NULL** or **CASCADE** independently for deletions and updates
- Follow the foreign-key declaration by:
ON [UPDATE, DELETE][SET NULL, CASCADE]
- Two such clauses may be used, otherwise, the default (reject) is used.



23

Example

- Delete the 'Bud' tuple from Beers
 - **DEFAULT:** do not change any tuple from Sells that have beer = 'Bud'
 - **CASCADE:** delete all tuples from Sells that have beer = 'Bud'
 - **SET NULL:** Change all tuples of Sells that have beer = 'Bud' to have beer = NULL
- Update the 'Bud' tuple to 'Budweiser'
 - **DEFAULT:** do not change any tuple from Sells that have beer = 'Bud'
 - **CASCADE:** change all Sells tuples with beer = 'Bud' to beer = 'Budweiser'
 - **SET NULL:** Same change as for deletions



25

Example: Setting a Policy

```
CREATE TABLE Sells (  
    bar      CHAR(20),  
    beer     CHAR(20),  
    price    REAL,  
    FOREIGN KEY(beer)  
        REFERENCES Beers(name)  
        ON DELETE SET NULL  
        ON UPDATE CASCADE  
);
```



26

Attribute-Based Checks

- Constraints on the value of a particular attribute
- Add **CHECK**(<condition>) to the declaration for the attribute
- The condition may use the name of the attribute, but **any other relation or attribute name must be in a subquery**



28

Timing of Checks

- Attribute-based checks are performed **only when** a value for that attribute is inserted or updated
- Example:
 - **CHECK** (price <= 5.00)
Checks every new price and rejects the modification (for that tuple) if the price is more than \$5
 - **CHECK** (beer **IN** (**SELECT** name **FROM** Beers))
Not checked if a beer is later deleted from Beers (unlike foreign-keys)



27

Example: Attribute-based Check

```
CREATE TABLE Sells (
  bar      CHAR(20),
  beer     CHAR(20) CHECK (beer IN(SELECT name FROM Beers)),
  price    REAL CHECK (price <= 5.00 )
);
```



29

Tuple-Based Checks

- **CHECK** (<condition>) may be added as a relation-schema element
 - The condition may refer to any attribute of the relation, but other attributes or relations require a subquery
 - Checked on insert or update only

Example: Only Joe's Bar can sell beer for more than \$5:

```
CREATE TABLE Sells (
  bar      CHAR(20),
  beer     CHAR(20),
  price    REAL,
  CHECK    (bar = 'Joe's Bar' OR price <= 5.00));
```



Assertions

- Permit the definition of constraints over whole tables, rather than individual tuples
 - useful in many situations -- e.g., to express generic inter-relational constraints
 - An assertion associates a name to a check clause. Syntax:

CREATE ASSERTION AssertName **CHECK** (Condition)

Example:

"There must always be at least one tuple in table Employee":

```
CREATE ASSERTION AlwaysOneEmployee  
CHECK (1 <= (SELECT count(*) FROM Employee))
```



DATA MANIPULATION LANGUAGE (DML)



Enforcement Policies

- Integrity constraints (checks, assertions) may be checked immediately when a change takes place to a relation, or at the end of a **transaction**
 - The first case may result in a **partial rollback**
 - the latter in a **(full) rollback**.
- This topic is discussed in more detail in **CSC443**



Data Manipulation Language (DML)

- Syntax elements used for inserting, deleting and updating data in a database
- Modification statements include:
 - **INSERT** - for inserting data in a database
 - **DELETE** - for deleting data in a database
 - **UPDATE** - for updating data in a database
- All modification statements operate on **a set** of tuples (no duplicates)



Example

Employee(FirstName,Surname,Dept,Office,Salary,City)

Department(DeptName, Address, City)

Product(Code,Name,Description,ProdArea)

LondonProduct(Code,Name,Description)



Notes on Insertions

- The ordering of attributes (if present) and of values is meaningful -- first value for the first attribute, etc.
- If *AttributeList* is omitted, all the relation attributes are considered, in the order they appear in the table definition
- If *AttributeList* does not contain all the relation attributes, left-out attributes are assigned default values (if defined) or the NULL value



Insertions

Syntax varies:

- Using only *values*:

```
INSERT INTO Department VALUES ('Production', 'Rue du Louvre 23', 'Toulouse')
```

- Using both *column names* and *values*:

```
INSERT INTO Department(DeptName, City)
VALUES ('Production', 'Toulouse')
```

- Using a *subquery*:

```
INSERT INTO LondonProducts
(SELECT Code, Name, Description FROM Product WHERE ProdArea = 'London')
```



Deletions

Syntax:

```
DELETE FROM TableName [WHERE Condition]
```

- "Remove the Production department":

```
DELETE FROM Department
WHERE DeptName = 'Production'
```

- "Remove departments with no employees":

```
DELETE FROM Department
WHERE DeptName NOT IN (SELECT Dept FROM Employee)
```



Notes on Deletions

- The DELETE statement removes from a table all tuples that satisfy a condition
- If the WHERE clause is omitted, DELETE removes all tuples from the table (keeps the table schema):

DELETE FROM Department

- The removal may produce deletions from other tables – (see referential integrity constraint with **cascade** policy later)
- To remove table **Department** completely (content and schema) :

DROP TABLE Department **CASCADE**



VIEWS



Updates

Syntax:

UPDATE *TableName*

SET *Attribute* = < *Expression* | *SelectSQL* | null | default >
 {, *Attribute* = < *Expression* | *SelectSQL* | null | default >}
 [**WHERE** *Condition*]

Examples:

UPDATE Employee

SET Salary = Salary + 5

WHERE RegNo = 'M2047'

UPDATE Employee

SET Salary = Salary * 1.1

WHERE Dept = 'Administration'



Views

- A **view** is a relation defined in terms of stored tables (called **base tables**) and other views.
- Two kinds:
 - **Virtual** = not stored in the database; just a query for constructing the relation

CREATE VIEW <name> **AS** <query>;
 - **Materialized** = actually constructed and stored

CREATE MATERIALIZED VIEW <name> **AS** <query>;



Running Example

Beers(name, manf)
Bars(name, addr, license)
Drinkers(name, addr, phone)
Likes(drinker, beer)
Sells(bar, beer, price)
Frequents(drinker, bar)

42



Example: View Definition

CanDrink(drinker, beer) is a view “containing” the drinker-beer pairs such that the drinker frequents at least one bar that serves the beer:

```
CREATE VIEW CanDrink AS  
  SELECT drinker, beer  
  FROM Frequents, Sells  
  WHERE Frequents.bar = Sells.bar;
```



Example: Accessing a View

Query a view as if it were a base table:

```
SELECT beer  
FROM CanDrink  
WHERE drinker = 'Sally';
```

44



Notes on Views

- Data independence (hide schema from apps)
 - DB team splits CustomerInfo into Customer and Address
 - View accommodate changes with web apps
- Data hiding (access data on need-to-know basis)
 - Doctor outsources patient billing to third party
 - View restricts access to billing-related patient info
- Code reuse
 - Very similar subquery appears multiple times in a query
 - View shortens code, improves readability, reduces bugs, ...
 - Bonus: query optimizer often does a better job!

45

Example: Views and Queries

Employee(RegNo,FirstName,Surname,Dept,Office,Salary,City)
Department(DeptName,Address,City)

"Find the department with highest salary expenditures" (**without using a view**):

```
SELECT Dept
FROM Employee
GROUP BY Dept
HAVING sum(Salary) >= ALL (
    SELECT sum(Salary) FROM Employee GROUP BY Dept)
```

Updates on Views

- Generally, it is impossible to modify a virtual view because it doesn't exist
- Can't we "translate" updates on views into "equivalent" updates on base tables?
 - Not always (in fact, not often)
 - Most systems prohibit most view updates

Example: Views and Queries (cont.)

"Find the department with highest salary expenditures" (**using a view**):

```
CREATE VIEW SalBudget (Dept, SalTotal) AS
SELECT Dept, sum(Salary)
FROM Employee
GROUP BY Dept

SELECT Dept
FROM SalBudget
WHERE SalTotal = (SELECT max(SalTotal) FROM SalBudget)
```

Example:The View

```
CREATE VIEW Synergy AS
SELECT Likes.drinker, Likes.beer, Sells.bar
FROM Likes, Sells, Frequents
WHERE Likes.drinker = Frequents.drinker
AND Likes.beer = Sells.beer
AND Sells.bar = Frequents.bar;
```

Pick one copy of
each attribute

Join of Likes,
Sells, and Frequents



50

Interpreting a View Insertion

- We cannot insert into Synergy - it is a virtual view
- **Idea:** Try to translate a (drinker, beer, bar) triple into three insertions of projected pairs, one for each of Likes, Sells, and Frequent.
 - Sells.price will have to be NULL.
 - There isn't always a unique translation

Need for **SQL Triggers** - Not discussed



52

Example: A Data Warehouse

- Wal-Mart stores every sale at every store in a database
- Overnight, the sales for the day are used to update a data warehouse = **materialized views** of the sales
- The warehouse is used by analysts to **predict trends** and move goods to where they are selling best



51

Materialized Views

- **Problem:** each time a base table changes, the materialized view may change
 - Cannot afford to recompute the view with each change
- **Solution:** Periodic reconstruction of the materialized view, which is otherwise “out of date”



53

INDEXES (INDICES)



54

Index

- **Problem:** needle in haystack
 - Find all phone numbers with first name 'Mary'
 - Find all phone numbers with last name 'Li'
- **Index:** auxiliary database structure which provides random access to data
 - Index a set of attributes. No standard syntax! Typical is:
`CREATE INDEX indexName ON TableName(AttributeList);`
 - Random access to any indexed attribute
(e.g., retrieve a single tuple out of billions in <5 disk accesses)
 - Similar to a hash table, but in a DBMS it is a **balanced search tree** with giant nodes (a full disk page) called a **B-tree**



56

Another Example: Using Index

```
CREATE INDEX BeerInd ON Beers(manf);  
CREATE INDEX SellInd ON Sells(bar, beer);
```

Query: Find the prices of beers manufactured by Pete's and sold by Joe's bar

```
SELECT price FROM Beers, Sells  
WHERE manf = 'Pete's' AND Beers.name = Sells.beer  
AND bar = 'Joe's Bar';
```

DBMS uses:

- **BeerInd** to get all the beers made by Pete's **fast**
- **SellInd** to get prices of those beers, with bar = 'Joe's Bar' **fast**



55

Example: Using Index

```
SELECT fname  
FROM people  
WHERE lname = 'Sina'
```

- **Without an index:**
The DBMS must look at the *lname* column on every row in the table (this is known as a **full table scan**)
- **With an index** (defined on attribute lname):
The DBMS simply follows the **B-tree** data structure until the 'Sina' entry has been found

This is much less computationally expensive than a full table scan



57

Database Tuning

- How to make a database run fast?
 - Decide which indexes to create
- **Pro:** An index speeds up queries that can use it
- **Con:** An index slows down all modifications on its relation as the index must be modified too



58

Example: Database Tuning

- Suppose the only things we did with our beers database was:
 - Insert new beers into a relation (10%).
 - Find the price of a given beer at a given bar (90%).
- Then
 - **SellInd** on Sells(bar, beer) would be wonderful
 - **BeerInd** on Beers(manf) would be harmful

Make common case fast



60

What's Next?

- Embedded SQL
 - Part of Assignment 2



59

Tuning Advisors

- A major research thrust
 - Because hand tuning is so hard
- An advisor gets a query load, e.g.:
 - Choose random queries from the history of queries run, or
 - Designer provides a sample workload
- The advisor generates candidate indexes and evaluates each on the workload
 - Feed each sample query to the query optimizer, which assumes only this one index is available
 - Measure the improvement/degradation in the average running time of the queries.