

Relational Algebra

csc343, Introduction to Databases

Nosayba El-Sayed (based on slides from Diane Horton)

Fall 2015




UNIVERSITY OF
TORONTO

DCS50

Announcements – Week 2

- If you're missing pre-requisites for CSC343 and we didn't discuss it yet → email me (if you emailed, I'll get back to you)
- Assignment #1 will be posted next week
(Relational Algebra queries)
 - Don't worry if you can't login to MarkUs yet, your accounts will be added soon
 - Plan your partnership!
- Today: introducing Relational Algebra + simple exercises
- Next week: more advanced RA examples ;-)

Questions from last week



- Is every Key also a Superkey, by definition?
 - Yes – it's just a special type of superkeys (minimal)
- Movies Schema: “Is there a limit to the # of directors a movie can have?”
 - Can't we just add a new tuple for the same movie, only with different director names?
 - Problem: from the database system's point of view, the only way of identifying a unique movie is using its mID field.

Movies(mID, title, director, year, length)

mID	title	director	year	length
1	Shining	Kubrick	1980	146
2	Player	Altman	1992	146
3	Chinatown	Polanski	1974	131
4	Repulsion	Polanski	1965	143
5	Star Wars IV	Lucas	1977	126
6	American Graffiti	Lucas	1973	110
7	Full Metal Jacket	Kubrick	1987	156

8		Star Wars IV		Jack		1977		126
---	--	--------------	--	------	--	------	--	-----

RA Basics

(covered by your week 2 Prep)

Elementary Algebra

- You did algebra in high school
 - $27y^2 + 8y - 3$
- Operands?
- Operators?

Relational Algebra

- Operands? *tables*
- Operators?
 - choose only the *rows* you want
 - choose only the *columns* you want
 - *combine* tables
 - and a few other things..

A schema for our examples

Movies(mID, title, director, year, length)

Artists(aID, aName, nationality)

Roles(mID, aID, character)

Foreign key constraints:

- Roles[mID] \subseteq Movies[mID]
- Roles[aID] \subseteq Artists[aID]

Select: choose rows

- Notation: $\sigma_c(R)$
 - R is a table.
 - Condition c is a boolean expression.
 - ✓ It can use comparison operators and boolean operators
 - ✓ The operands are either *constants* or *attributes* of R .
- The result is a relation
 - with the *same schema* as the operand
 - but with only the *tuples* that satisfy the *condition*

Exercise

- Write queries to find:
 - All British actors?
 - All movies from the 1970s?

Movies Schema

Movies(mID, title, director, year, length)
Artists(aID, aName, nationality)
Roles(mID, aID, character)

Roles[mID] \subseteq Movies[mID]
Roles[aID] \subseteq Artists[aID]

$\sigma_c(R)$

All British actors?

All movies from the 1970s?

Movies:

mID	title	director	year	length
1	Shining	Kubrick	1980	146
2	Player	Altman	1992	146
3	Chinatown	Polanski	1974	131
4	Repulsion	Polanski	1965	143
5	Star Wars IV	Lucas	1977	126
6	American Graffiti	Lucas	1973	110
7	Full Metal Jacket	Kubrick	1987	156

Artists:

aID	aName	nat
1	Nicholson	American
2	Ford	American
3	Stone	British
4	Fisher	American

Roles:

mID	aID	character
1	1	Jack Torrance
3	1	Jake 'J.J.' Gittes
1	3	Delbert Grady
5	2	Han Solo
6	2	Bob Falfa
5	4	Princess Leia Organa

Exercise

- Write queries to find:

- All British actors?

$\sigma_{\text{nat}=\text{"British"}}(\text{Artists})$

- All movies from the 1970s?

$\sigma_{\text{year} > 1969 \wedge \text{year} < 1980}(\text{Movies})$

$\sigma_{\text{nat}=\text{"British"}}(\text{Artists})$

aID	aName	nat
3	Stone	British

- What if we only want the **names** of all British actors?
We need a way to pare down the columns.

Movies Schema

Movies(mID, title, director, year, length)
Artists(aID, aName, nationality)
Roles(mID, aID, character)

Roles[mID] \subseteq Movies[mID]
Roles[aID] \subseteq Artists[aID]

Movies:

mID	title	director	year	length
1	Shining	Kubrick	1980	146
2	Player	Altman	1992	146
3	Chinatown	Polanski	1974	131
4	Repulsion	Polanski	1965	143
5	Star Wars IV	Lucas	1977	126
6	American Graffiti	Lucas	1973	110
7	Full Metal Jacket	Kubrick	1987	156

Artists:

aID	aName	nat
1	Nicholson	American
2	Ford	American
3	Stone	British
4	Fisher	American


Roles:

mID	aID	character
1	1	Jack Torrance
3	1	Jake 'J.J.' Gittes
1	3	Delbert Grady
5	2	Han Solo
6	2	Bob Falfa
5	4	Princess Leia Organa

Project: choose columns

- Notation: $\pi_L(R)$
 - R is a table.
 - L is a subset of the attributes of R .
- The result is a relation
 - with all the tuples from R
 - but with only the **attributes** in L , and in that order

About project π



- Why is it called “project”?
- What is the value of $\pi_{\text{director}}(\text{Movies})$?
 - (a) { Kubrick, Altman, Polanski, Polanski, Lucas, Lucas, Kubrick }, OR

(b) { Kubrick, Altman, Polanski, Lucas }

← *project removes duplicates*
(think sets)


- Exercise: Write an RA expression to find the names of all directors of movies from the 1970s

$\pi_{\text{director}} \sigma_{\text{year} > 1969 \wedge \text{year} < 1980}(\text{Movies})$

- Now, suppose you want the names of all **characters** in movies from the 1970s..
- We need to be able to **combine** tables.

Cartesian Product

- Notation: $R1 \times R2$
- The result is a **relation** with
 - every combination of a tuple from $R1$ concatenated to a tuple from $R2$
 - Its schema is every attribute from $R1$ followed by every attribute of $R2$, in order
- How many tuples are in $R1 \times R2$?
 - $|R1| \times |R2|$
- Example: Movies x Roles
- If an **attribute** occurs in both relations, it occurs twice in the result (prefixed by relation name)
 - E.g.: **Movies.mID**, **Roles.mID**



$|R|$ = cardinality of R

Example of Cartesian product

profiles:

twitterID	name
hford	Harrison Ford
gclooney	George Clooney

follows:

aID	bID
hford	gclooney
gclooney	amal
marissamayer	lpage

profiles X follows:

twitterID	name	aID	bID
hford	Harrison Ford	hford	gclooney
hford	Harrison Ford	gclooney	amal
hford	Harrison Ford	marissamayer	lpage
gclooney	George Clooney	hford	gclooney
gclooney	George Clooney	gclooney	amal
gclooney	George Clooney	marissamayer	lpage



Cartesian product can be inconvenient

- It can introduce **nonsense** tuples.
- You can get rid of them with **selects**.
- But this is so highly common, an operation was defined to make it easier..
 - **Joins**..

Just like you did
in the lecture prep
PCRS exercise ;-)

Just like you tried
to do but PCRS
didn't accept ;-(



? question ★

Do `natural_join` and `theta_join` work?

When doing number three, or going back and joining unnecessary tal of these functions fail. Can any instructor please confirm/deny this?

Joining two relations

Natural Join

- Notation: $R \bowtie S$
- The result is formed by
 1. taking the **Cartesian product**
 2. **select** to ensure equality on attributes that are in both relations (determined *by name*)
 3. **projecting** to remove duplicate attributes.
- Example:
Artists \bowtie **Roles** gets rid of the nonsense tuples.

Natural Join Examples

- The following examples show what **natural join** does when the tables have:
 - no attributes in common
 - one attribute in common
 - a different attribute in common
- (Note that we change the attribute names for relation **follows** to set up these scenarios.)

profiles:

twitterID	name
hford	Harrison Ford
gclooney	George Clooney

follows:

aID	bID
hford	gclooney
gclooney	amal
marissamayer	lpage

profiles X follows:

twitterID	name	aID	bID
hford	Harrison Ford	hford	gclooney
hford	Harrison Ford	gclooney	amal
hford	Harrison Ford	marissamayer	lpage
gclooney	George Clooney	hford	gclooney
gclooney	George Clooney	gclooney	amal
gclooney	George Clooney	marissamayer	lpage

profiles:

twitterID	name
hford	Harrison Ford
gclooney	George Clooney

follows:

twitterID	bID
hford	gclooney
gclooney	amal
marissamay	lpage

*redundant attribute
is omitted from result*

profiles ⋈ follows?

twitterID	name	twitterID	bID
hford	Harrison Ford	hford	gclooney
hford	Harrison Ford	gclooney	amal
hford	Harrison Ford	marissamayer	lpage
gclooney	George Clooney	hford	gclooney
gclooney	George Clooney	gclooney	amal
gclooney	George Clooney	marissamayer	lpage

profiles:

twitterID	name
hford	Harrison Ford
gclooney	George Clooney

follows:

alD	twitterID
hford	gclooney
gclooney	amal
marissamayer	lpage

profiles ✕ follows:

twitterID	name	alD	twitterID
hford	Harrison Ford	hford	gclooney
hford	Harrison Ford	gclooney	amal
hford	Harrison Ford	marissamayer	lpage
gclooney	George Clooney	hford	gclooney
gclooney	George Clooney	gclooney	amal
gclooney	George Clooney	marissamayer	lpage

Properties of Natural Join

- Commutative:

$$R \bowtie S = S \bowtie R$$

(although attribute order may vary)

- Associative:

$$R \bowtie (S \bowtie T) = (R \bowtie S) \bowtie T$$

- So when writing n-ary joins, brackets are irrelevant.
We can just write:

$$R_1 \bowtie R_2 \bowtie \dots \bowtie R_n$$

Questions

1. How many tuples are in Artists \times Roles?

– 24

2. How many tuples are in Artists \bowtie Roles?

– 6

3. What is the result of:

$\Pi_{aName} \sigma_{director="Kubrick"} (Artists \bowtie Roles \bowtie Movies)$

4. What is the result of:

$\Pi_{aName} ((\sigma_{director="Kubrick"} Artists) \bowtie Roles \bowtie Movies)$

Movies Schema

Movies(mID, title, director, year, length)
Artists(aID, aName, nationality)
Roles(mID, aID, character)

Roles[mID] \subseteq Movies[mID]
Roles[aID] \subseteq Artists[aID]

Movies:

mID	title	director	year	length
1	Shining	Kubrick	1980	146
2	Player	Altman	1992	146
3	Chinatown	Polanski	1974	131
4	Repulsion	Polanski	1965	143
5	Star Wars IV	Lucas	1977	126
6	American Graffiti	Lucas	1973	110
7	Full Metal Jacket	Kubrick	1987	156

Artists:

aID	aName	nat
1	Nicholson	American
2	Ford	American
3	Stone	British
4	Fisher	American

Roles:

mID	aID	character
1	1	Jack Torrance
3	1	Jake 'J.J.' Gittes
1	3	Delbert Grady
5	2	Han Solo
6	2	Bob Falfa
5	4	Princess Leia Organa

(Artists ⋈ Roles ⋈ Movies)

Artists. aID	Artists. aName	Artists. nat	Roles. MID	Roles. character	Movies. title	Movies. director	Movies. year	Movies. length
1	Nicholson	US	1	Jack Torrance	Shining	Kubrick	1980	146
1	Nicholson	US	3	Jack Gittes	Chinatown	Polanski	1974	131
2	Ford	US	5	Han Solo	Star Wars IV	Lucas	1977	126
2	Ford	US	6	Bob Falfa	American Graffiti	Lucas	1973	110
3	Stone	UK	1	Delbert Grady	Shining	Kubrick	1980	146
4	Fisher	US	5	Process Organa	Star Wars IV	Lucas	1977	126

$\Pi_{aName} \sigma_{\text{director}="Kubrick"} (\text{Artists} \bowtie \text{Roles} \bowtie \text{Movies})$?

aName


Nicholson

Stone

(Artists ⋈ Roles ⋈ Movies)

Artists. aID	Artists. aName	Artists. nat	Roles. MID	Roles. character	Movies. title	Movies. director	Movies. year	Movies. length
1	Nicholson	US	1	Jack Torrance	Shining	Kubrick	1980	146
1	Nicholson	US	3	Jack Gittes	Chinatown	Polanski	1974	131
2	Ford	US	5	Han Solo	Star Wars IV	Lucas	1977	126
2	Ford	US	6	Bob Falfa	American Graffiti	Lucas	1973	110
3	Stone	UK	1	Delbert Grady	Shining	Kubrick	1980	146
4	Fisher	US	5	Process Organa	Star Wars IV	Lucas	1977	126

$\Pi_{aName} ((\sigma_{director="Kubrick"} Artists) \bowtie Roles \bowtie Movies) \quad ??$



Special cases for natural join

Imagine this scenario

Result of natural join

Artist	Name
9132	William Shatner
8762	Harrison Ford
5555	Patrick Stewart
1868	Angelina Jolie

Artist	Name
5555	Patrick Stewart
1868	Angelina Jolie

Exactly
the same
attributes

Artist	Name
1234	Brad Pitt
1868	Angelina Jolie
5555	Patrick Stewart

What about this one?

Artist	Name
1234	Brad Pitt
1868	Angelina Jolie
5555	Patrick Stewart

No attributes
in common

mID	Title	Director	Year	Length
1111	Alien	Scott	1979	152
1234	Sting	Hill	1973	130

Result? Cartesian product of the two relations!

Natural join can “over-match”

- Natural join bases the matching on **attribute names**.
- What if two attributes have the same name, but we don't want them to have to match?
- Example: if **Artists** used “name” for *actors'* names and **Movies** used “name” for *movies'* names.
 - Can rename one of them (we'll see how).
 - Or?
...use a Cartesian product + select ;-)

Natural join can “under-match”

- What if two attributes don't have the same *name* and we do want them to match?
- Example: Suppose we want **aName** and **director** to match! (which Artists are also Directors?)
- Solution?

Theta Join

- It's common to use σ to check conditions after a Cartesian product.

$$\sigma_{condition} (R \times S)$$

- **Theta Join** makes this easier.

- Notation: $R \bowtie_{condition} S$

- The result is

- the same as Cartesian product (not natural join!) followed by select.

- In other words, $R \bowtie_{condition} S = \sigma_{condition} (R \times S)$.

Theta Join

$$R \bowtie_{condition} S$$

- The word “theta” has no special connotation. It is an artifact of a definition in an early paper.
- You still have to write out the *conditions*, since they are not inferred.
- Exercise: Use Theta join to find **artists** who have also *directed* movies. Display artist name, movie title.

$\Pi_{aName, title} (Artists \bowtie_{Artists.aName=Movies.director} Movies)$



Outer Joins

Dangling tuples

- If a tuple in one relation has **no match** in the other, natural join leaves that tuple **out**.
- Example schema:
 - People(phone, name, address)
 - Donations(phone, charity, amount, date)

Phone is a foreign key referencing People.
- What if **someone** has made **no** donations?
- Natural join leaves out those tuples, but you may want them!

Outer Join

- Outer Join leaves those tuples in, and adds *null* values whenever no value exists.
- Three variants:
 - Left Outer Join: only do this “null padding” for left operand
 - Right Outer Join: only do this null padding for right operand
 - Full Outer Join: pad both operands!
(see next slide for example)

Example of Left Outer Join

<i>Employee</i>		
Name	EmpId	DeptName
Harry	3415	Finance
Sally	2241	Sales
George	3401	Finance
Harriet	2202	Sales
Tim	1123	Executive

<i>Dept</i>	
DeptName	Manager
Sales	Harriet
Production	Charles

<i>Employee ⋈ Dept</i>			
Name	EmpId	DeptName	Manager
Harry	3415	Finance	ω
Sally	2241	Sales	Harriet
George	3401	Finance	ω
Harriet	2202	Sales	Harriet
Tim	1123	Executive	ω



Full Outer Join - Example

Employee

Name	EmpId	DeptName
Harry	3415	Finance
Sally	2241	Sales
George	3401	Finance
Harriet	2202	Sales
Tim	1123	Executive

Dept

DeptName	Manager
Sales	Harriet
Production	Charles

Employee \bowtie Dept

Name	EmpId	DeptName	Manager
Harry	3415	Finance	ω
Sally	2241	Sales	Harriet
George	3401	Finance	ω
Harriet	2202	Sales	Harriet
Tim	1123	Executive	ω
ω	ω	Production	Charles

Composing larger expressions
(plus a few new operators)

Assignment operator

- Notation:

$R := \text{Expression}$

- Alternate notation:

$R(A_1, \dots, A_n) := \text{Expression}$

- Lets you name all the attributes of the new relation
- Sometimes you don't want the name they would get from Expression.

- R must be a temporary variable, not one of the relations in the schema.

I.e., you are not updating the content of a relation!

■ Example:

- $\text{Temp1} := Q \times R$
- $\text{Temp2} := \sigma_{a=99}(\text{Temp1}) \times S$
- $\text{Answer}(\text{part}, \text{price}) := \pi_{b,c}(\text{Temp2})$

- Whether / how small to break things down is up to you. It's all for readability.
- Assignment helps us break a problem down
- It also allows us to change the *names* of relations [and attributes].
- *There is another way to rename things ...*

Rename operation

- Notation: $\rho_{RI}(R2)$
- Alternate notation: $\rho_{RI(A1, \dots, An)}(R2)$
 - Lets you rename all the *attributes* as well as the relation.
- Note that these are equivalent:
$$RI(A1, \dots, An) := R2$$
$$RI := \rho_{RI(A1, \dots, An)}(R2)$$
- ρ is useful if you want to rename *within* an expression.

Precedence

- Expressions can be composed recursively.
- Make sure attributes match as you wish.
 - It helps to annotate each subexpression, showing the attributes of its resulting relation.
- Parentheses and precedence rules define the **order** of evaluation.
- Precedence, from highest to lowest, is:
 - σ, π, ρ
 - \times, \bowtie
 - \cap
 - $\cup, -$

- Unless very sure, use brackets!


Breaking down expressions

- Complex nested expressions can be hard to read.
- Two alternative notations allow us to break them down:
 - Expression trees.
 - Sequences of assignment statements.

Expression Trees


- Leaves are relations.
- Interior nodes are operators.
- Exactly like representing arithmetic expressions as trees.

$3 + 7 * x$



- If interested, see Ullman and Widom, section 2.4.10.

Expression Trees – RA Example



Summary of operators

Operation	Name	Symbol
choose rows	select	σ
choose columns	project	π
combine tables	Cartesian product	\times
	natural join	\bowtie
	theta join	$\bowtie_{condition}$
rename relation [and attributes]	rename	ρ
assignment	assignment	$:=$

“Syntactic sugar”

- Some operations are not *necessary*.
 - You can get the same effect using a combination of other operations.
- Examples: natural join, theta join.
- We call this “syntactic sugar”.
- This concept also comes up in logic and programming languages.



Set operations

- Because relations are sets, we can use set **intersection**, **union** and **difference**.
- But only if the operands are relations over the **same** attributes (in number, name, and order).

Graduates

Number	Surname	Age
7274	Robinson	37
7432	O'Malley	39
9824	Darkes	38

Managers

Number	Surname	Age
9297	O'Malley	56
7432	O'Malley	39
9824	Darkes	38

Union

Graduates \cup Managers

Number	Surname	Age
7274	Robinson	37
7432	O'Malley	39
9824	Darkes	38
9297	O'Malley	56

Graduates

Number	Surname	Age
7274	Robinson	37
7432	O'Malley	39
9824	Darkes	38

Managers

Number	Surname	Age
9297	O'Malley	56
7432	O'Malley	39
9824	Darkes	38

Intersection

Graduates \cap Managers

Number	Surname	Age
7432	O'Malley	39
9824	Darkes	38

Set operations

- Because relations are sets, we can use set **intersection**, **union** and **difference**.
- But only if the operands are relations over the same attributes (in number, name, and order).

Graduates			Difference		
Number	Surname	Age			
7274	Robinson	37			
7432	O'Malley	39			
9824	Darkes	38			
Managers			Graduates - Managers		
Number	Surname	Age	Number	Surname	Age
9297	O'Malley	56	7274	Robinson	37
7432	O'Malley	39			
9824	Darkes	38			

Set operations

- Because relations are sets, we can use set **intersection**, **union** and **difference**.
- But only if the operands are relations over the same attributes (in number, name, and order).
- If the names or order *mismatch*?

Paternity		Maternity	
Father	Child	Mother	Child
Adam	Cain	Eve	Cain
Adam	Abel	Eve	Seth
Abraham	Isaac	Sarah	Isaac
Abraham	Ishmael	Hagar	Ishmael

Paternity \cup Maternity ???

Set operations

- Because relations are sets, we can use set **intersection**, **union** and **difference**.
- But only if the operands are relations over the same attributes (in number, name, and order).
- If the names or order *mismatch*?

Paternity		Maternity	
Father	Child	Mother	Child
Adam	Cain	Eve	Cain
Adam	Abel	Eve	Seth
Abraham	Isaac	Sarah	Isaac
Abraham	Ishmael	Hagar	Ishmael

$$\rho_{\text{Father} \rightarrow \text{Parent}}(\text{Paternity}) \cup \rho_{\text{Mother} \rightarrow \text{Parent}}(\text{Maternity})$$

Parent	Child
Adam	Cain
Adam	Abel
Abraham	Isaac
Abraham	Ishmael
Eve	Cain
Eve	Seth
Sarah	Isaac
Hagar	Ishmael

Example from: http://www.cs.toronto.edu/~faye/343/f07/lectures/wk3/03_RAAlgebra.pdf

Summary of techniques for writing queries in relational algebra

Approaching the problem

- Ask yourself which **relations** need to be involved. Ignore the rest.
- Every time you **combine** relations, confirm that
 - attributes that should match will be ***made*** to match and
 - attributes that will be made to match ***should*** match.
- Annotate each subexpression, to show the attributes of its resulting relation.

Breaking down the problem

- Remember that you must look one tuple at a time.
 - If you need info from two different tuples, you must make a new relation where it's in one tuple.
- Is there an *intermediate* relation that would help you get the final answer?
 - Draw it out with actual data in it.
- Use assignment to define those intermediate relations.
 - Use good names for the new relations.
 - Name the attributes on the LHS each time, so you don't forget what you have in hand.
 - Add a comment explaining exactly what's in the relation.

Specific types of query

- **Max** (min is analogous):
 - Not directly supported in relational algebra
 - Pair tuples and find those that are not the max.
 - Then subtract from all to find the maxes.
- **“k or more”**:
 - Make all combos of k different tuples that satisfy the condition.
- **“exactly k”**:
 - “k or more” - “(k+1) or more”.
- **“every”**:
 - Make all combos that should have occurred.
 - Subtract those that did occur to find those that didn't always.
These remaining are the **failures**.
 - Subtract the failures from **all** to get the answer.

Relational algebra wrap-up

RA is procedural

- An RA query itself suggests a procedure for constructing the result (*i.e.*, how one could implement the query).
- We say that it is “procedural.”

Evaluating queries

- Any problem has multiple RA solutions.
 - Each solution suggests a “query execution plan”.
 - Some may seem more efficient.
- But in RA, we won't care about efficiency; it's an algebra.
- In a DBMS, queries actually are executed, & efficiency matters!
 - Which query execution plan is most efficient depends on the data in the database and what indices you have.
 - Fortunately, the DBMS optimizes our queries.
 - We can focus on what we want, not how to get it.