

# SQL:

# Structured Query Language

## Data Manipulation Language (DML) – Part II

CSC343, Introduction to Databases

Nosayba El-Sayed (based on slides from Diane Horton)

Fall 2015



DCS50

# SQL Recap

**SELECT** sID, AVG(grade) as stdavg  
**FROM** Took  
**WHERE** sID > 22222  
**GROUP BY** sID  
**HAVING** AVG(grade) > 70  
**ORDER BY** AVG(grade) **DESC**

If you're not using **DISTINCT**  
or **GROUPING**, you can  
**ORDER BY**  
**UN-SELECTED** attributes.

Order of execution:

1. **FROM**
2. **ON**
3. **JOIN**
4. **OUTER**
5. **WHERE**
6. **GROUP BY**
7. **HAVING**
8. **SELECT**
9. **DISTINCT**
10. **ORDER BY**

# SQL Recap

```
SELECT sID
FROM Student
WHERE sID > 22222
ORDER BY cgpa DESC
```



Order of execution:

1. FROM
2. ON
3. JOIN
4. OUTER
5. WHERE
6. GROUP BY
7. HAVING
8. SELECT
9. DISTINCT
10. ORDER BY

If you're not using **DISTINCT**  
or **GROUPING**, you can  
**ORDER BY**  
**UN-SELECTED** attributes.

# SQL Joins

# The joins you know from RA

Expression within SQL statements	Meaning (RA)
<code>R, S</code>	$R \times S$
<code>R cross join S</code>	
<code>R natural join S</code>	$R \bowtie S$
<code>R join S on &lt;Condition&gt;</code>	$R \bowtie_{condition} S$
<code>R natural left [outer] join S</code>	$R \ltimes S$
<code>R natural right [outer] join S</code>	$R \ltimes S$
<code>R natural full [outer] join S</code>	$R \ltimes S$

# The joins you know from RA (Cont.)

Expression within SQL statements	Meaning (RA)
<code>R left [outer] join S on &lt;cond&gt;</code>	$R \bowtie_{condition} S$
<code>R right [outer] join S on &lt;cond&gt;</code>	$R \ltimes_{condition} S$
<code>R full [outer] join S on &lt;cond&gt;</code>	$R \Join_{condition} S$

# Dangling tuples

- With joins that require some attributes to match, tuples lacking a match are left out of the results.
- We say that they are “dangling”.
- An **outer join** preserves dangling tuples by padding them with **NULL** in the other relation.
- A join that doesn't pad with **NULL** is called an **inner join**.

# Example

R

A	B
1	2
4	5

S

B	C
2	3
6	7

R NATURAL LEFT JOIN S

A	B	C
1	2	3
4	5	NULL



# Example

R

A	B
1	2
4	5

S

B	C
2	3
6	7

R NATURAL RIGHT JOIN S

A	B	C
1	2	3
NULL	6	7

# Example

R

A	B
1	2
4	5

S

B	C
2	3
6	7

R NATURAL FULL JOIN S

A	B	C
1	2	3
4	5	NULL
NULL	6	7

# In practice natural join is dangerous

- Attributes with matching **names** don't necessarily mean matching **meanings**!
- Having **implicit** comparisons impairs **readability**.
- Also: if the schema changed, a query that *looks* fine may actually be broken, without being able to tell.
- **Best practise: Don't use natural join.**

# Summary of join expressions

- Cartesian product

A cross join B

same as A, B

- Theta-join

A join B on C

no padding of tuples

A {left|right|full} join B on C

padding

- Natural join

A natural join B

no padding of tuples

A natural {left|right|full} join B

padding

# Keywords INNER and OUTER

- There are keywords `INNER` and `OUTER`, but you never need to use them.
- Your intentions are clear anyway:
  - You get an `OUTER` join `iff` you use the keywords `LEFT`, `RIGHT`, or `FULL`.
  - If you don't use the keywords `LEFT`, `RIGHT`, or `FULL` you get an `INNER` join.

# Impact of having null values



# Missing Information

- Two common scenarios:
  - **Missing value.**  
E.g., we know a student has some *email* address, but we don't know what it is.
  - **Inapplicable attribute.**  
E.g., the value of attribute *spouse* for an unmarried person.

# Representing missing information

- One possibility: use a special value as a placeholder. E.g.,
  - If age unknown, use -1.
  - If StNum unknown, use 999999999.
- Pros and cons?
- Better solution: use a value not in any domain. We call this a **null value**.
- Tuples in SQL relations can have **NULL** as a value for one or more components.



# Checking for null values

- You can compare an attribute value to `NULL` with

- `IS NULL`

- `IS NOT NULL`

- Example:

```
SELECT *  
FROM Course  
WHERE breadth IS NULL;
```

- Note: do not use `WHERE breadth = NULL;`

# Impact of **null** values on SQL **expressions**?

- Arithmetic expressions?

Assume:  $x$  is **NULL**

- Result is always **NULL**
- Example:  $(x + \text{grade}) \rightarrow \text{NULL}$
- Even if 'grade' is 0!  
i.e.  $(x * 0) \rightarrow \text{NULL}$
- Also:  $(x - x) \rightarrow \text{NULL}$

- Comparison operators? ( $>$ ,  $<$ ,  $=$ , ...)

- $(x < 32) \rightarrow \text{UNKNOWN}$
- Result is **UNKNOWN**
- This **UNKNOWN** is a *truth-value*!
- Truth-values in SQL are: (**TRUE**, **FALSE**, **UNKNOWN**)



- What if we have: `WHERE (x<32 OR name='Alex')`?

# Evaluating Logic Expressions with UNKNOWN

- Logic: TRUE, FALSE, UNKNOWN
  - UNKNOWN OR FALSE → UNKNOWN
  - UNKNOWN OR TRUE → TRUE
  - UNKNOWN AND TRUE → UNKNOWN
  - UNKNOWN AND FALSE → FALSE
  - NOT UNKNOWN → UNKNOWN

- In SQL
  - A tuple is in a query result **iff** the result of the WHERE clause is TRUE
  - Demo: *where-null*

## Ternary logic tricks:

TRUE = 1  
FALSE = 0  
UNKNOWN =  $\frac{1}{2}$

AND =  $\min(\dots)$   
OR =  $\max(\dots)$   
NOT =  $1-x$

# Thinking of the truth-values as numbers

A	B	as nums	A and B	min	A or B	max
T	T	1, 1	T	1	T	1
TF or FT		1, 0	F	0	T	1
F	F	0, 0	F	0	F	0
TU or UT		1, 0.5	U	0.5	T	1
FU or UF		0, 0.5	F	0	U	0.5
U	U	0.5, 0.5	U	0.5	U	0.5

# Impact of null values on aggregation

- “Aggregation ignores **NULL**.”
- **NULL** never contributes to a sum, average, or count, and can never be the minimum or maximum of a column (unless every value is **NULL**).
- If **ALL** values are **NULL** in a column, then the result of the aggregation is **NULL**.
  - Exception: **COUNT** of an empty set is 0.

# Impact of null values on aggregation

- COUNT()

- COUNT(R.\*) = 2

- COUNT(R.x) = 1

- COUNT(S.\*) = 1

- COUNT(S.x) = 0

- COUNT(T.\*) = 0

- COUNT(T.x) = 0

**R**

x
NULL
1

- Other aggregations (e.g. MIN/MAX)

- MIN(R.x) = 1

- MAX(R.x) = 1

- MIN(S.x) = NULL

- MAX(S.x) = NULL

- MIN(T.x) = NULL

- MAX(T.x) = NULL

**S**

x
NULL

**T** x

# Summary

	<u>Some</u> nulls in A	<u>All</u> nulls in A
<code>min(A)</code>	ignore the nulls	NULL
<code>max(A)</code>		
<code>sum(A)</code>		
<code>avg(A)</code>		
<code>count(A)</code>		0
<code>count(*)</code>	all tuples count	

*\*Demo*

## Subqueries





# Subqueries in a FROM clause

- Instead of a **relation** name in the **FROM** clause, we can use a **subquery**.
- The subquery must be parenthesized.
- Must name the result, so you can refer to it in the outer query.

# Example

- What does this do?

```
SELECT sid, dept||cnum as course, grade
FROM Took,
    (SELECT *
     FROM Offering
     WHERE instructor='Horton') Hofferings
WHERE Took.oid = Hofferings.oid;
```

- This FROM is analogous to:

$\text{Took} \times \rho_{\text{Hofferings}} (\llbracket \text{subquery} \rrbracket)$

- Can you suggest another version?

# Subquery as a value in a **WHERE**

- If a subquery is guaranteed to produce exactly one tuple, then the subquery can be used as a *value*.
- Simplest situation: that one tuple has only one attribute.

# Example

- Find all students with a cgpa greater than that of student 99999.

```
SELECT sid, surname  
FROM Student  
WHERE cgpa >  
      (SELECT cgpa  
       FROM Student  
       WHERE sid = 99999);
```

# Special cases

- What if the subquery returns **NULL**?
  - Evaluates to UNKNOWN, tuple not returned
- What if the subquery could return more than one value?
- When a subquery can return multiple values, we can make comparisons using a quantifier:
  - $\text{cgpa} > \text{at least one of them}$  (ANY)
  - $\text{cgpa} > \text{all of them}$  (ALL)

# The Operator ANY

- Syntax:

$x$  «comparison» ANY («subquery»)

or equivalently

$x$  «comparison» SOME («subquery»)

- Semantics:

Its value is true iff the comparison holds for at least one tuple in the subquery result, i.e.,

$\exists y \in \text{«subquery results»} \mid x \text{ «comparison» } y$

- $x$  can be a list of attributes,  
but this feature is not supported by psql.

# The Operator ALL

- Syntax:

$x \text{ «comparison» ALL («subquery»)$

- Semantics:

Its value is true iff the comparison holds for every tuple in the subquery result, i.e.,

$\forall y \in \text{«subquery results»} \mid x \text{ «comparison» } y$

- $x$  can be a list of attributes,  
but this feature is not supported by psql.

# The Operator IN

- Syntax:  
     $x \text{ IN } (\text{«subquery»})$
- Semantics:  
    Its value is true iff  $x$  equals at least one of the tuples in the subquery result.
- $x$  can be a list of attributes, and psql does support this feature.



# Example – Q2 in Class Exercises

What does this do?

```
SELECT sid, dept || cnum AS course, grade
FROM Took NATURAL JOIN Offering
WHERE
    grade >= 80 AND
    (cnum, dept) IN (
        SELECT cnum, dept
        FROM Took NATURAL JOIN Offering
            NATURAL JOIN Student
        WHERE surname = 'Lakemeyer' );
```

## Q3 in Class Exercises

Suppose we have tables  $R(a, b)$  and  $S(b, c)$ .

1. What does this query do?

```
SELECT a  
FROM R  
WHERE b IN (SELECT b FROM S);
```

2. Can we express this query without using IN?

# The Operator EXISTS

- Syntax:  
EXISTS («*subquery*»)
- Semantics:  
Its value is true iff the subquery **has** at least one tuple.

# Example (Q4) : NOT EXISTS

What does this do?

```
SELECT instructor
FROM Offering Off1
WHERE NOT EXISTS (
    SELECT *
    FROM Offering
    WHERE
        oid <> Off1.oid AND
        instructor = Off1.instructor );
```

Instructors who have  
exactly one offering!

# Scope

- Queries are evaluated from the inside out.
  - If a name might refer to more than one thing, use the most closely nested one.
  - If a subquery refers only to names defined inside it, it can be evaluated **once** and used repeatedly in the outer query.
  - If it refers to any name defined outside of itself, it must be evaluated **once for each tuple in the outer query**.
- These are called **correlated subqueries**.

# Renaming can make scope explicit

```
SELECT instructor
FROM Offering Off1
WHERE NOT EXISTS (
    SELECT *
    FROM Offering Off2
    WHERE
        Off2.oid <> Off1.oid AND
        Off2.instructor = Off1.instructor );
```

# Summary: where subqueries can go

- As a *relation* in a **FROM** clause.
- As a *value* in a **WHERE** clause.
- With **ANY**, **ALL**, **IN** or **EXISTS** in a **WHERE** clause.
- As operands to **UNION**, **INTERSECT** or **EXCEPT**.
- Reference: textbook, section 6.3.

# Q5

- For each **course** find the **instructor** who has taught **the most offerings** of it. If there are ties, include them all.  
Report the course (eg csc343), instructor, and number of offerings of the course by that instructor.  
**Use one or more views to hold intermediate steps.**

## ROSI Schema

Students(sID, surName, campus)

Courses(dept, cNum, name, breadth)

Offerings(oID, dept, cNum, term, instructor)

Took(sID, oID, grade)