

# SQL

## Data Definition Language (DDL)

CSC343, Introduction to Databases

Nosayba El-Sayed (based on slides from Diane Horton)

Fall 2015



DCS50

# Modifying a Database

# Midterm next week

- Time: Oct 27 6-7pm (Be there 10 min early)
- Location: Brennan Hall (BR200)
- Material: everything up to Embedded SQL (so embedded sql is not included)
- Practice *writing*:
  - RA Queries
  - SQL Queries + Aggregates + Subqueries
  - Views
- Understand the rest of the material well, of course!

# Database Modifications

- Queries return a relation.
- A **modification** command does not; it **changes** the database in some way.
- Three kinds of modifications:
  - Insert a tuple or tuples.
  - Delete a tuple or tuples.
  - Update the value(s) of an existing tuple or tuples.

# Two ways to insert

- We've already seen two ways to insert tuples into an empty table:

`INSERT INTO «relation» VALUES «list of tuples» ;`

`INSERT INTO «relation» («subquery») ;`

- These can also be used to add tuples to a non-empty table.

## Naming attributes in INSERT

- Sometimes we want to insert tuples, but we don't have values for *all* attributes.
- If we name the attributes we *are* providing values for, the system will use **NULL** or a default for the rest.
- Convenient!

## Example

```
CREATE TABLE Invite (  
    name TEXT,  
    campus TEXT DEFAULT 'StG',  
    email TEXT,  
    age INT);  
  
INSERT INTO Invite VALUES  
( 'Mark', 'StG', 'm@m.com', 18 );  
  
INSERT INTO Invite(name, email)  
( SELECT firstname, email  
    FROM Student  
    WHERE cgpa > 3.4 ) ;
```

Here, name and email get values from the query, **campus** gets the *default* value, and **age** gets **NULL**.

# Deletion

- Delete tuples satisfying a condition:

```
DELETE FROM «relation»  
WHERE «condition»;
```

- Delete all tuples:

```
DELETE FROM «relation»;
```



# Example 1: Delete Some Tuples

```
DELETE FROM Course
WHERE NOT EXISTS (
    SELECT *
    FROM Took JOIN Offering
              ON Took.oid = Offering.oid
    WHERE
        grade > 50 AND
        Offering.dept = Course.dept AND
        Offering.cnum = Course.cnum
);
```

# Updates

- To change the value of certain attributes in certain tuples to given values:

UPDATE *«relation»*

SET *«list of attribute assignments»*

WHERE *«condition on tuples»* ;

## Example: update one tuple

- Updating one tuple:

```
UPDATE Student  
SET campus = 'UTM'  
WHERE sid = 999999;
```

- Updating several tuples:

```
UPDATE Took  
SET grade = 50  
WHERE grade >= 47 and grade < 50;
```

Types

# Table attributes have types

- When creating a table, you must define the **type** of each attribute.
- Analogous to declaring a variable's type in a program. Eg, “int num;” in Java or C.
- Some programming languages don't require type declarations. Eg, Python.
- Pros and cons?
- Why are type declarations required in SQL?

# Built-in types

- `CHAR ( n )`: fixed-length string of *n* characters. Padded with blanks if necessary.
- `VARCHAR ( n )`: variable-length string of *up to* *n* characters.
- `TEXT`: variable-length, unlimited. Not in the SQL standard, but psql and others support it.
- `INT = INTEGER`
- `FLOAT = REAL`
- `BOOLEAN`
- `DATE`; `TIME`; `TIMESTAMP` (date plus time)

# Values for these types

- Strings: `'Shakespeare's Sonnets'`  
Must surround with single quotes.
- INT: `37`
- FLOAT: `1.49`, `37.96e2`
- BOOLEAN: `TRUE`, `FALSE`
- DATE: `'2011-09-22'`
- TIME: `'15:00:02'`, `'15:00:02.5'`
- TIMESTAMP: `'Jan-12-2011 10:25'`

# And much more

- These are all defined in the SQL standard.
- There is much more, e.g.,
  - specifying the precision of numeric types
  - other formats for data values
  - more types
- For what psql supports, see chapter 8 of the documentation.



# User-defined types

- Defined in terms of a built-in type.
- You make it more specific by defining **constraints** (and perhaps a **default** value).

- Examples:

```
create domain Grade as int  
    default null  
    check (value>=0 and value <=100);
```

```
create domain Campus as varchar(4)  
    default 'StG'  
    check (value in ('StG', 'UTM', 'UTSC'));
```

# Semantics of type constraints

- Constraints on a type are checked every time a **value** is assigned to an **attribute** of that type.
- You can use these to create a powerful type system.

# Semantics of default values

- The *default* value for a **type** is used when no value has been specified.
- Useful! You can run a query and insert the resulting tuples into a relation -- even if the query does not give values for all attributes.
- Table **attributes** can also have default values.
- The difference:
  - **attribute default**: for that one attribute in that one table
  - **type default**: for every attribute defined to be of that type

# Keys and Foreign Keys

# Key constraints

- Declaring that a set of one or more attributes are the **PRIMARY KEY** for a relation means:
  - they form a **key** (unique, and no subset is)
  - their values will **never be null** (you don't need to separately declare that)
- Big hint to the DBMS: optimize for searches by this set of attributes!
- Every table must have 0 or 1 primary key.
  - A table can have no primary key, but in practise, every table should have one. Why?
  - You cannot declare more than one primary key.  
(Think of the PK as the *identity* of a row..)

# Declaring primary keys

- For a single-attribute key, can be part of the attribute definition.

```
create table Blah (  
    ID integer primary key,  
    name varchar(25));
```

- Or can be at the end of the table definition. (This is the only way for multi-attribute keys.) The brackets are required.

```
create table Blah (  
    ID integer,  
    name varchar(25),  
    primary key (ID));
```

# Uniqueness constraints

- Declaring that a set of one or more attributes is **UNIQUE** for a relation means:
  - they form a **key**
  - their values *can* be **null**; if they mustn't, you need to separately declare that
- You can declare more than one set of attributes to be **UNIQUE**.

# Declaring UNIQUE

- If only one attribute is involved, can be part of the attribute definition.

```
create table Blah (  
    ID integer unique,  
    name varchar(25));
```

- Or can be at the end of the table definition.  
(This is the only way if multiple attributes are involved.) The brackets are required.

```
create table Blah (  
    ID integer,  
    name varchar(25),  
    unique (ID));
```



# We saw earlier how nulls affect “unique”

- For uniqueness constraints, no two nulls are considered equal.

- E.g., consider:

```
create table Testunique (  
    first varchar(25),  
    last varchar(25),  
    unique(first, last))
```

- This would prevent two insertions of  
( 'Stephen', 'Cook' )
- But it would allow two insertions of  
( null, 'Rackoff' )

This can't occur with a **primary key**. Why not?

# Foreign key constraints

- Eg in table Took:  
`foreign key (sID) references Student`
- Means that attribute sID in this table is a foreign key that references the **primary key** of table Student.
  - Every value for sID in this table must actually occur in the Student table.
- Requirements:
  - Must be declared either **primary key** or **unique** in the “home” table (i.e. table “Student” in the above example)

# Declaring foreign keys

- Again, declare with the attribute (only possible if just a single attribute is involved) or as a separate table element.
- Can reference attribute(s) that are not the primary key as long as they are unique; just name them.

```
create table People (  
    SIN integer primary key,  
    name text,  
    OHIP text unique);  
  
create table Volunteers (  
    email text primary key,  
    OHIPnum text references People(OHIP));
```


# Enforcing foreign-key constraints

- Suppose there is a foreign-key constraint from relation R to relation S.
- How/when can the DBMS ensure that:
  1. the referenced attributes are **PRIMARY KEY** or **UNIQUE**?
  2. the values actually exist?
- What could cause a violation?
  - Example: a row is deleted from **Course**; **Offering** is now referring to a course that doesn't exist.
- You get to define what the DBMS should do.
- This is called specifying a “**reaction policy**.”

# Reaction Policies

What's the DBMS *default*?

Can we *change* it? How?



# Example

- Suppose  $T = \text{Took}$  and  $S = \text{Student}$ .
  1. What sorts of action must simply be rejected?
    - Inserting in  $T$  a student whose SID is not in  $S$
  2. How about a deletion or update of a Student whose sid occurs in  $\text{Took}$ ?

# Possible policies

- **cascade**: propagate the change to the referring table
  - E.g. if a *Student* leaves university, delete all their referrals in *Took*!
- **set null**: set the referring attribute(s) to null
- If you say nothing, the *default* is to **forbid** the change in the referred-to table.



# Reaction policy example

- In the University schema, what should happen in these situations:
  - csc343 changes number to be 543
  - student 99132 is deleted
  - student 99132's grade in csc148 is raised to 85.
  - csc148 is deleted



# Note the asymmetry

- Suppose table **R** refers to table S.
- You can define “fixes” that propagate changes backwards from S to **R**.
- (You define them in table **R** because it is the table that will be **affected**.)
- You cannot define fixes that propagate forward from **R** to S.

# Syntax for specifying a reaction policy

- Add your reaction policy where you specify the foreign key constraint.

- Example:

```
create table Took (  
    ...  
    foreign key (sID) references Student  
        on delete cascade  
    ...  
);
```

# What you can react to

- Your reaction policy can specify what to do either
  - `on delete`, i.e., when a deletion creates a dangling reference,
  - `on update`, i.e., when an update creates a dangling reference,
  - **or both.** Just put them one after the other.

Example:

`on delete restrict on update cascade`

# What your reaction can be

- Your policy can specify one of these reactions (there are others):
  - `restrict`: Don't allow the deletion/update.
  - `cascade`: Make the same deletion/update in the referring tuple.
  - `set null`: Set the corresponding value in the referring tuple to null.

# Semantics of Deletion

- What if deleting one tuple affects the outcome for a tuple encountered later?
- To prevent such interactions, deletion proceeds in two stages:
  - Mark all tuples for which the WHERE condition is satisfied.
  - Go back and delete the marked tuples.

# Example

- In the **Guesses** table, some students made more than one guess:
  - Who are these students?
  - Delete their guesses from the table!

```
csc343h-nosayba=> select * from guesses;
```

number	name	guess
1	Cole	365
2	Avery	500
3	Sam	502
4	Madeleine	390
5	Cole	450
6	Michael	1000
7	Mackenzie	700
8	Mackenzie	701

(8 rows)

## Other Constraints and Assertions

# “check” constraints

- We’ve seen a check clause on a user-defined domain:

```
create domain Grade as smallint
    default null
    check (value >= 0 and value <= 100);
```

- You can also define a check constraint
  - on an **attribute**
  - on the **tuples** of a relation
  - *across* **relations**



# Attribute-based “check” constraints

- Defined with a single attribute and constrain its value (in every tuple).
- Can only refer to that attribute.
- Can include a subquery.

- Example:

```
create table Student (  
    sID integer,  
    program varchar(5) check  
        (program in (select post from P)),  
    firstName varchar(15) not null, ...);
```

- Condition can be anything that could go in a WHERE clause.

# When they are checked

- Only when a tuple is **inserted** into that relation, or its value for that attribute is **updated**.
- If a change **somewhere** else **violates** the constraint, the DBMS will not notice. E.g.,
  - If a student's program changes to something not in table P, we get an error.
  - But if table **P drops a program** that some student has, there is **no** error.

# “not null” constraints

- You can declare that an attribute of a table is NOT NULL.

```
create table Course(  
    cNum integer,  
    name varchar(40) not null,  
    dept Department,  
    wr boolean,  
    primary key (cNum, dept));
```

- In practise, many attributes should be `not null`.
- This is a very specific kind of attribute-based constraint.

# Tuple-based “check” constraints

- Defined as a separate element of the table schema, so can refer to *any* attributes of the table.
- Again, condition can be anything that could go in a **WHERE** clause, and can include a subquery.

- Example:

```
create table Student (  
    sID integer,  
    age integer,  
    year integer,  
    college varchar(4),  
    check (year = age - 18),  
    check college in  
        (select name from Colleges));
```

# When they are checked

- Only when a tuple is **inserted** into that relation, or **updated**.
- Again, if a change somewhere else violates the constraint, the DBMS will not notice.

# How nulls affect “check” constraints

- A check constraint only fails if it evaluates to **false**.
- It is not picky like a WHERE condition.
- E.g.: `check (age > 0)`

age	Value of condition	CHECK outcome	WHERE outcome
19	TRUE	pass	pass
-5	FALSE	fail	fail
NULL	unknown	pass	fail

# Example

- Suppose you created this table:

```
create table Tester(  
    num integer,  
    word varchar(10),  
    check (num>5));
```

- It would allow you to insert (null, 'hello')
- If you need to prevent that, use a “not null” constraint.

```
create table Tester(  
    num integer not null,  
    word varchar(10),  
    check (num>5));
```

# Naming your constraints

- If you name your constraint, you will get more helpful error messages.
- This can be done with any of the types of constraint we've seen.
- Add

`constraint «name»`

before the

`check ( «condition» )`



# Examples

```
create domain Grade as smallint
    default null
    constraint gradeInRange
        check (value >= 0 and value <= 100));
```

```
create domain Campus as varchar(4)
    not null
    constraint validCampus
        check (value in ('StG', 'UTM', 'UTSC'));
```

```
create table Offering(...
    constraint validCourseReference
    foreign key (cNum, dept) references Course);
```

# Exercise

- Create table Student. sID is primary key.  
firstName, surName cannot be null.  
campus can only be StG, UTM, or UTSc.

```
CREATE TABLE Student (  
    sID INTEGER,  
    surName VARCHAR(25) NOT NULL,  
    firstName VARCHAR(25) NOT NULL,  
    campus VARCHAR(5)  
    check (campus in ('StG', 'UTM', 'UTSC')),  
    email VARCHAR(30),  
    cgpa FLOAT,  
    PRIMARY KEY (sID)  
);
```

- Order of constraints doesn't matter, and doesn't dictate the order in which they're checked.

# Assertions

- **Check** constraints apply to an attribute or table. They can't express constraints *across* tables, e.g.,
  - Every loan has at least one customer, who has an account with at least \$1,000.
  - For each branch, the sum of **all loan amounts** < the sum of all account balances.
- **Assertions** are schema elements at the top level, so *can* express cross-table constraints:  
`create assertion (<name>) check (<predicate>);`

# Powerful but costly

- SQL has a fairly **powerful** syntax for expressing the predicates, including quantification.
- Assertions are **costly** because
  - They have to be checked upon every database update (although a DBMS may be able to limit this).
  - Each check can be expensive.
- Testing and maintenance are also difficult.
- So assertions must be used with great care.

# Triggers

- Assertions are powerful, but costly.
- Check constraints are less costly, but less powerful.
- **Triggers** are a compromise between these extremes:
  - They are cross-table constraints, as powerful as assertions.
  - But you control the cost by having control over **when they are applied**.



# The basic idea

- You specify three things.
  - **Event:** Some type of database action, e.g.,  
after delete on Courses  
or  
before update of grade on Took
  - **Condition:** A boolean-valued expression, e.g.,  
when grade > 95
  - **Action:** Any SQL statements, e.g.,  
insert into Winners values (sID, course)

Using SQL “schemas”



# Schema: a kind of namespace


- “psql csc343h-mark” connects you to a database called csc343h-mark.  
(Substitute your cdf userid of course.)
- Everything defined (tables, types, etc.) goes into one big pot.
- Schemas let you create different namespaces.
- Useful for logical organization, and for avoiding name clashes.



# Creating a schema

- You already have a schema called “public”.
- You can also create your own. Example:  
`create schema University;`
- To refer to things inside a particular schema, you can use dot notation:

```
create table University.Student (...);  
select * from University.Student;
```



# When you don't use dot notation

- If you refer to a name without specifying what schema it is within:
  - Any new names you define go in the schema called “public”
  - E.g., if you create a table called `frindle`, you actually are defining `public.frindle`.
  - When referring to a name, there is a search path that finds it.

# The search path

- To see it the search path:  
`show search_path;`
- You can set the search path yourself. Example:  
`set search_path to University, public;`
- The default search path is: `"$user", public`
  - schema `"$user"` is not created for you, but if you create it, it's at the front of the search path.
  - schema `public` is created for you.

# Removing a schema

- Easy:  
`drop schema University cascade;`
- “`cascade`” means everything inside it is dropped too.
- To avoid getting an error message if the schema does not exist, add “`if exists`”.

# Usage pattern

- You can use this at the top of every DDL file:

```
drop schema if exists University cascade;  
create schema University;  
set search_path to University;
```

- Helpful during development, when you may want to change the schema, or test queries under different conditions.

# Updating the schema itself

- Alter: alter a domain or table  
`alter table Course`  
    `add column numSections integer;`  
`alter table Course`  
    `drop column breadth;`
- Drop: remove a domain, table, or whole schema  
`drop table course;`
- How is that different from this?  
`delete from course;`
- If you drop a table that is referenced by another table, you must specify “cascade”
- This removes all referring rows.